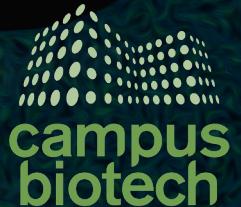


# DeepLabCut tutorial

JAX course

October 9<sup>th</sup>, 2022

Mathis Group  
computational neuroscience & AI



Alexander Mathis  
École Polytechnique Fédérale de Lausanne  
Swiss Federal Institute of Technology

Today's sunrise at JAX.

Happy DeepLabCutting!



# How to use DeepLabCut?



<https://github.com/DeepLabCut/DeepLabCut-Workshop-Materials>

# DeepLabCut workflow

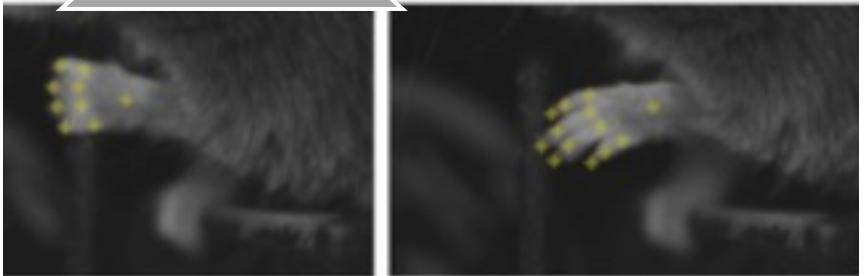
## Train DNN

Create a project,  
extract frames, +  
GUIs to label your data

Select + Train your  
deep neural network

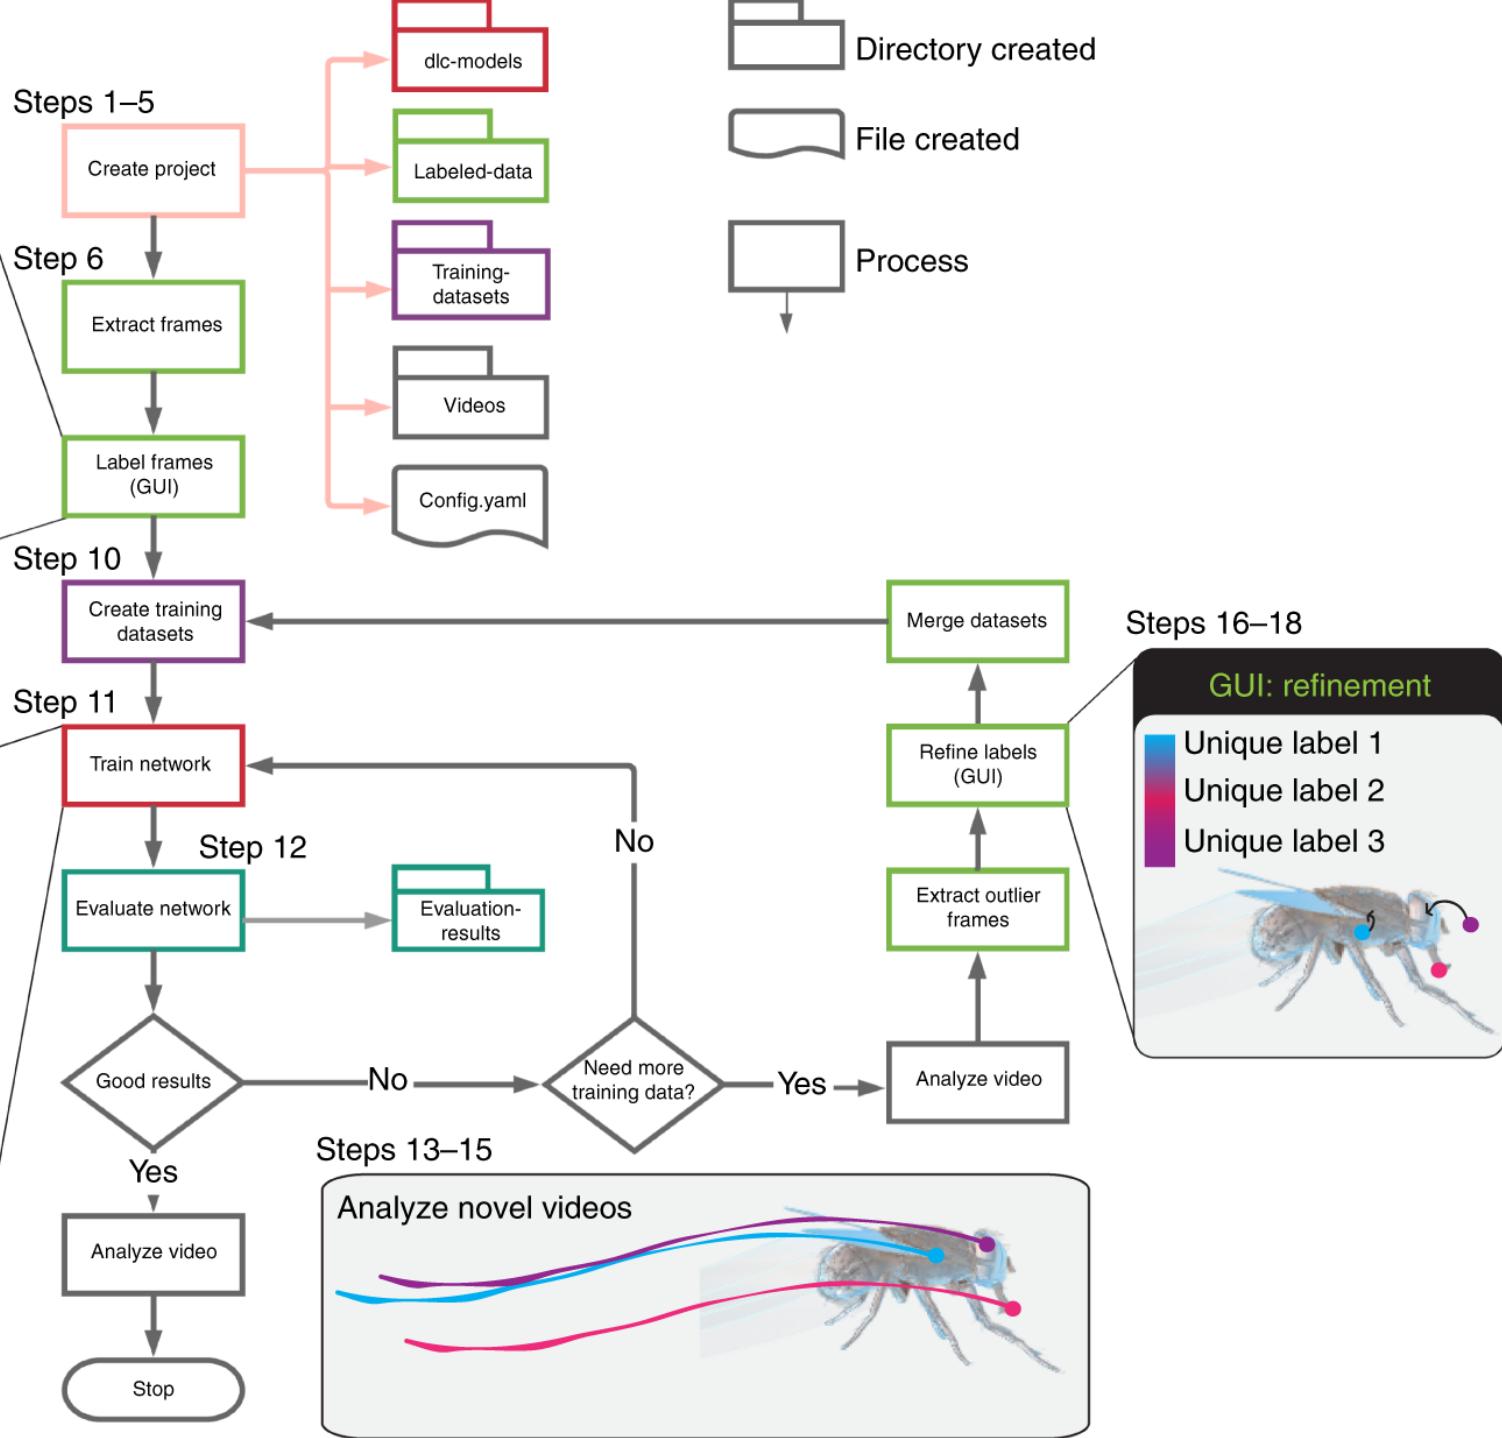
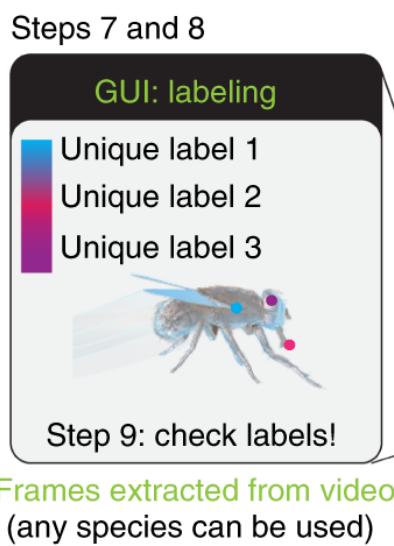
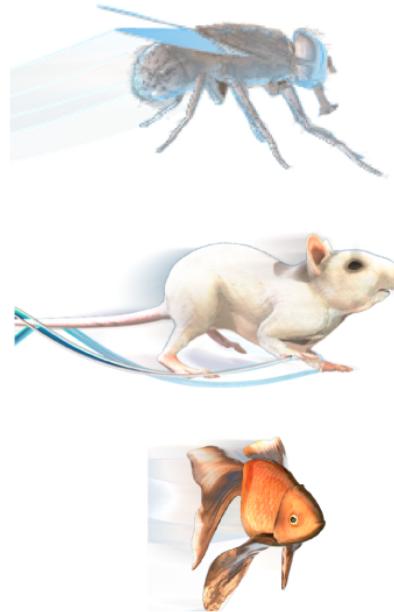
Evaluate network  
performance  
(active learning + GUIs  
if improvement needed)

Run inference on  
new videos,  
create labeled videos,  
+ plot your results!



refine?





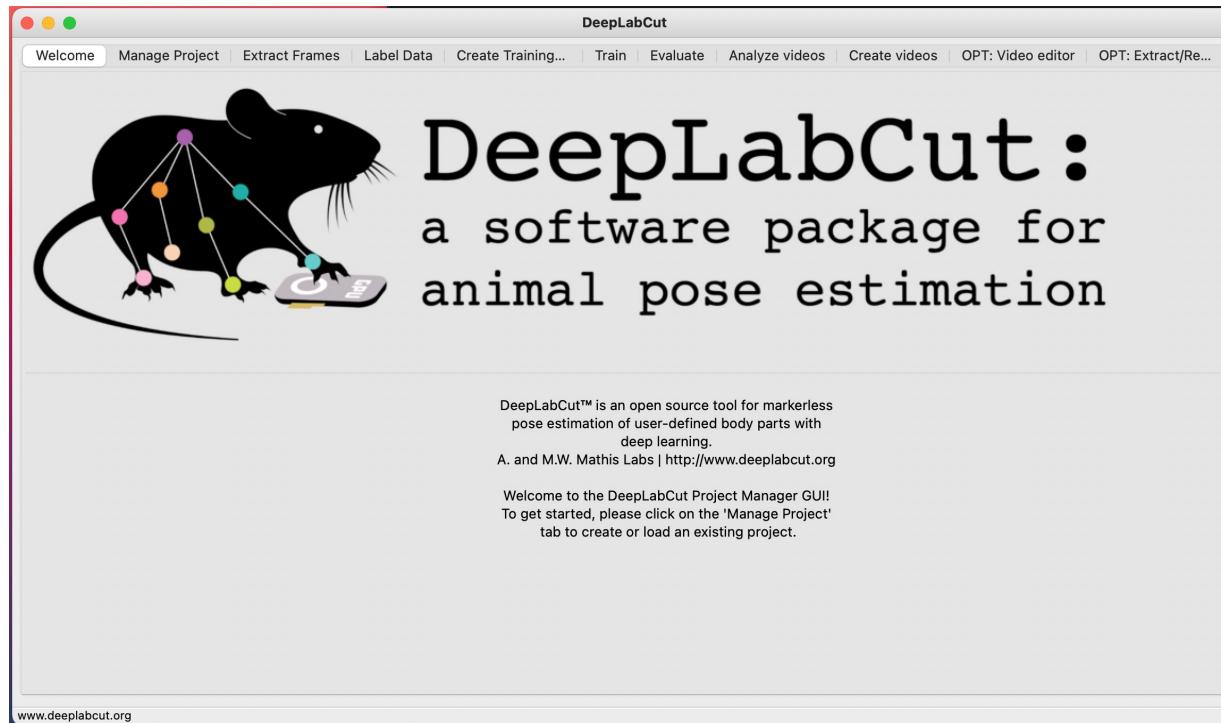
# Using DeepLabCut for 3D markerless pose estimation across species and behaviors

Tanmay Nath<sup>1,5</sup>, Alexander Mathis<sup>1,2,5</sup>, An Chi Chen<sup>3</sup>, Amir Patel<sup>3</sup>, Matthias Bethge<sup>4</sup> and Mackenzie Weygandt Mathis<sup>1\*</sup>

We are working on an expanded protocol that also includes maDLC  
But also see the online docs (which are already up-to-date) →  
<https://deeplabcut.github.io/DeepLabCut/README.html>

# Project GUI demo!

We have a user-friendly GUI that allows you to carry out the whole workflow.



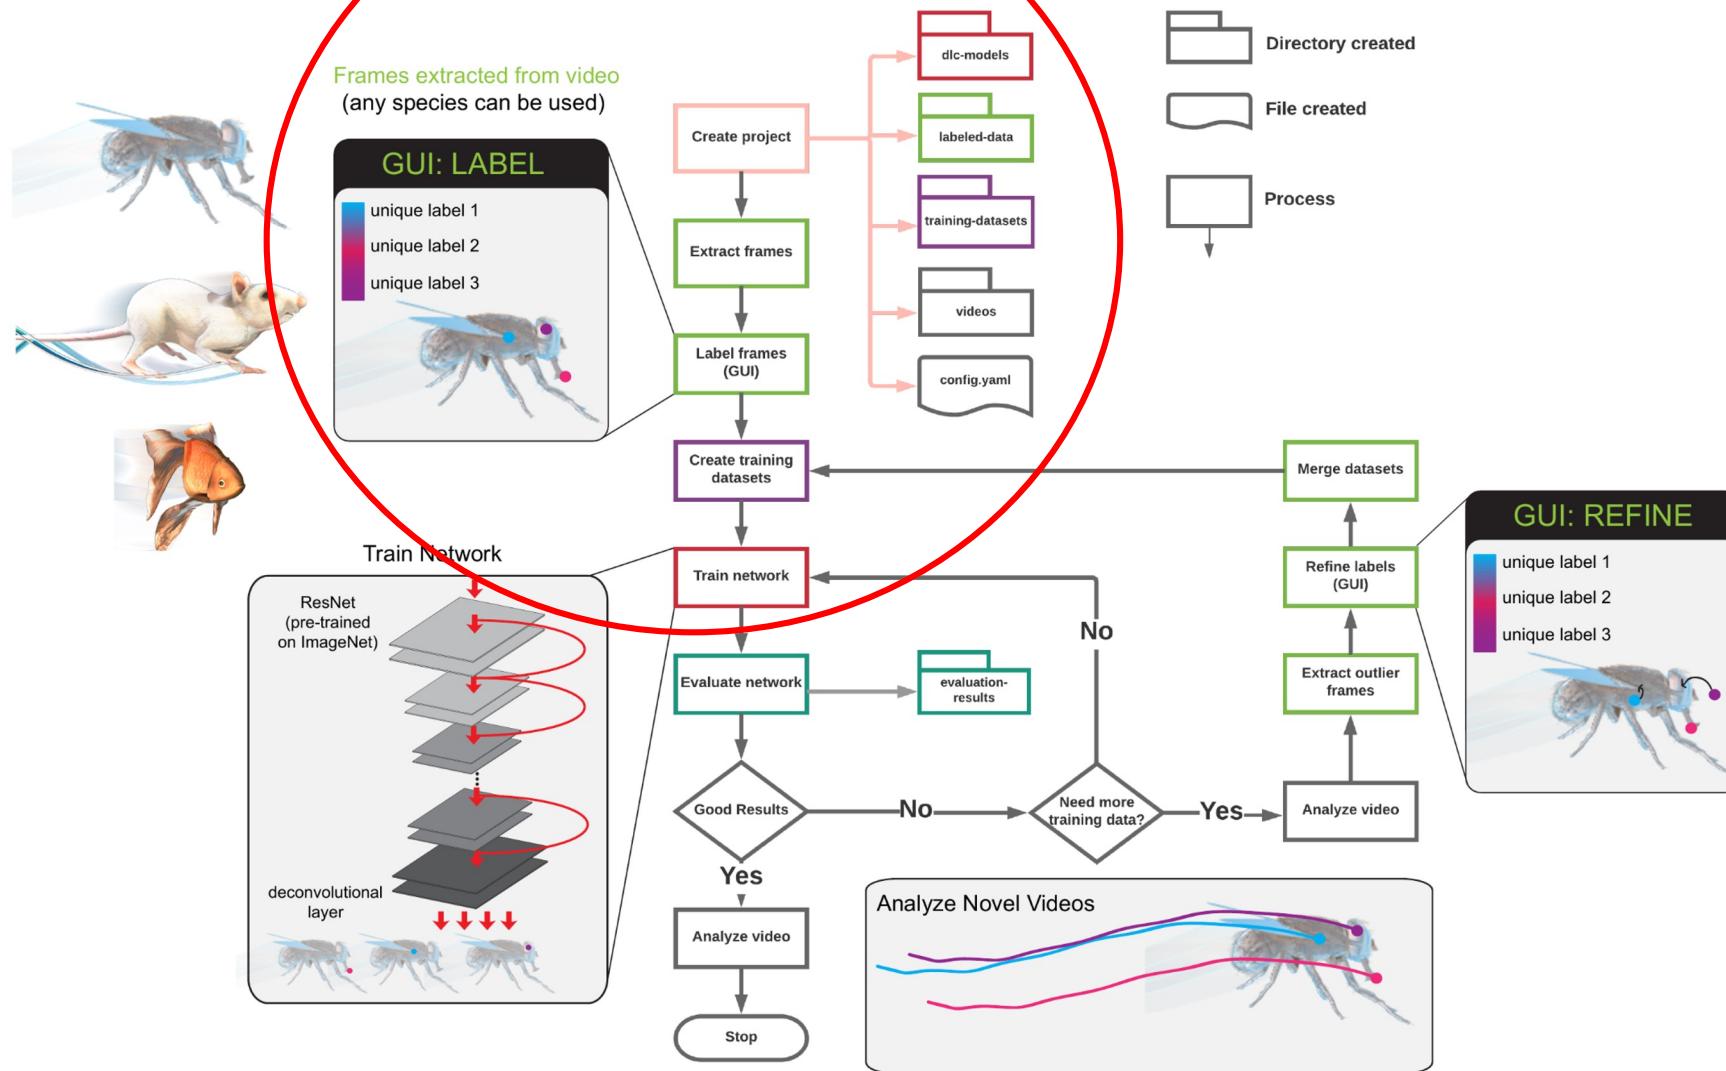
Note completely updated GUI forthcoming in DLC version 2.3; rc1 already out.

Nath\*, Mathis\* et al. Nature Protocols

# Key commands

| Operation                                   | Command  |
|---|--|
| Open IPython and import DeepLabCut (Step 1) | <code>ipython<br/>import deeplabcut</code>   |
| Create a new project (Step 2)               | <code>deeplabcut.create_new_project('project_name', 'experimenter', ['path of video 1', 'path of video2', ...])</code> |

# DeepLabCut 2.0 workflow



## 2D Project Folders



- ▶ training data,  
trained networks,  
`config.yaml` file

This is the master folder created when  
you create a project (Step 1)

## New (2D) videos for analysis



- ▶ batch process videos
- ▶ analyzed data

you can place new videos here, and then run:  
`deeplabcut.analyze_videos(config_path,  
folderpath, videotype='mp4')`

Your 2D or 3D project “entry point” is through the `config.yaml` file  
When you want to work on your project:

```
activate DLCenvName
ipython
import deeplabcut
config_path = '/home/yourprojectfolder/config.yaml'
```

# Frame extraction: deeplabcut.extract\_frames(..)

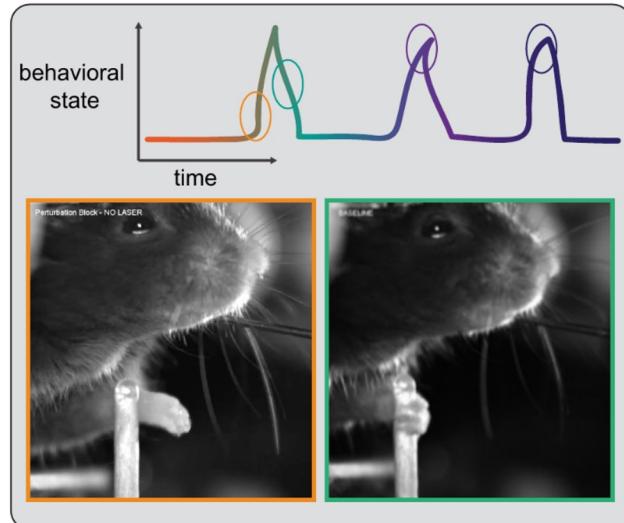
Select videos to grab frames:

Use videos with image from:

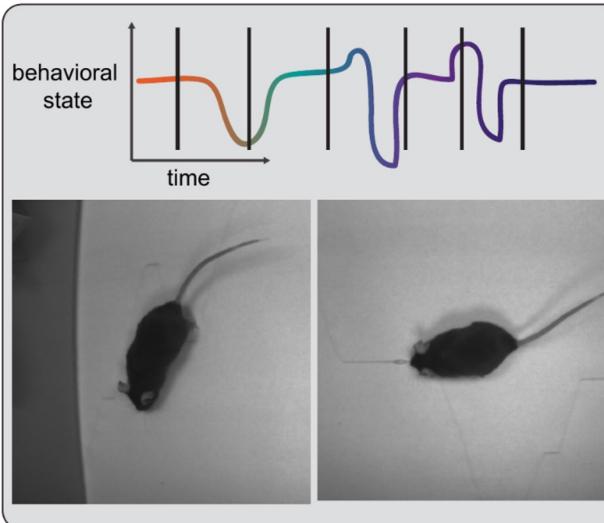
- different sessions reflecting (if the case) varying light conditions, backgrounds, setups, and camera angles (etc).
- different individuals, especially if they look different (i.e. brown + black mice)

3 methods for frame extraction to create a labeled train/test set

Image based clustering (k-means)



Random temporal sampling (uniform)

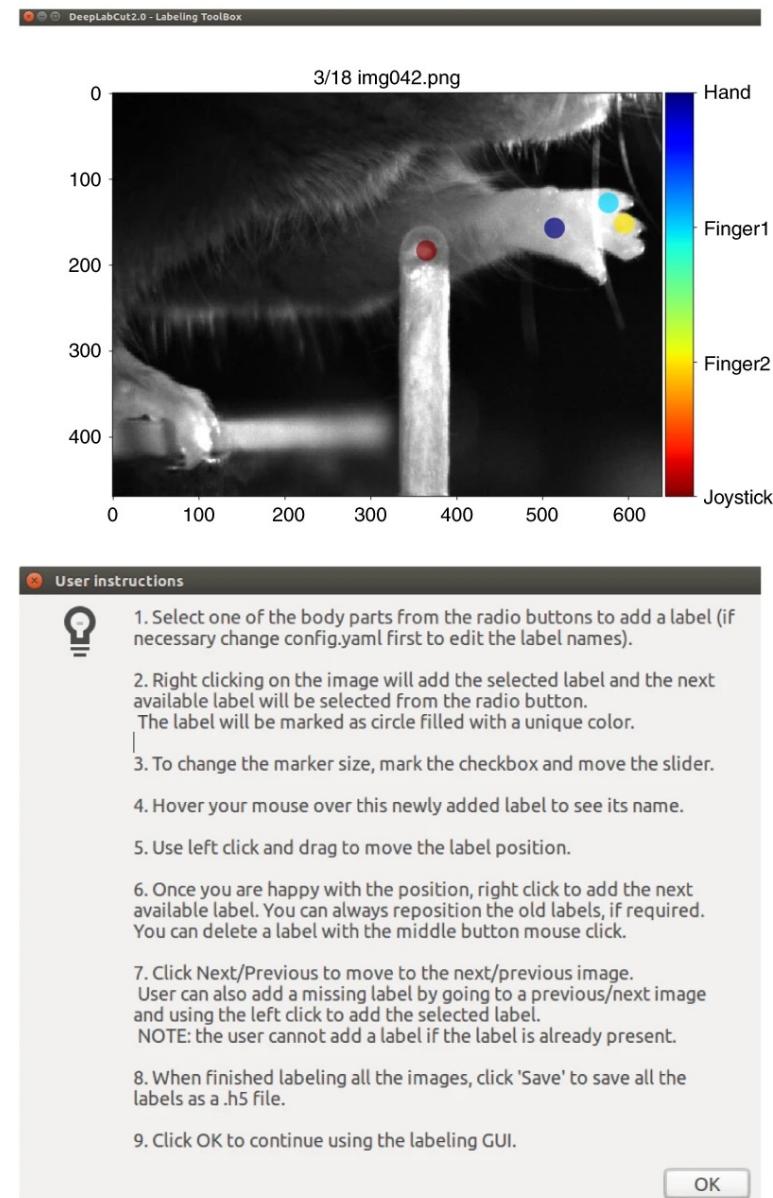
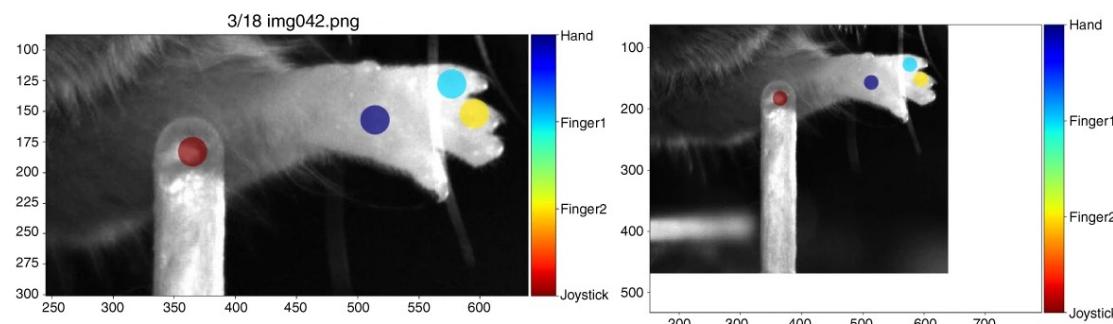
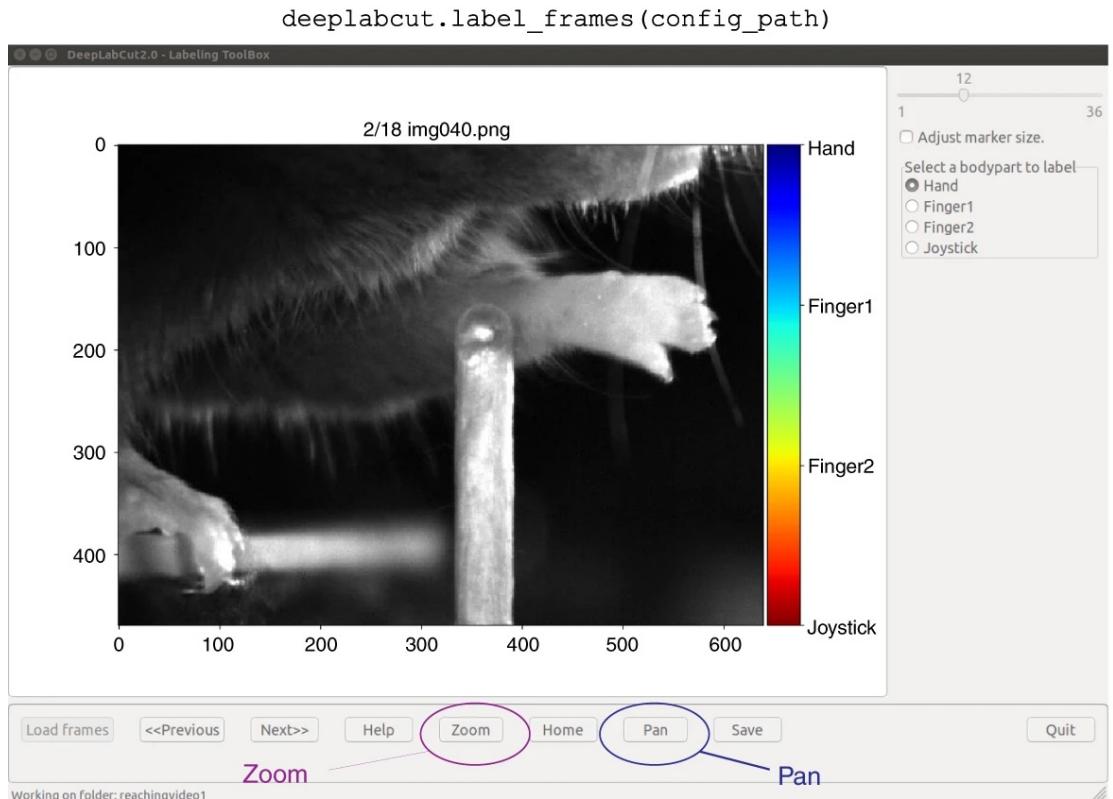


GUI for manual frame grabbing



# Label frames: deeplabcut.label\_frames(..)

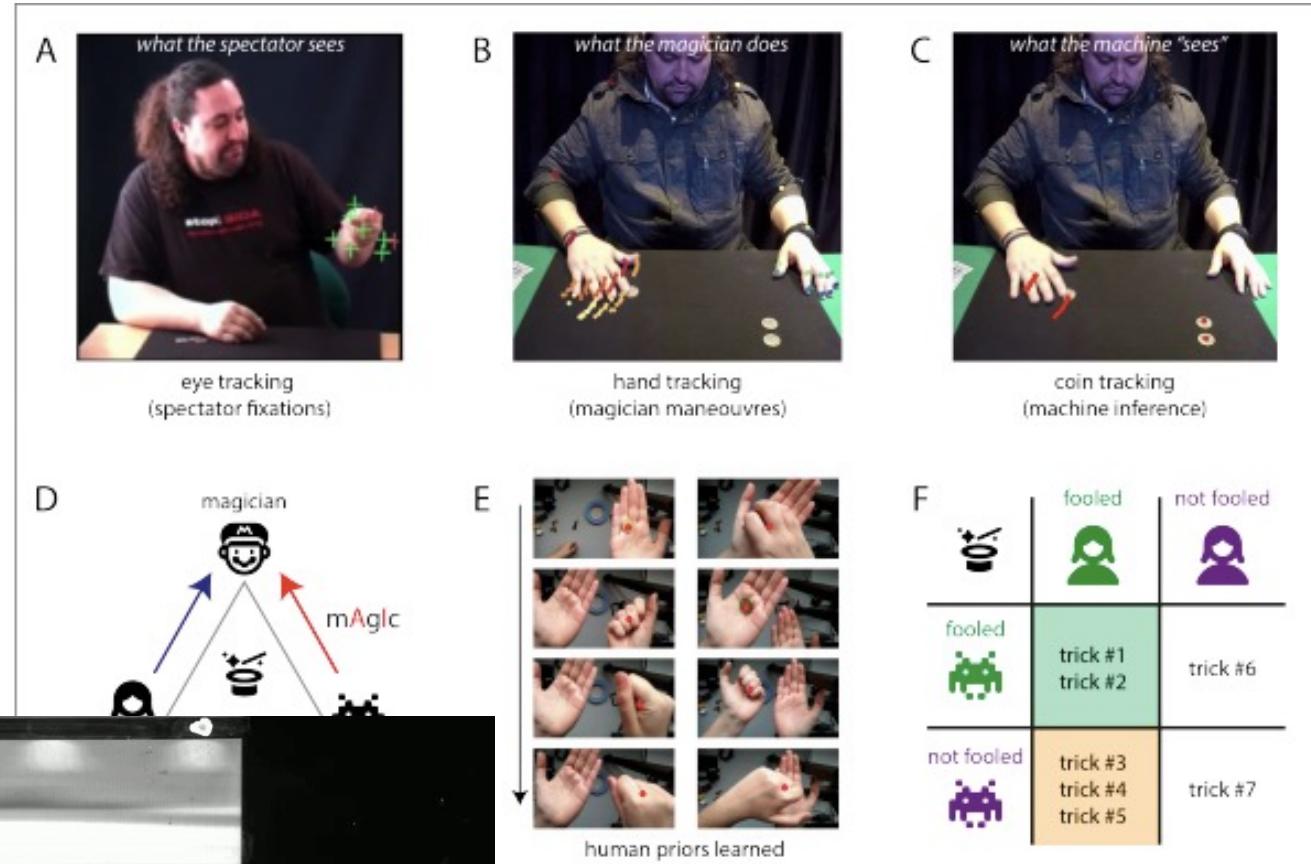
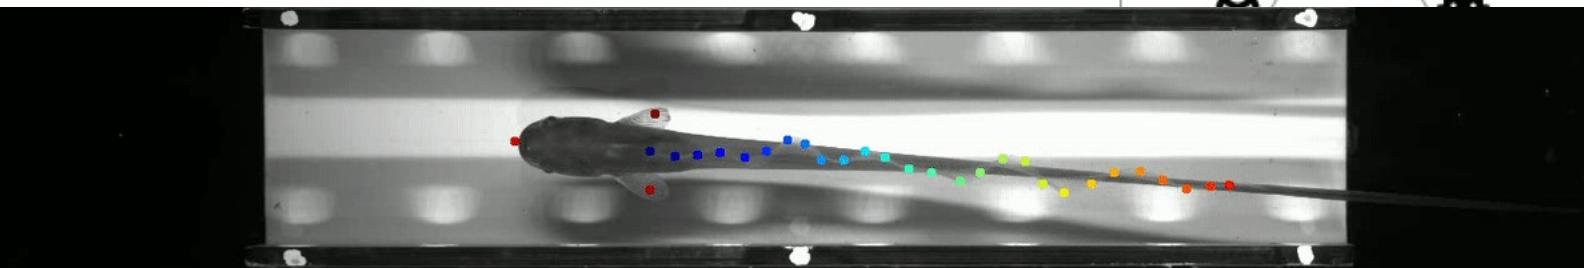
Label frames using the interactive GUI:



# What bodyparts can be detected?

DLC has very large receptive fields (due to deep ResNet), so many visual features + geometric layout can be considered:

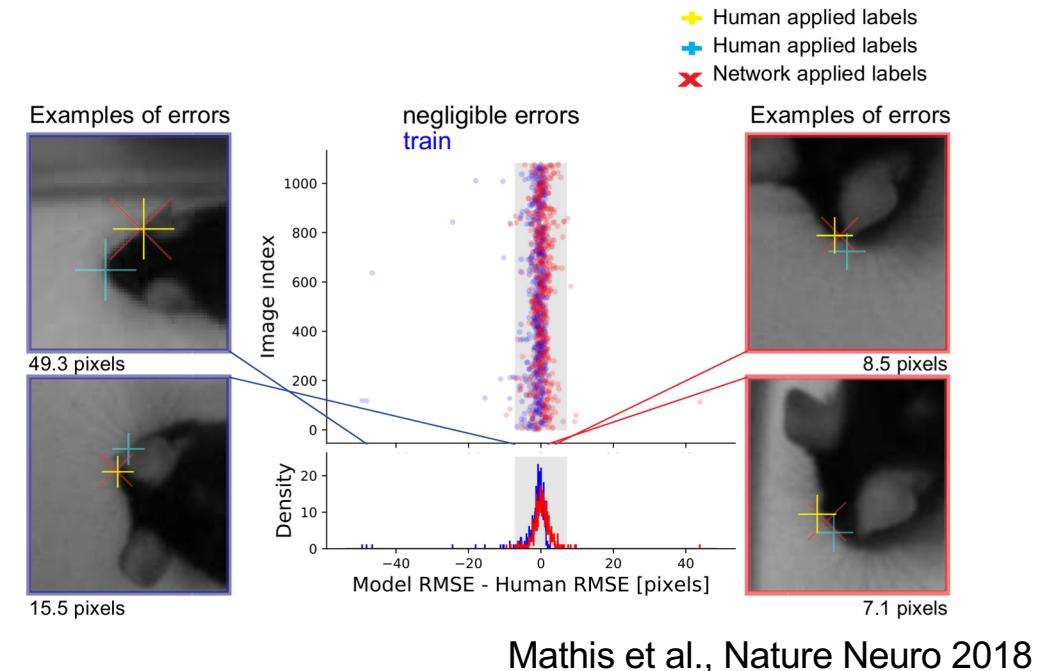
- Salient points (e.g. snout/tail base)
- Texture of objects
- Can teach the network to guess
- Geometric relationship
  - (e.g. left/right ear; whiskers!)



Playing magic tricks to deep neural networks  
untangles human deception  
by Zaghi-Lara et al. arxiv 2019

# Labeling considerations

- Label (feature locations) **consistently**
- Errors are bad! (especially when only a few images are labeled)
- Left vs. right body parts can be dist. if full animal visible (due to large rec. fields), but flipped labels are problematic
- You can check labels by plotting (do this before training!)
- You can correct wrong labels manually with the labeling GUI



Mathis et al., Nature Neuro 2018

# What network & augmentation method?

## Network architecture:

- Network backbones: MobileNets V2, *ResNets*, *DLC-Rnet* & EfficientNets,...
- Scale (level of deconvolution) -> decrease stride for closely interacting animals  
and and when you want high resolution
- Differences: **Inference & training speed, resolution**

## Pre-training (Weights):

- *ImageNet pretrained (the generalist)* -> see Mathis et al. 2021 WACV
- MPII-pose pretrained
- **DeepLabCut model zoo**

## Training method:

- Learning rate & optimizer (Adam/SGD, batch\_size, ...)
- **Augmentation transformations (rotation, scaling, ...)**

# Online demo ... let's start training

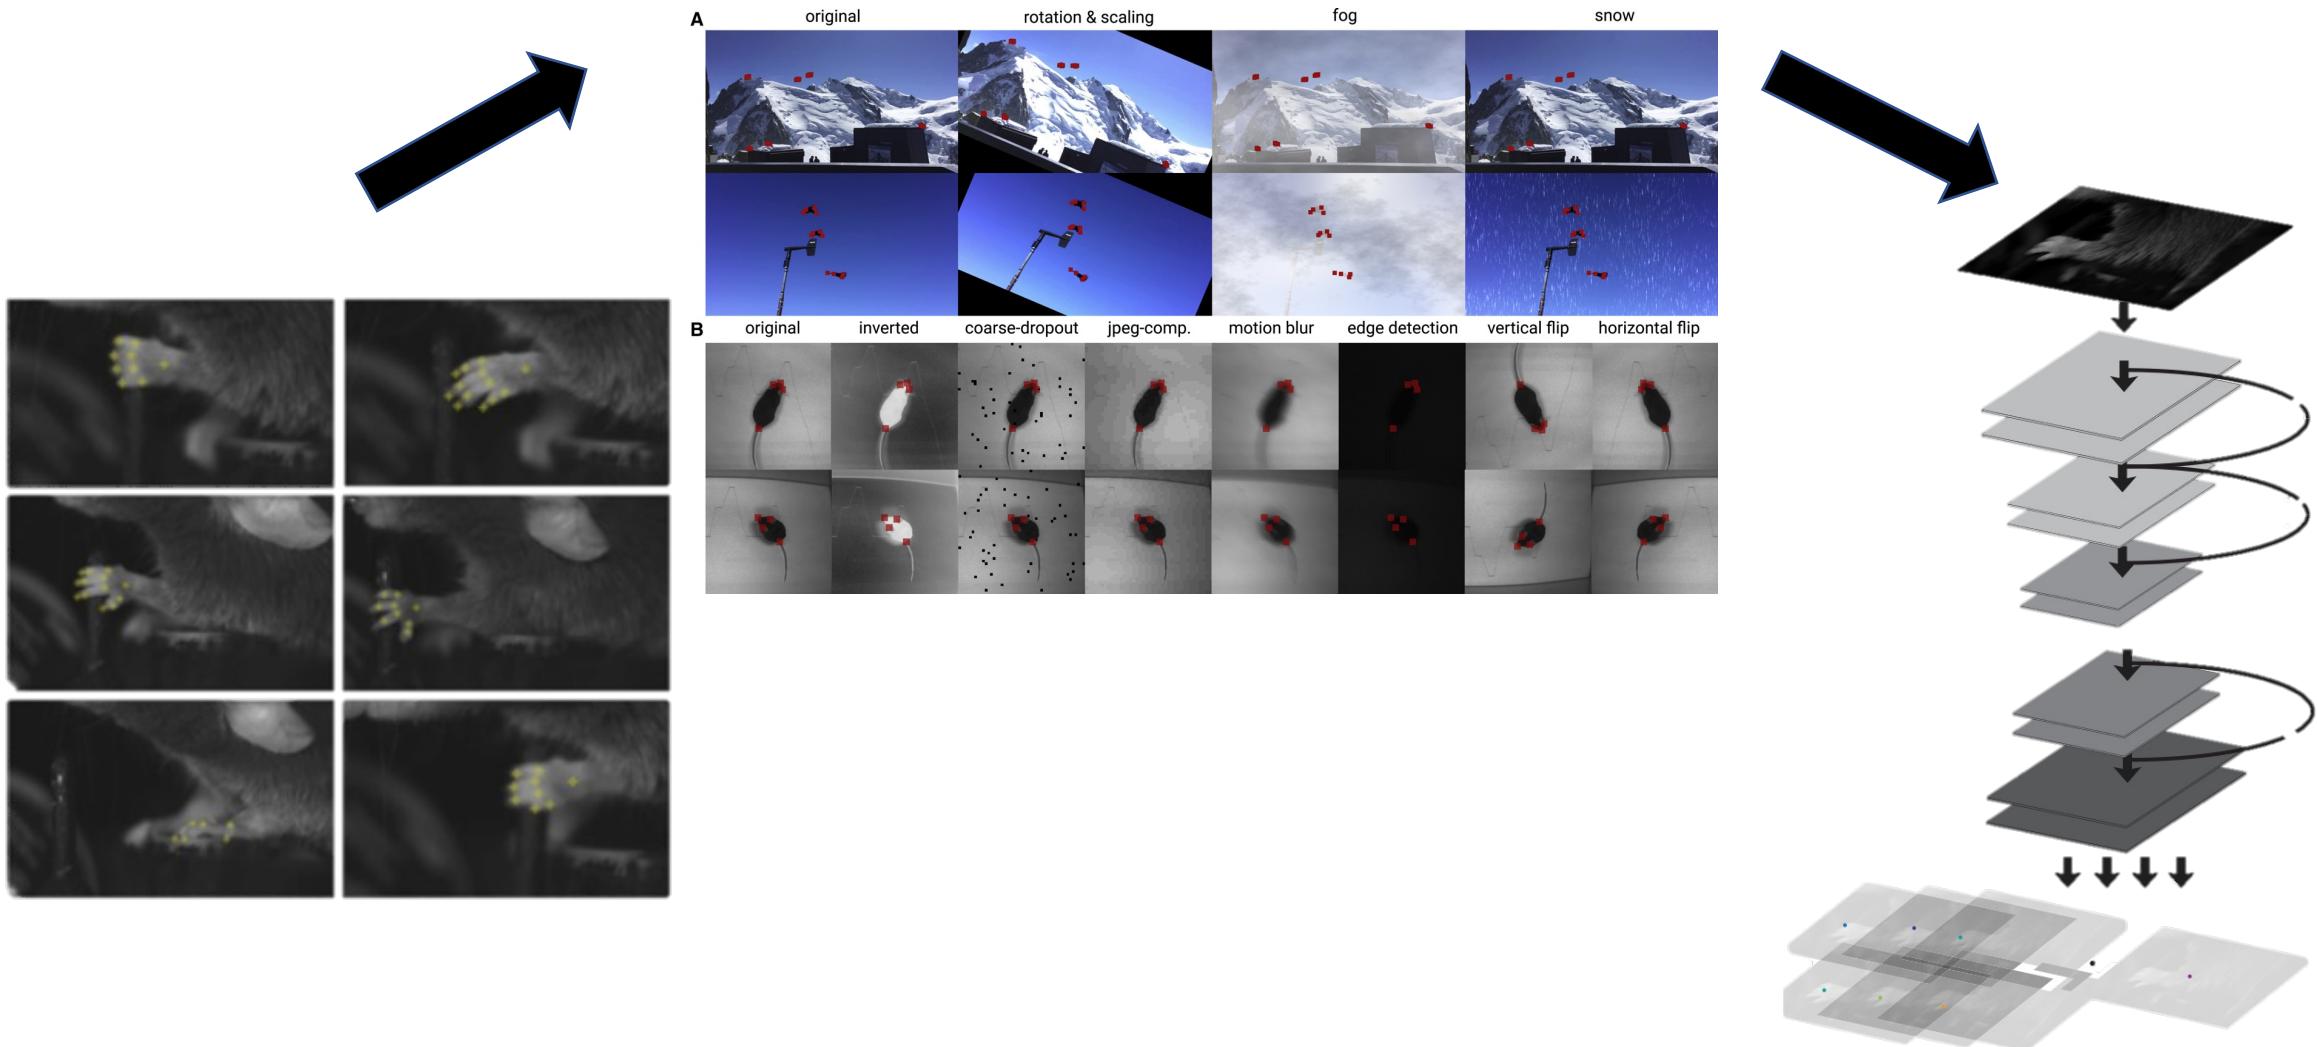
DEMO notebooks running on Google's resources can be found here:

<https://github.com/DeepLabCut/DeepLabCut/blob/master/examples/COLAB>

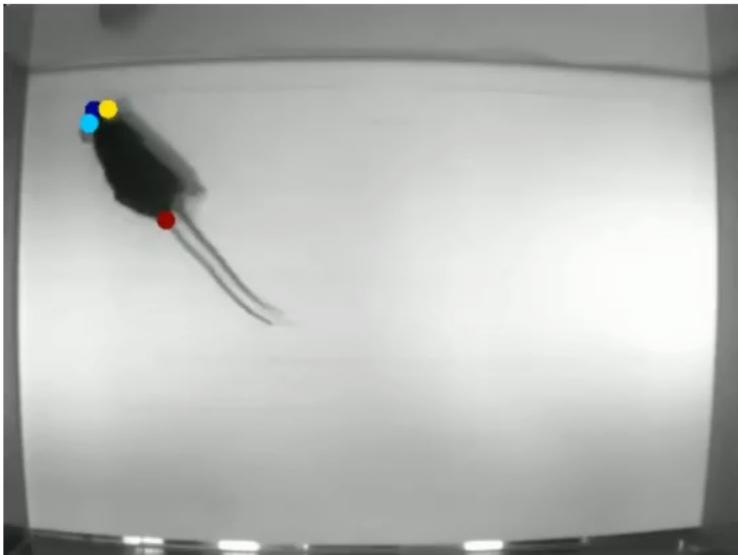
Contains examples for

- 3 mouse tracking with maDLC
- How to use a model from the model Zoo for inference
- How to go through all the steps of a DLC project
- And many more ...

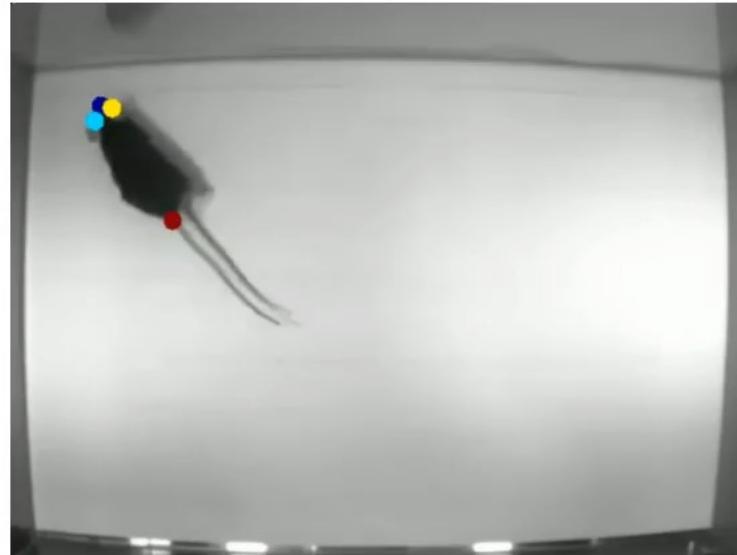
# Augmentation



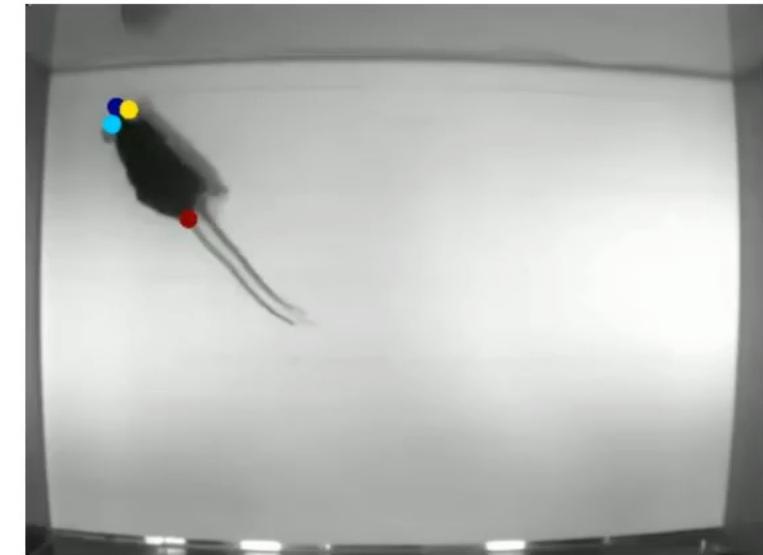
# Data Augmentation: how to get the most out of your data!



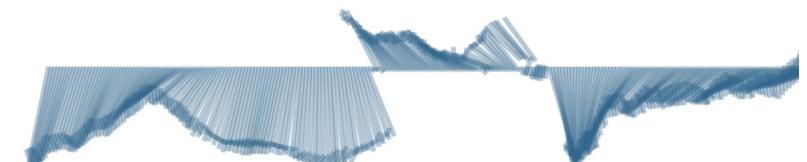
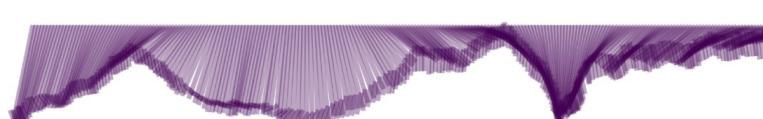
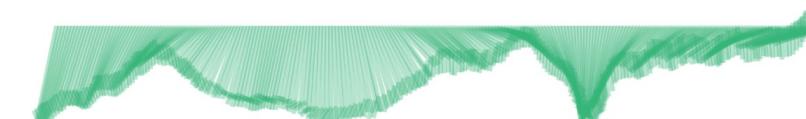
tensorpack



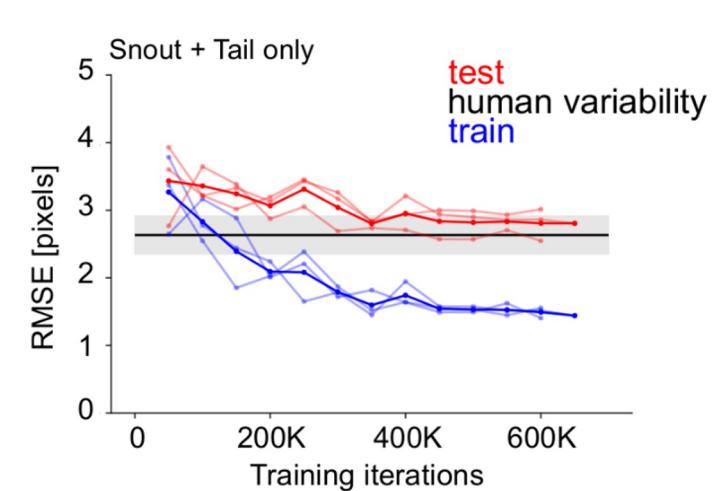
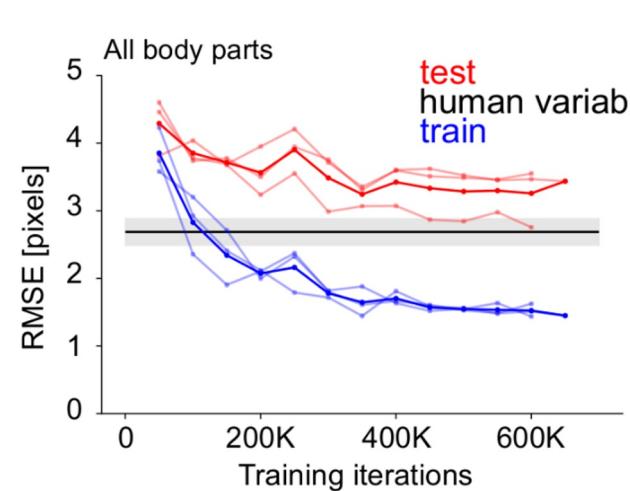
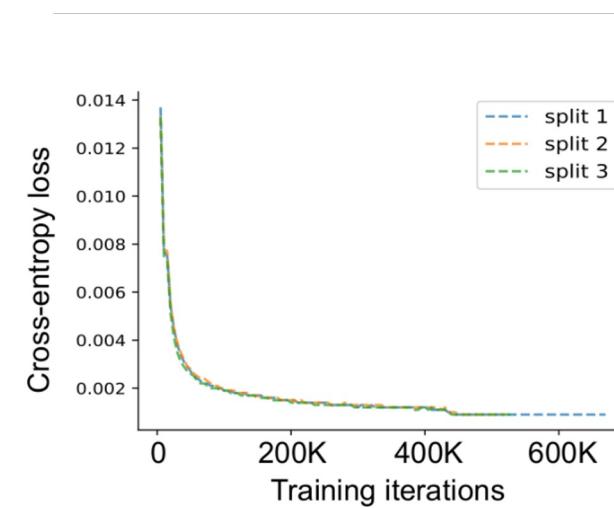
imgaug



scale-crop



# Evaluation: Compare test RMSE and ideally your labeling error!



# Evaluation => look at test images!

a

Example of the terminal output

```
deeplabcut.evaluate_network('<path of the proj. config file>', shuffle = [1], plotting=True)

Assessing accuracy of shuffle # 1 with 99 % training fraction.
Found the following training snapshots: [(200000, 0)]
You can choose among those for analysis of train/test performance.
Results for 200000 training iterations: 99 1 train error: 4.81 pixels. Test error: 7.02 pixels.
With putoff of 0.1 train error: 3.3 pixels. Test error: 5.2 pixels.
```

b

Examples of **training images** with Human and DeepLabCut labels



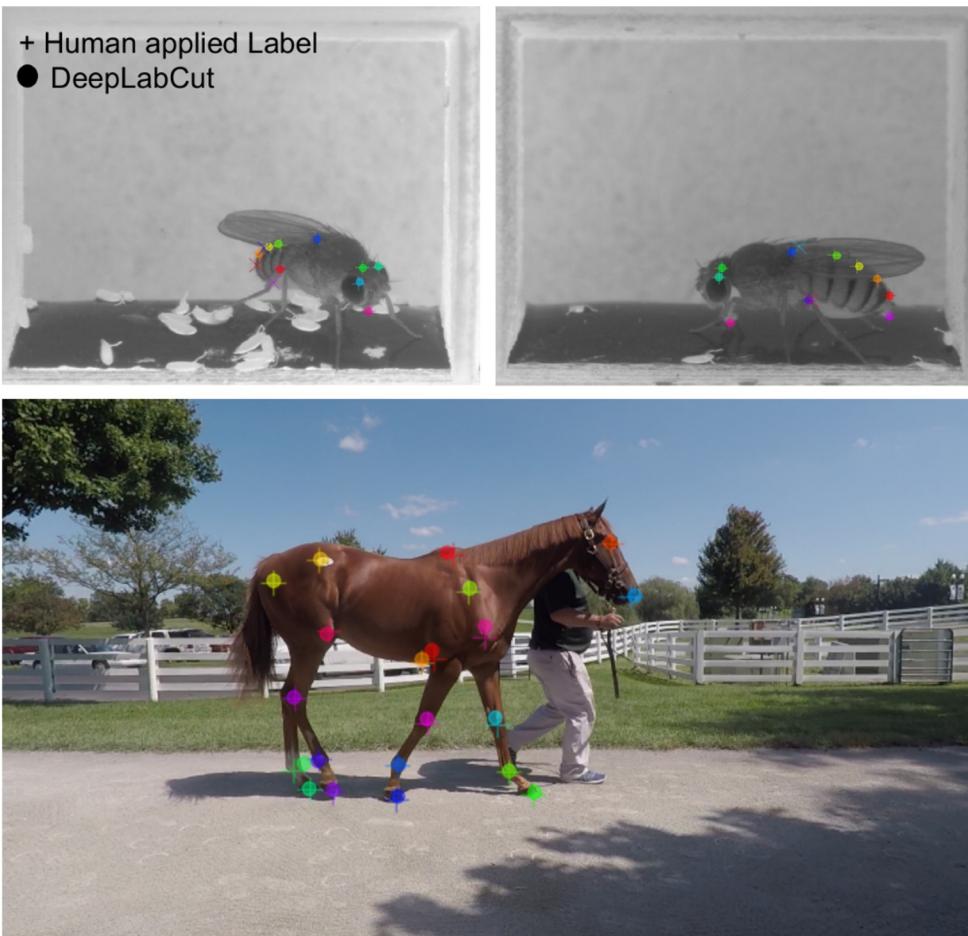
c

Examples of **test images** with Human and DeepLabCut labels



d

Examples of **test images** with Human and DeepLabCut labels



**Table 1 | Summary of commands**

|   |   |
|---|---|
| Evaluate the trained network (Step 11)        | <code>deeplabcut.evaluate_network(config_path)</code>   |
| Video analysis and plotting results (Step 11) | <code>deeplabcut.analyze_videos(config_path, ['path of video 1 or folder', 'path of video2', ...])</code> |
| Video analysis and plotting results (Step 12) | <code>deeplabcut.plot_trajectories(config_path, ['path of video 1', 'path of video2', ...])</code>        |
| Video analysis and plotting results (Step 13) | <code>deeplabcut.create_labeled_video(config_path, ['path of video 1', 'path of video2', ...])</code>     |
| Refinement: extract outlier frames (Step 14)  | <code>deeplabcut.extract_outlier_frames(config_path, ['path of video 1', 'path of video 2'])</code>       |
| Refine labels (Step 15)                       | <code>deeplabcut.refine_labels(config_path)</code>  |
| Combine datasets (Step 16)                    | <code>deeplabcut.merge_datasets(config_path)</code>   |

# Video analysis

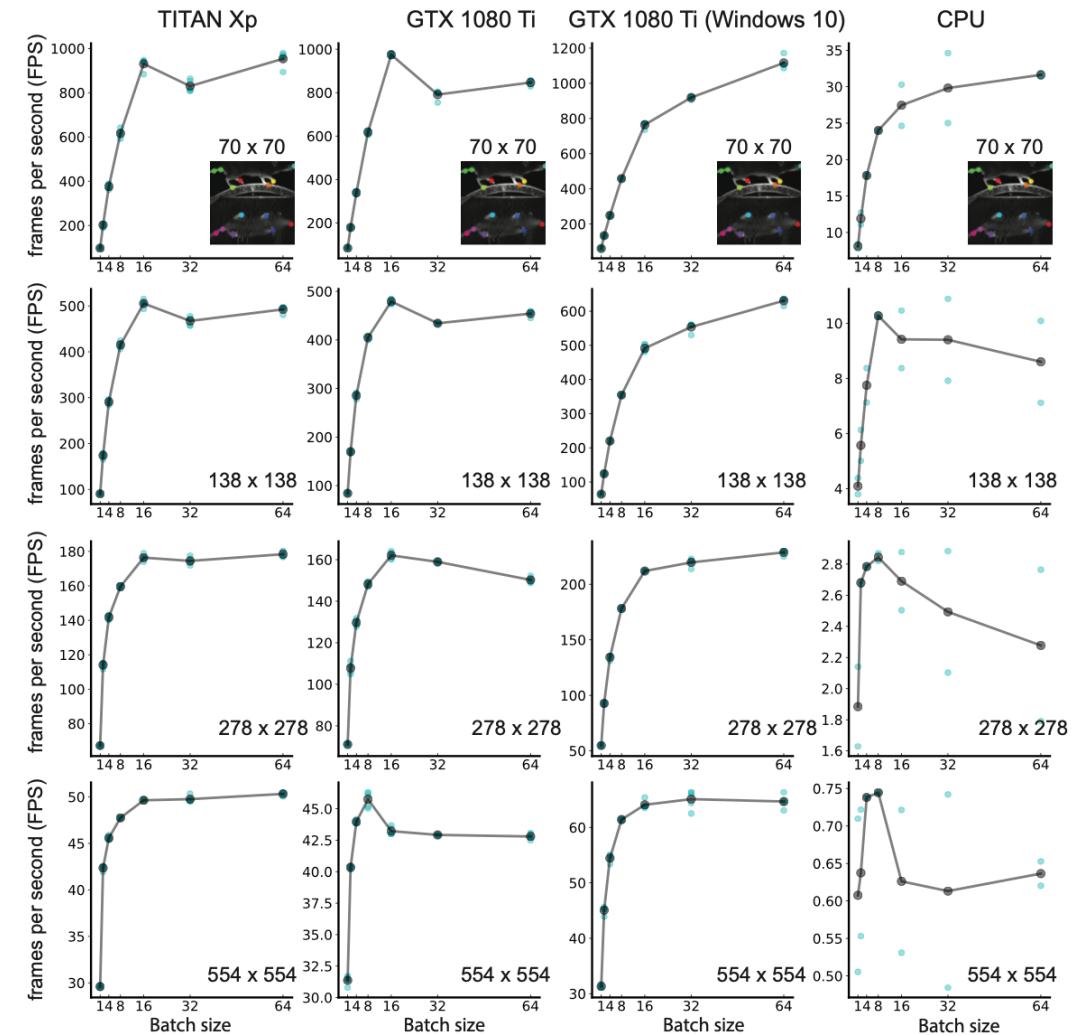
Videos for doing inference

need **not** be added to the config.yaml (only for labeling/active learning)!

Set *batch\_size* in the config.yaml!

Also check out the real-time DLC paper for speed benchmarking:

<https://elifesciences.org/articles/61909>

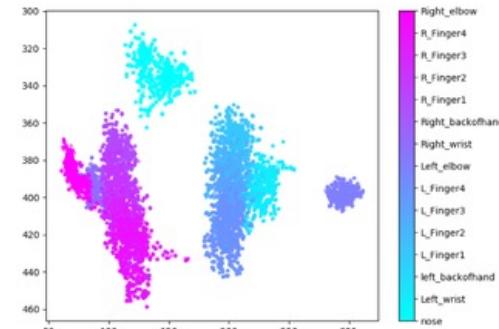


# Trajectory visualization

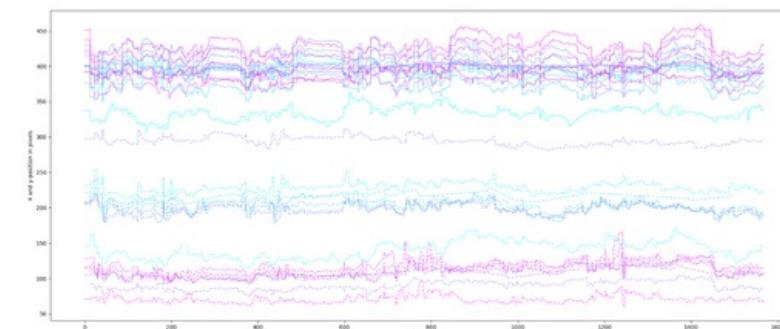
```
deeplabcut.plot_trajectories(config_path,['fullpath/analysis/project/videos/reachingvideo1.avi'])
```



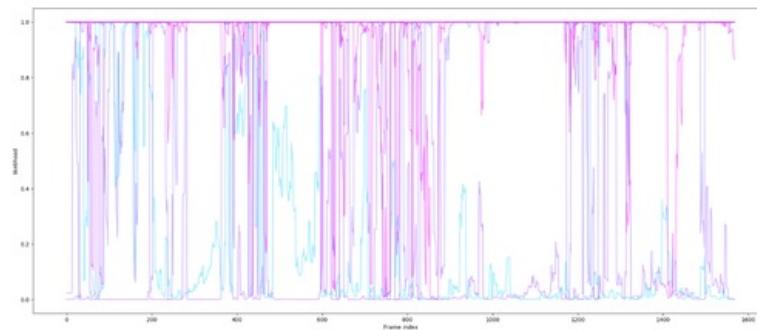
Body parts plotted in space  
(over all the frames)



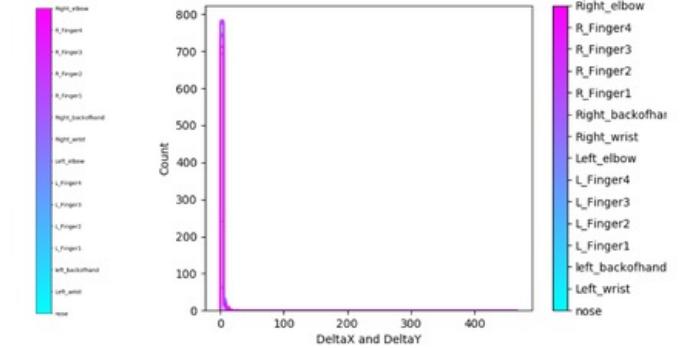
All body parts across time (frames).  
solid lines are Y and dashed lines are X



Every body part likelihood over time  
(over all the frames)



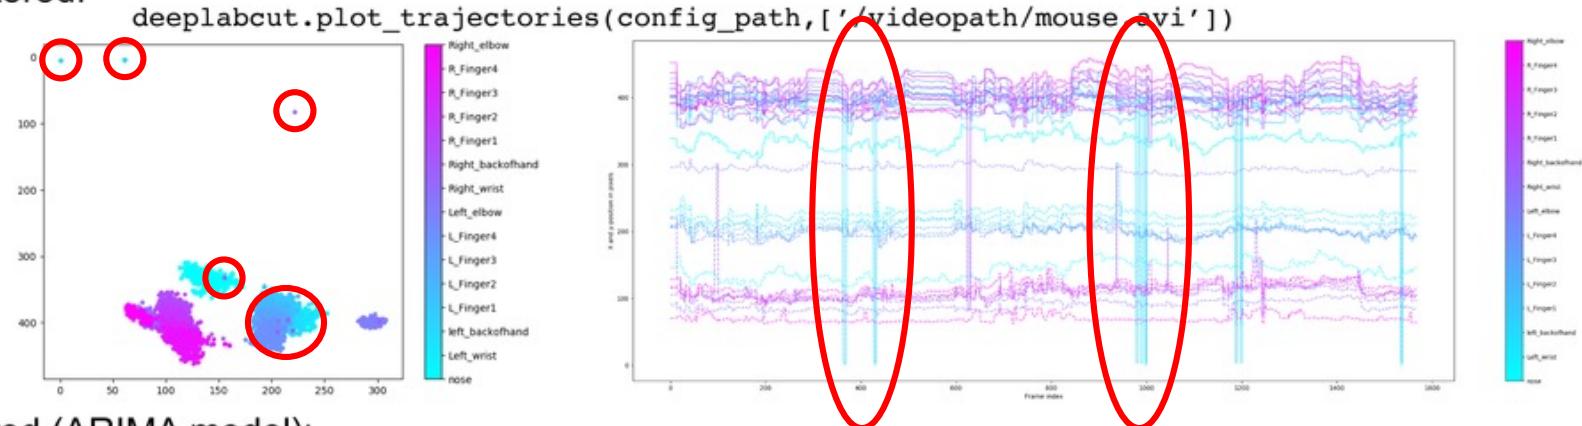
Consecutive coordinate differences  
(low values = minimal jumps across frames)



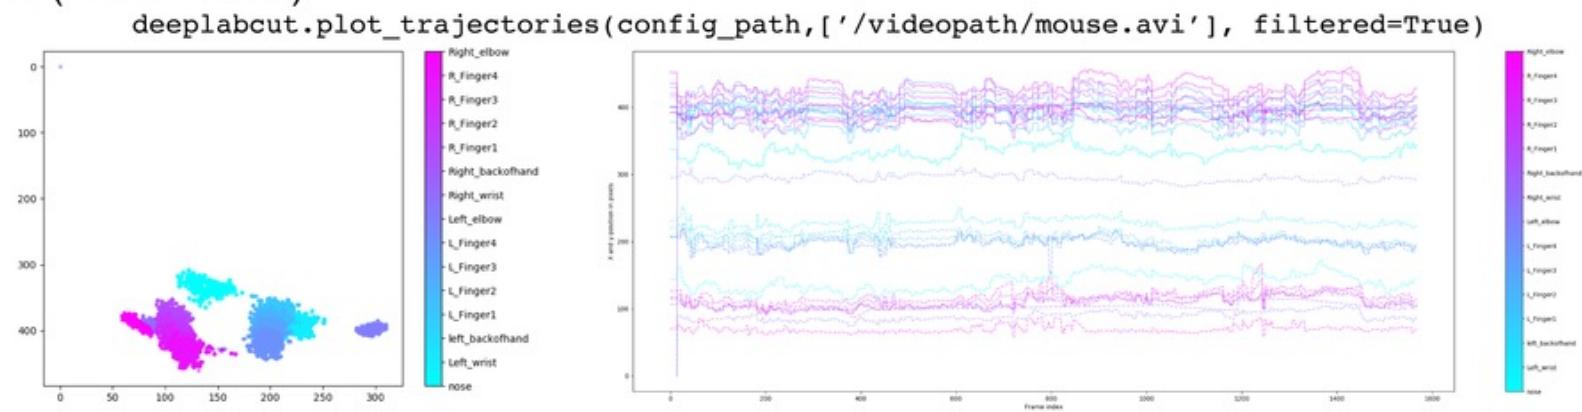
# Trajectory filtering

```
deeplabcut.filterpredictions(config_path,['/videopath/mouse.avi'], p_bound=0.1, ARdegree=5, MAdegree=1)
```

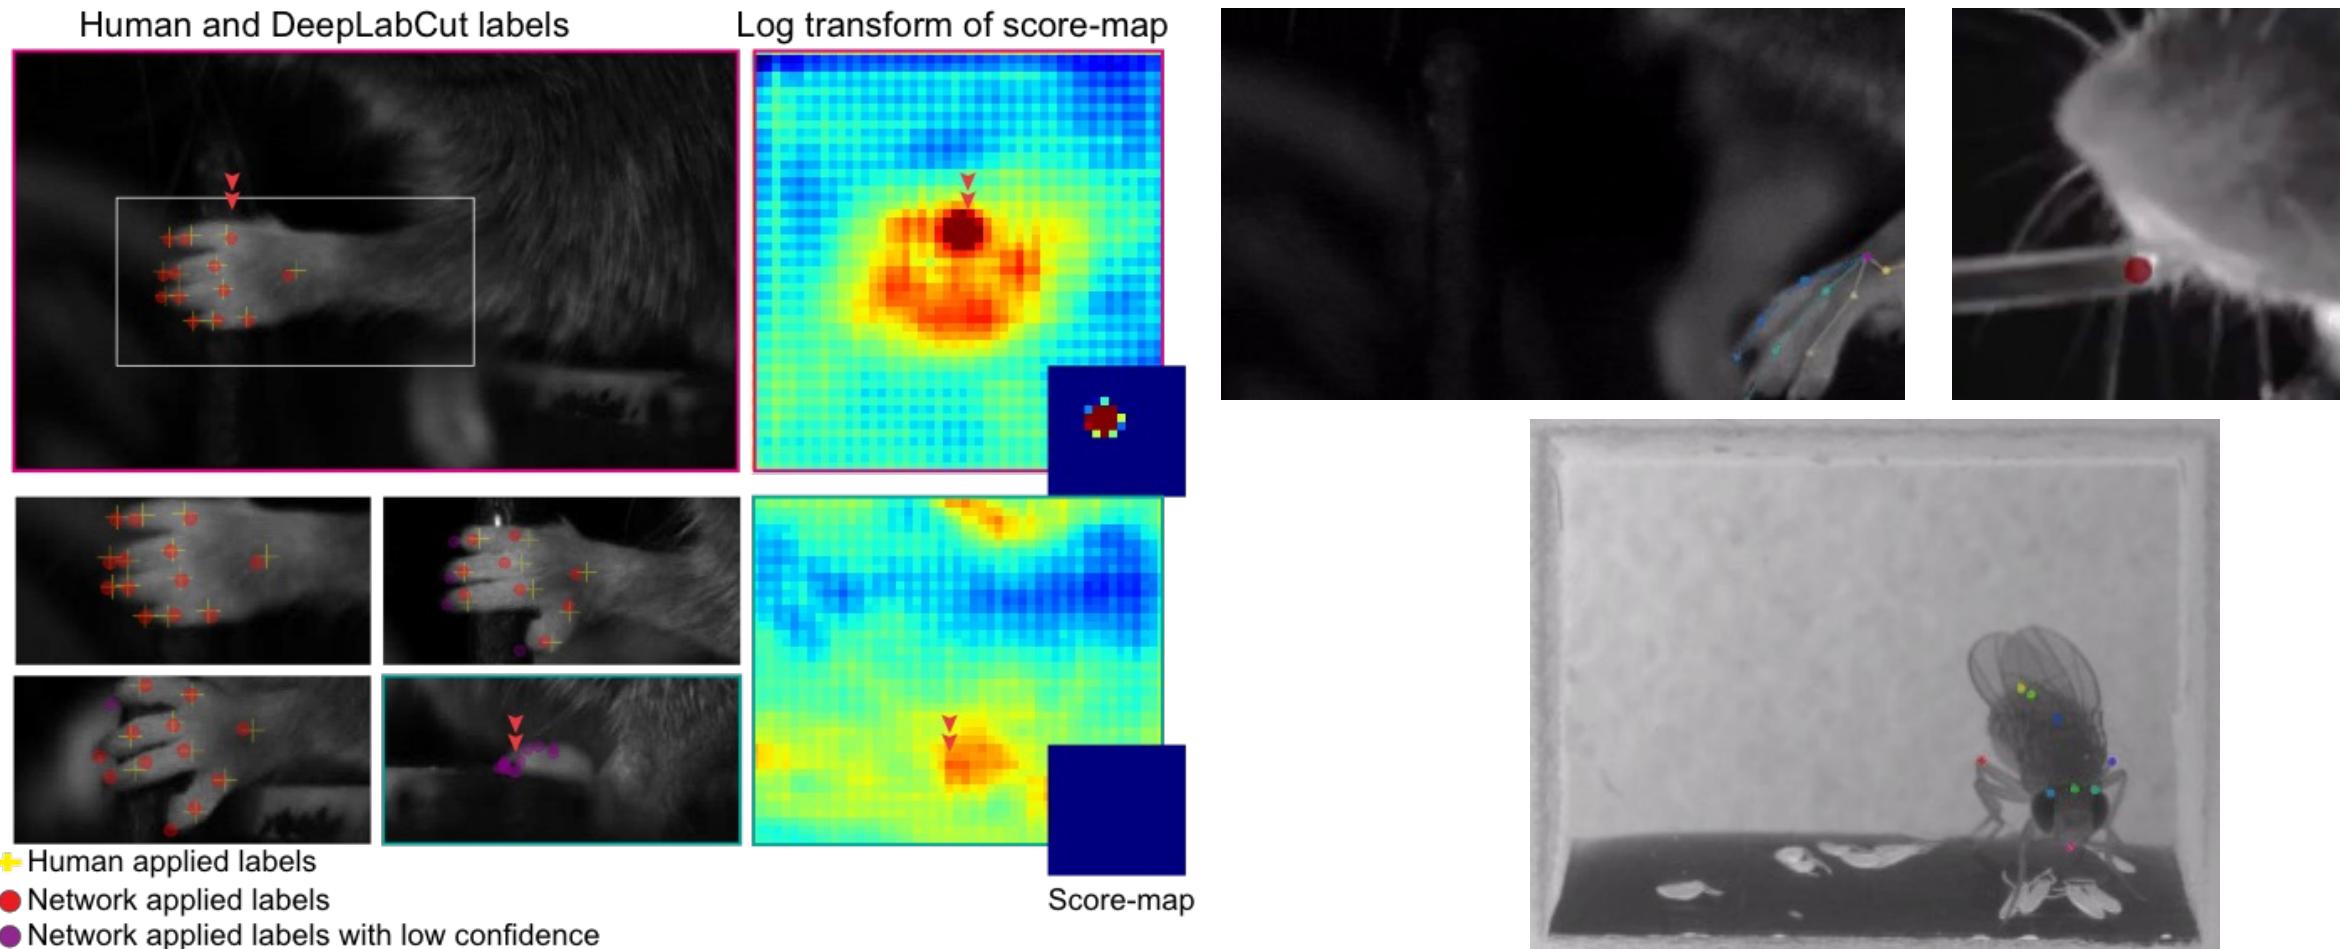
Unfiltered:



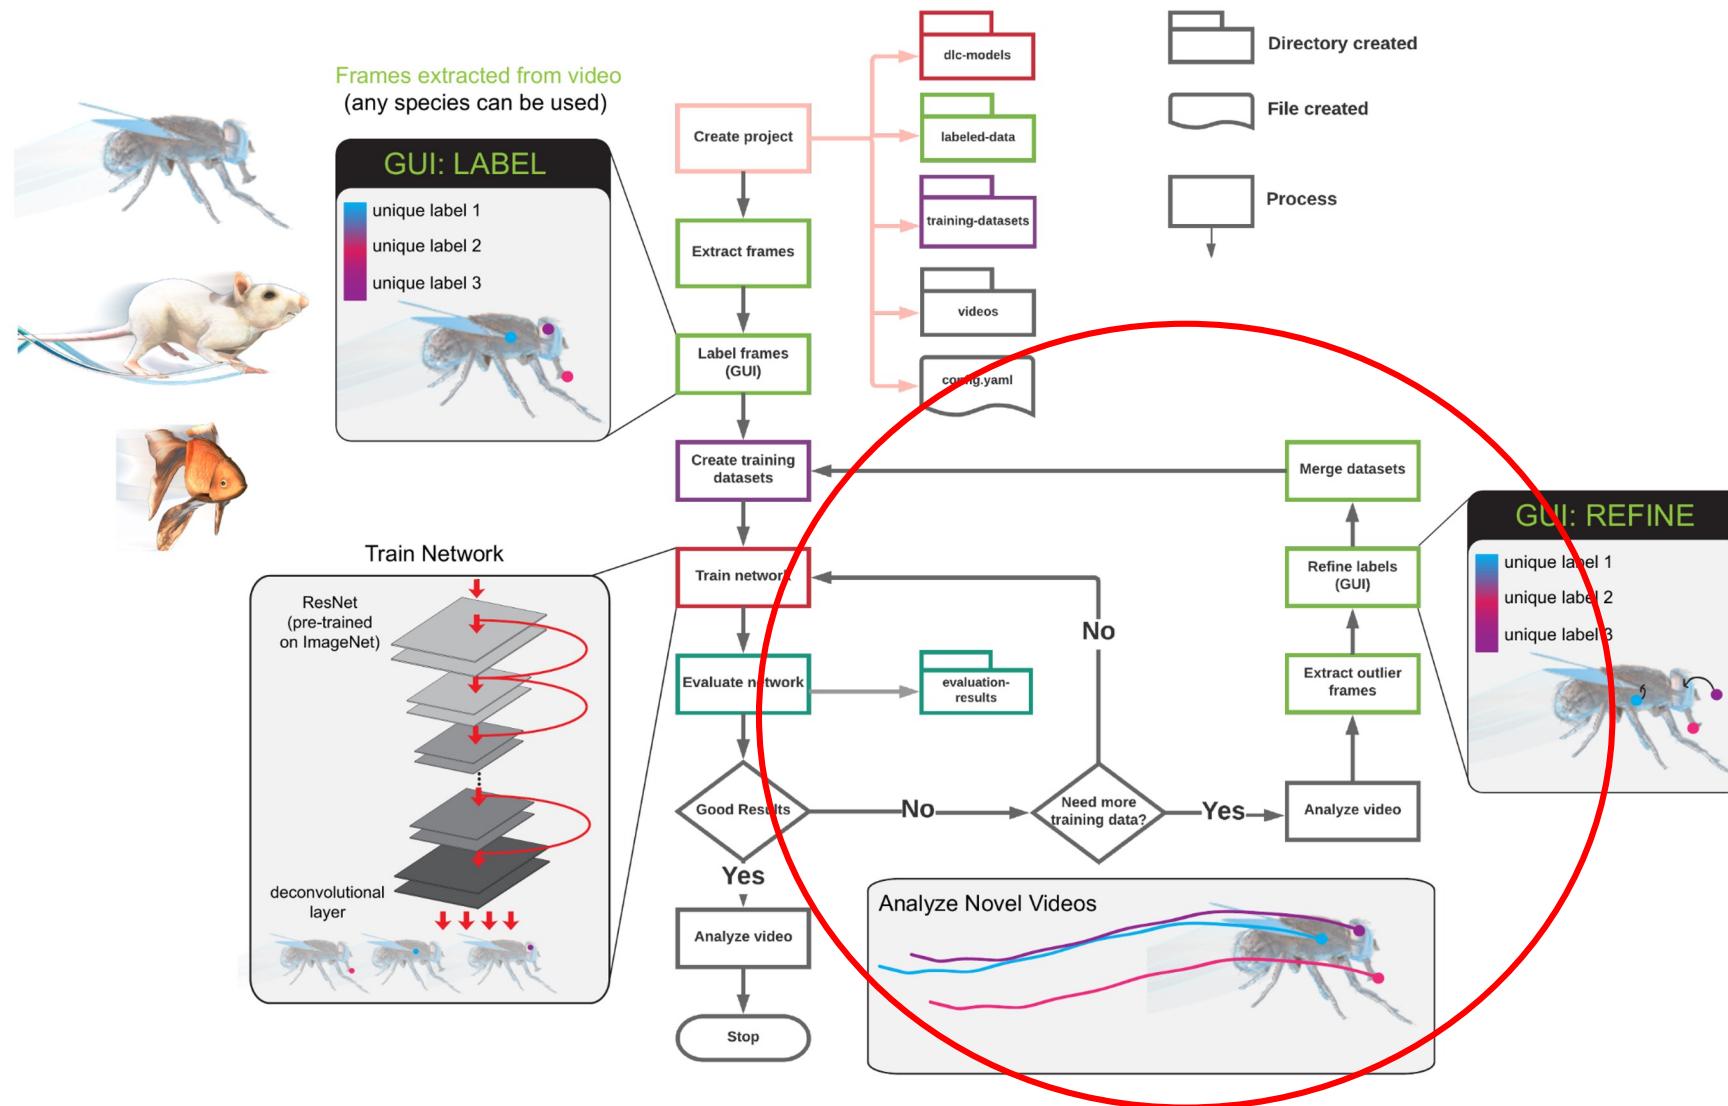
Filtered (ARIMA model):



# Score-maps provide network confidence readout



# DeepLabCut 2.0 workflow



# Active learning

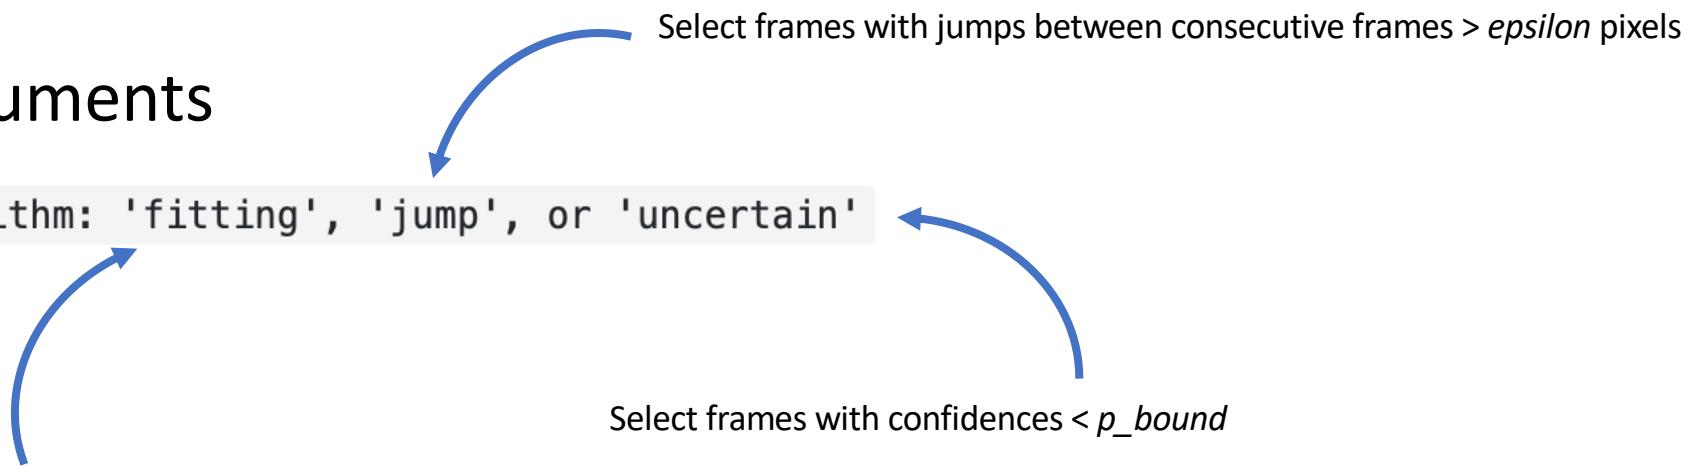
- If there are errors in analyzed videos, you can correct those manually and re-train!
- Note you do not need to correct all errors (presumably some are fixed due to generalization)

# Active learning

```
deeplabcut.extract_outlier_frames(config_path, ['videofile_path'])
```

- Key arguments

`outlieralgorithm: 'fitting', 'jump', or 'uncertain'`



Select frames deviating from  
a statistical model fit to the data

Select frames with confidences <  $p\_bound$

Select frames with jumps between consecutive frames >  $\epsilon_{pixel}$

Random sampling

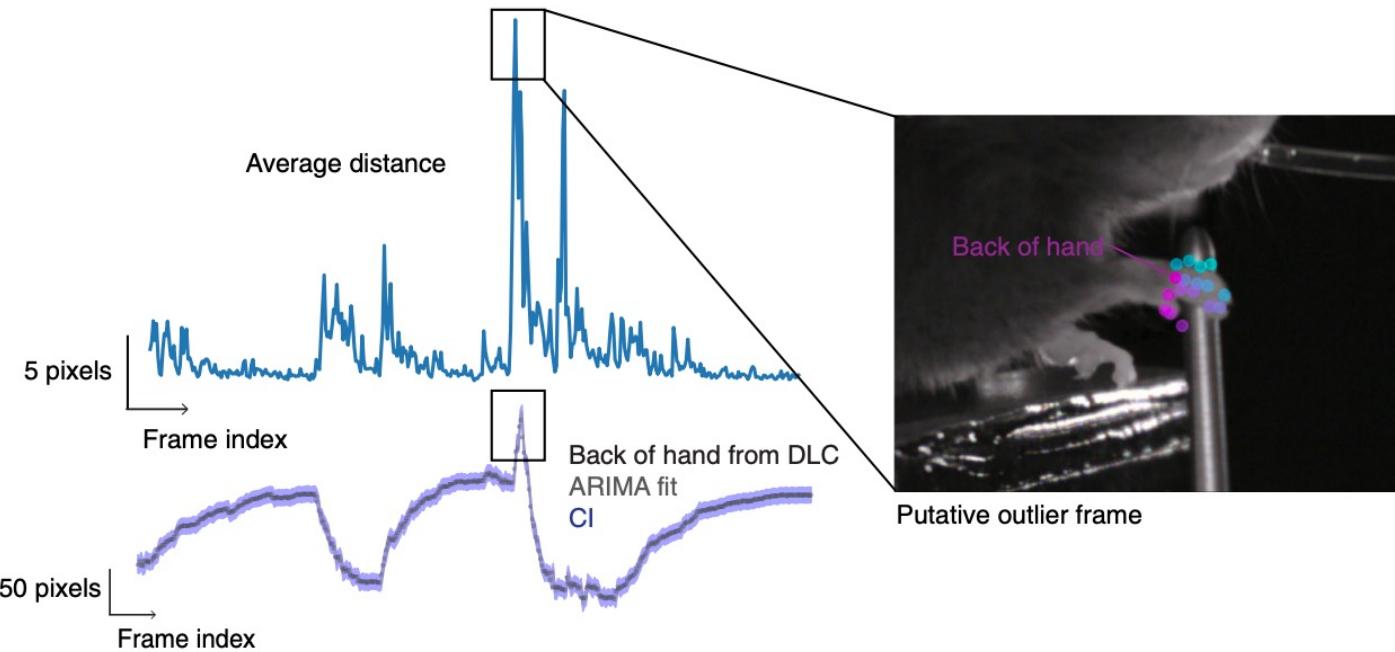
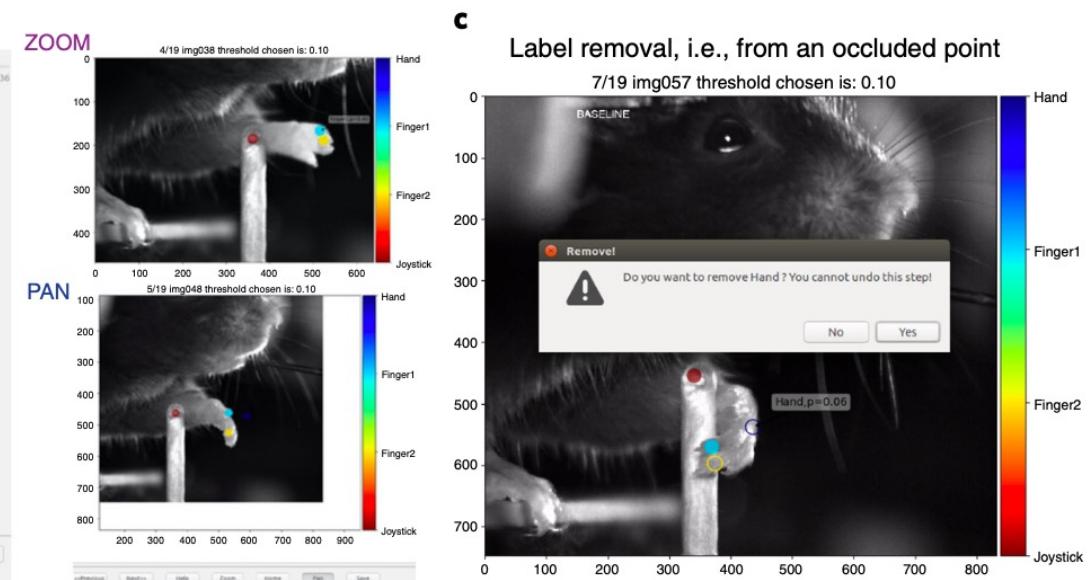
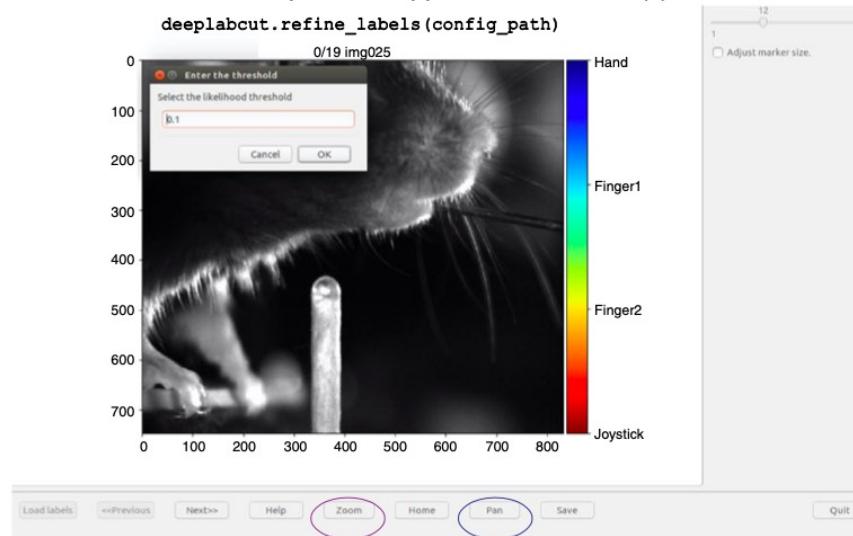
Clustering

`extractionalgorithm='uniform'`

`extractionalgorithm='k-means'`

**a**

## Identification of outlier frames

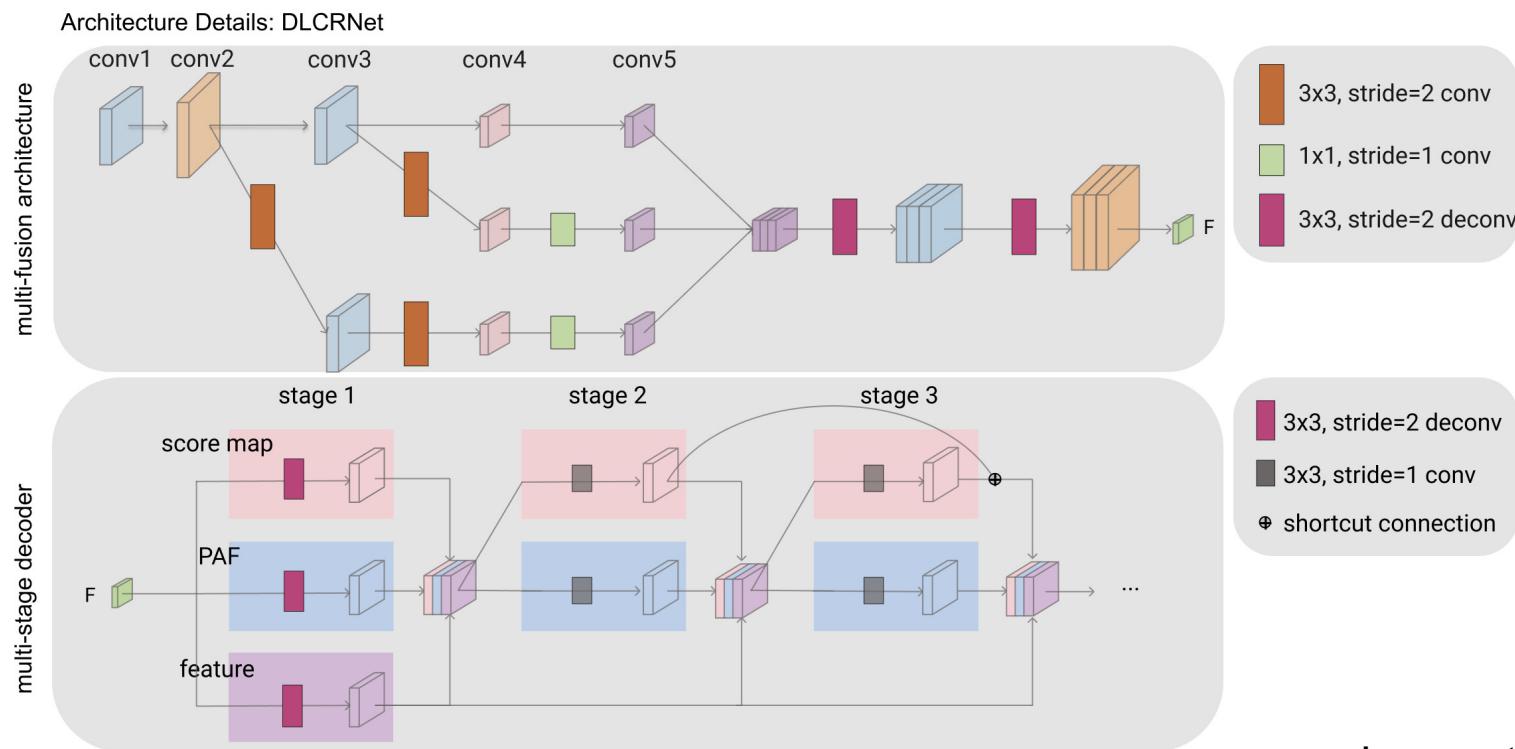
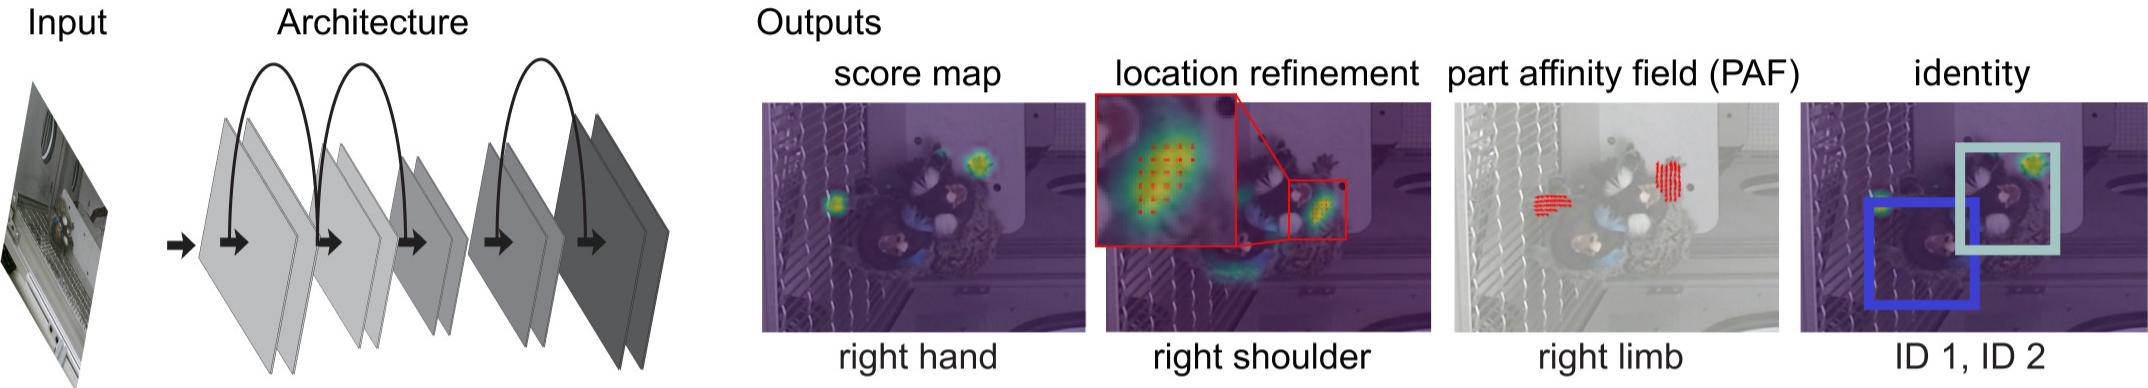
**b** Refinement of DeepLabCut-applied label location(s)**c** Label removal, i.e., from an occluded point

# Multi-animal DLC

DEMO notebook for 3 mice:

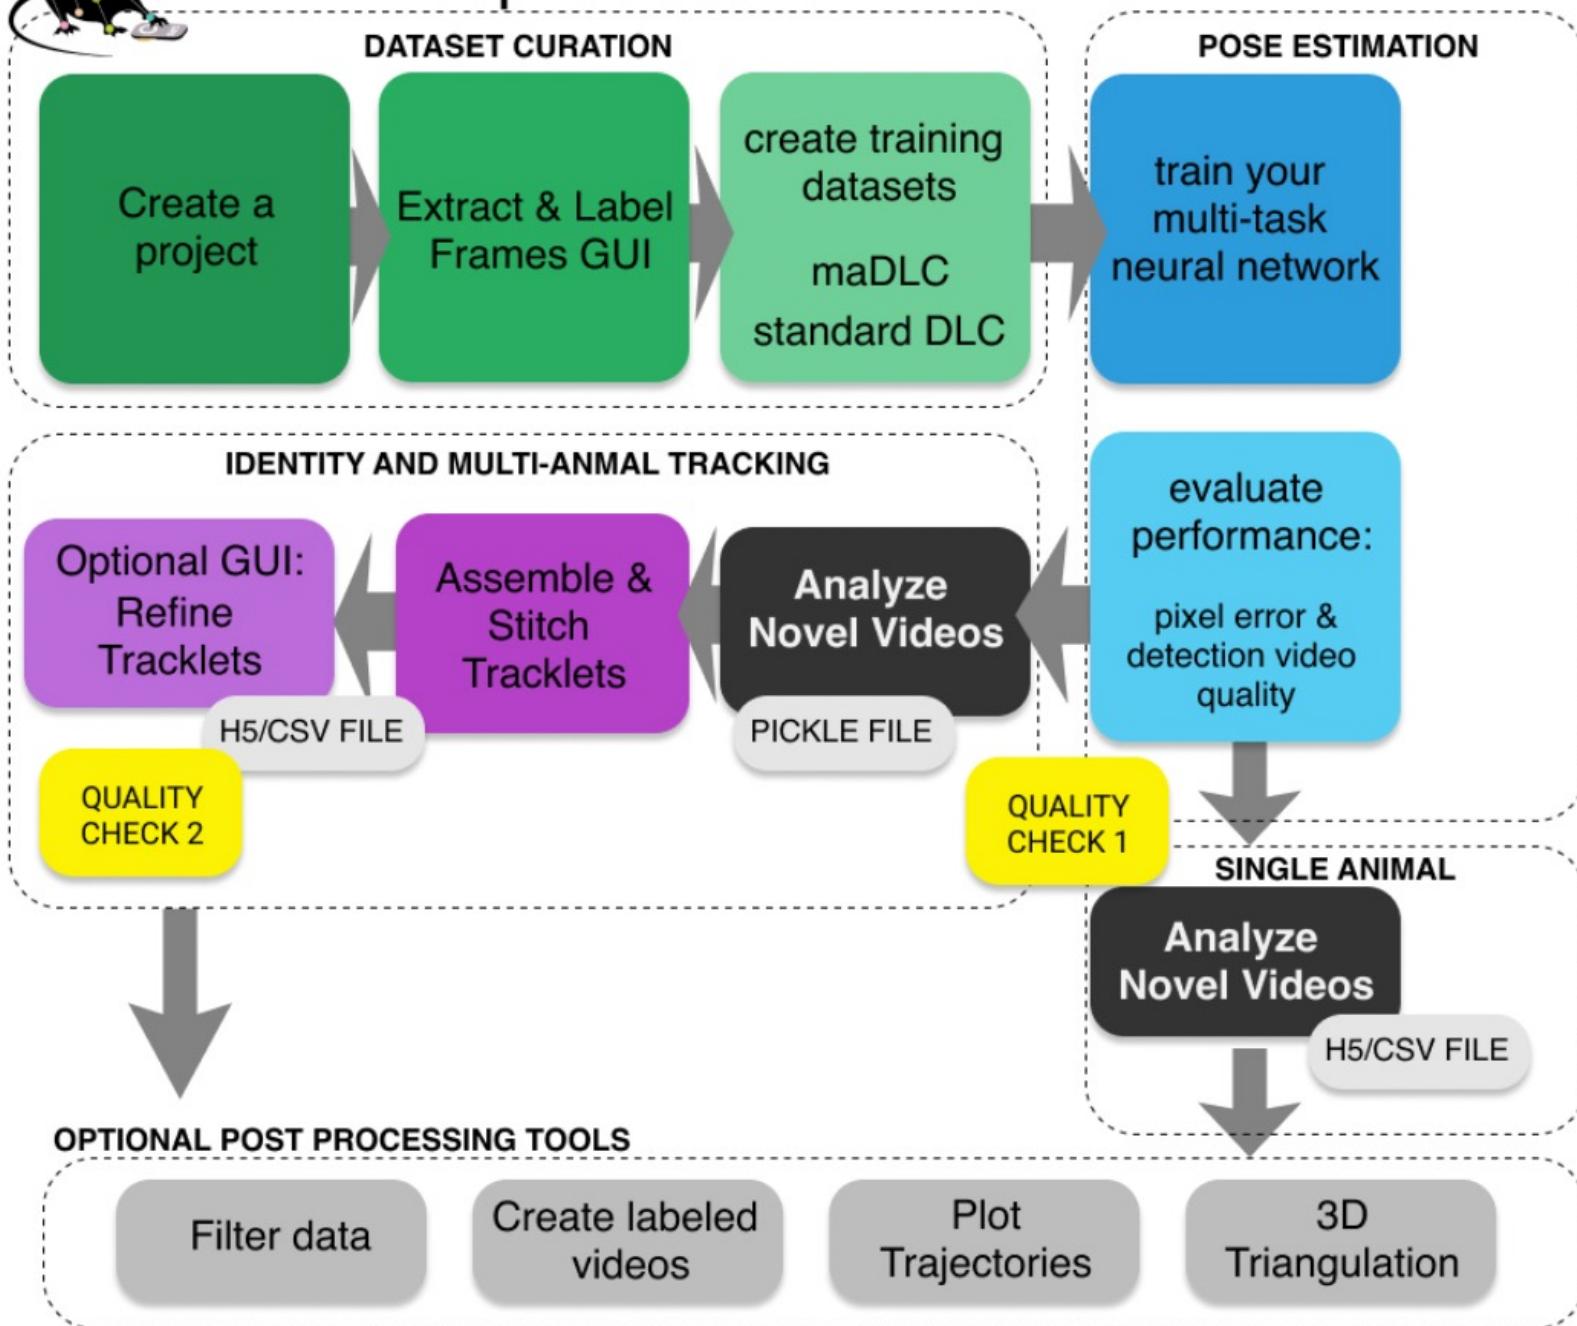
<https://github.com/DeepLabCut/DeepLabCut/blob/master/examples/COLAB>

# Multi-animal pose estimation & tracking with DeepLabCut



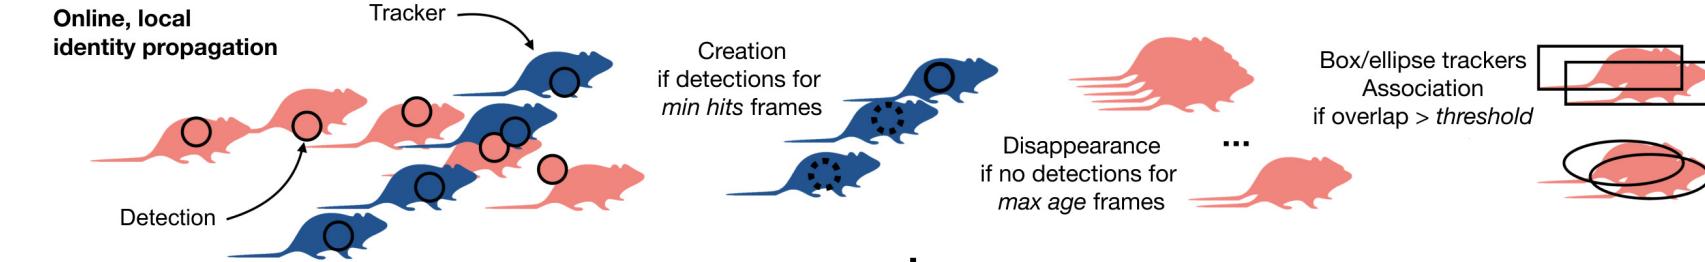


# DeepLabCut 2.2+ workflow

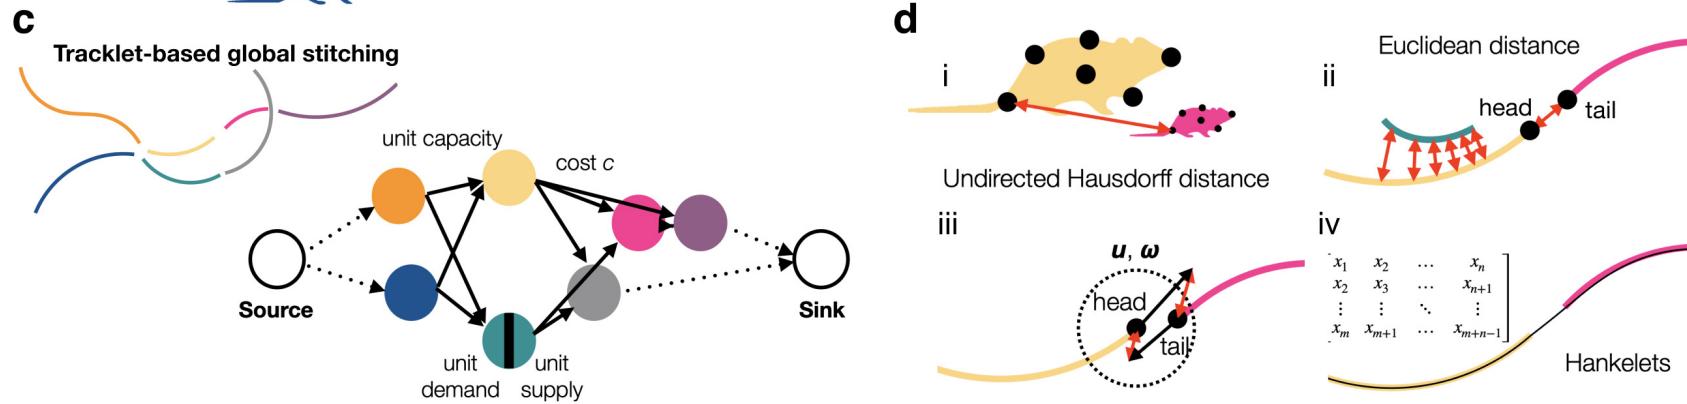


The workflow is available in the GUI or via Python.

# Tracking



# Stitching



Corresponding commands:

```
dlc.analyze_videos( ... )
```

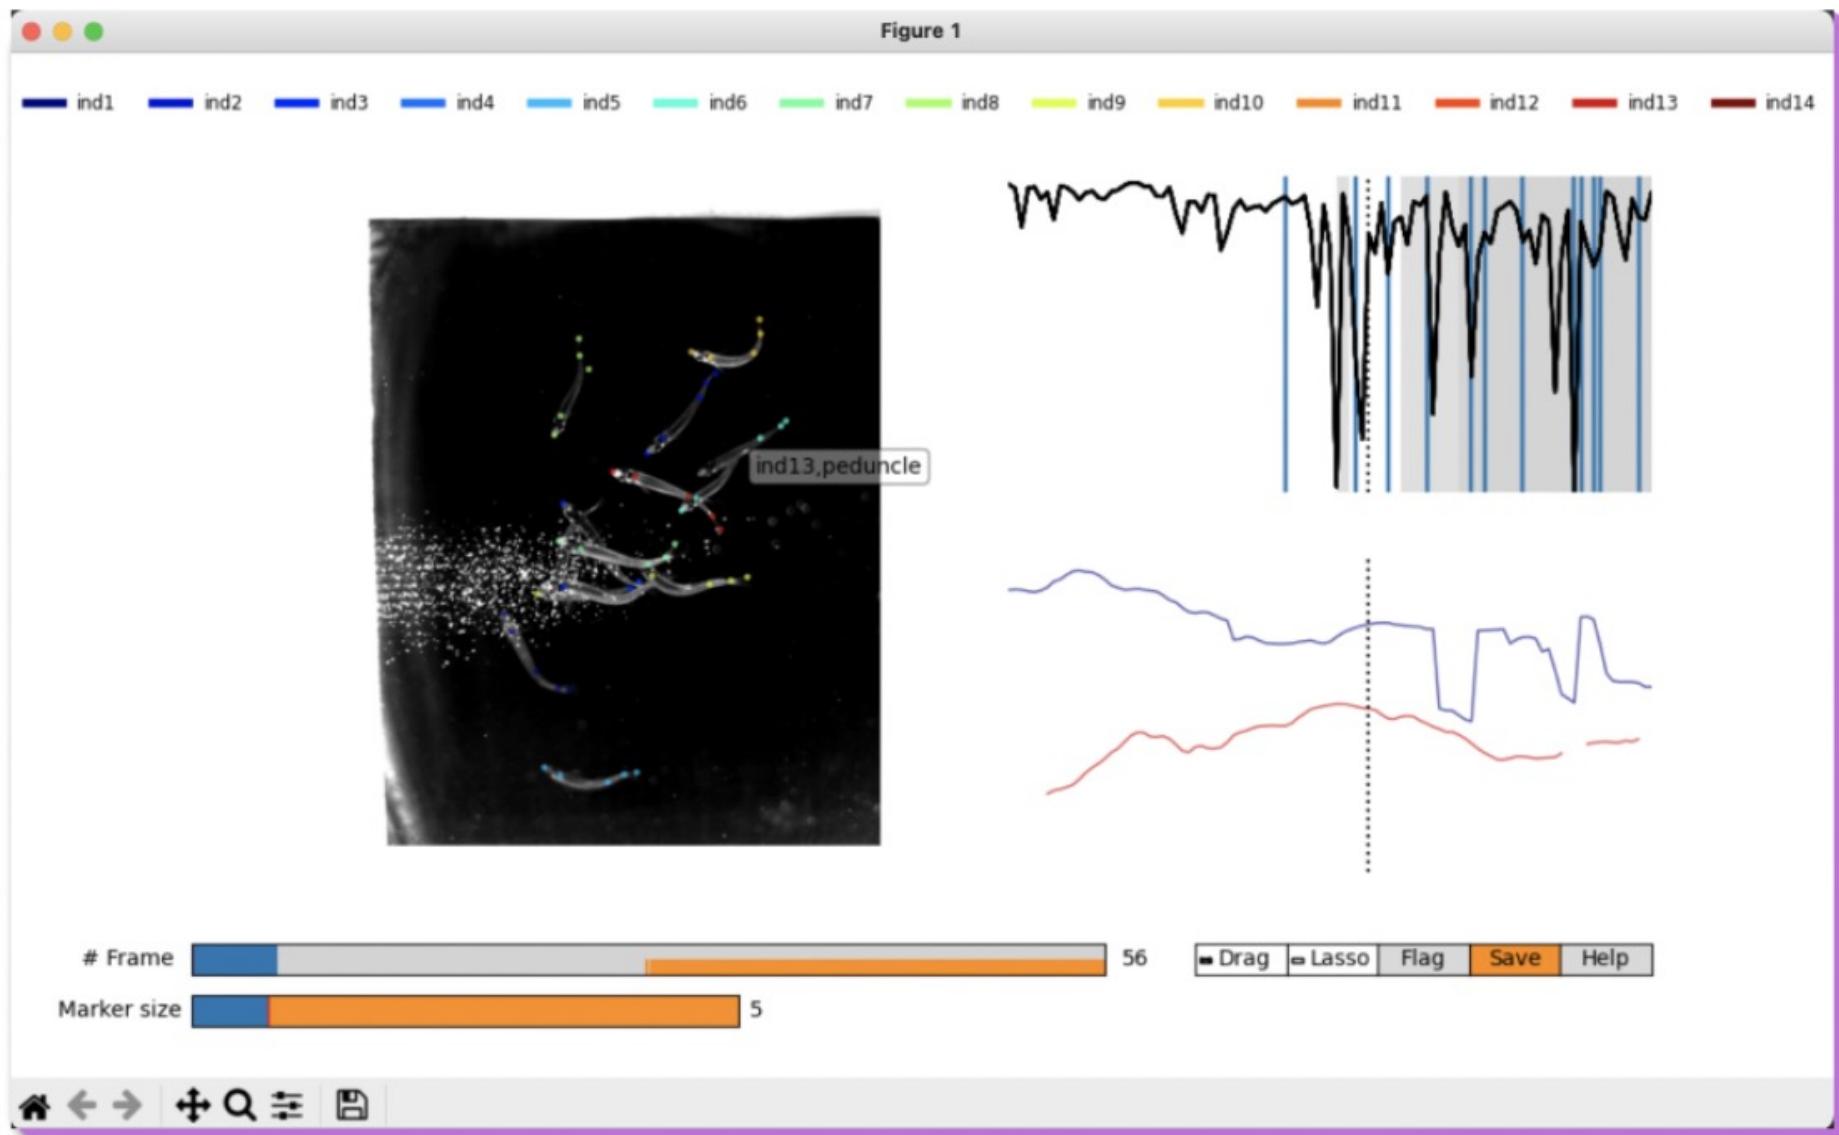
```
dlc.convert_detections2tracklets(...)
```

```
dlc.stitch_tracklets(...)
```

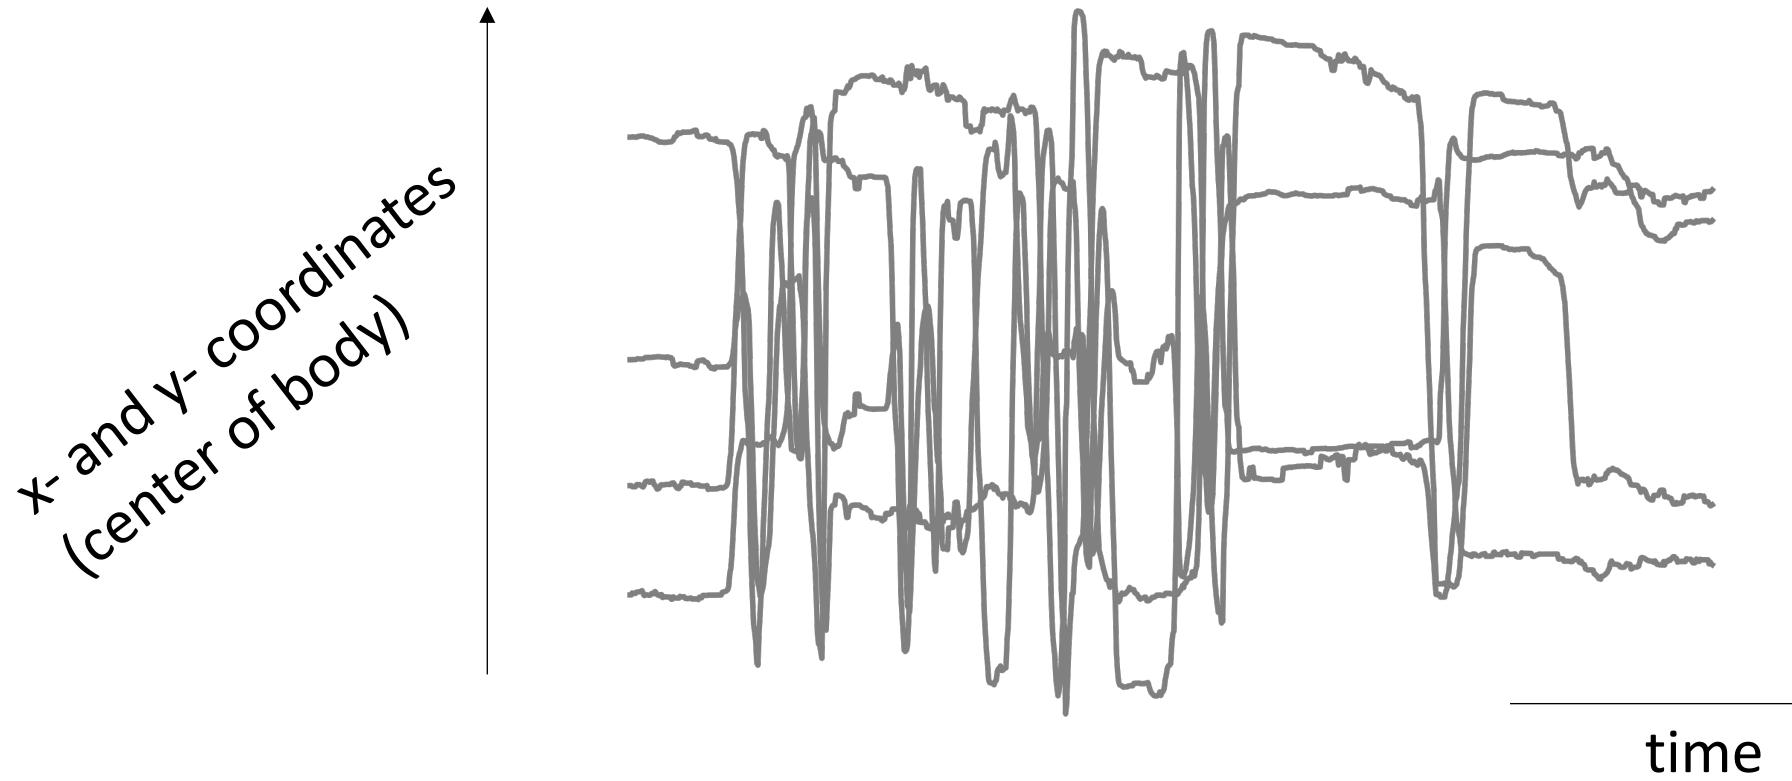
**DEMO:**

[https://github.com/DeepLabCut/DeepLabCut/blob/master/examples/COLAB/COLAB\\_3](https://github.com/DeepLabCut/DeepLabCut/blob/master/examples/COLAB/COLAB_3)

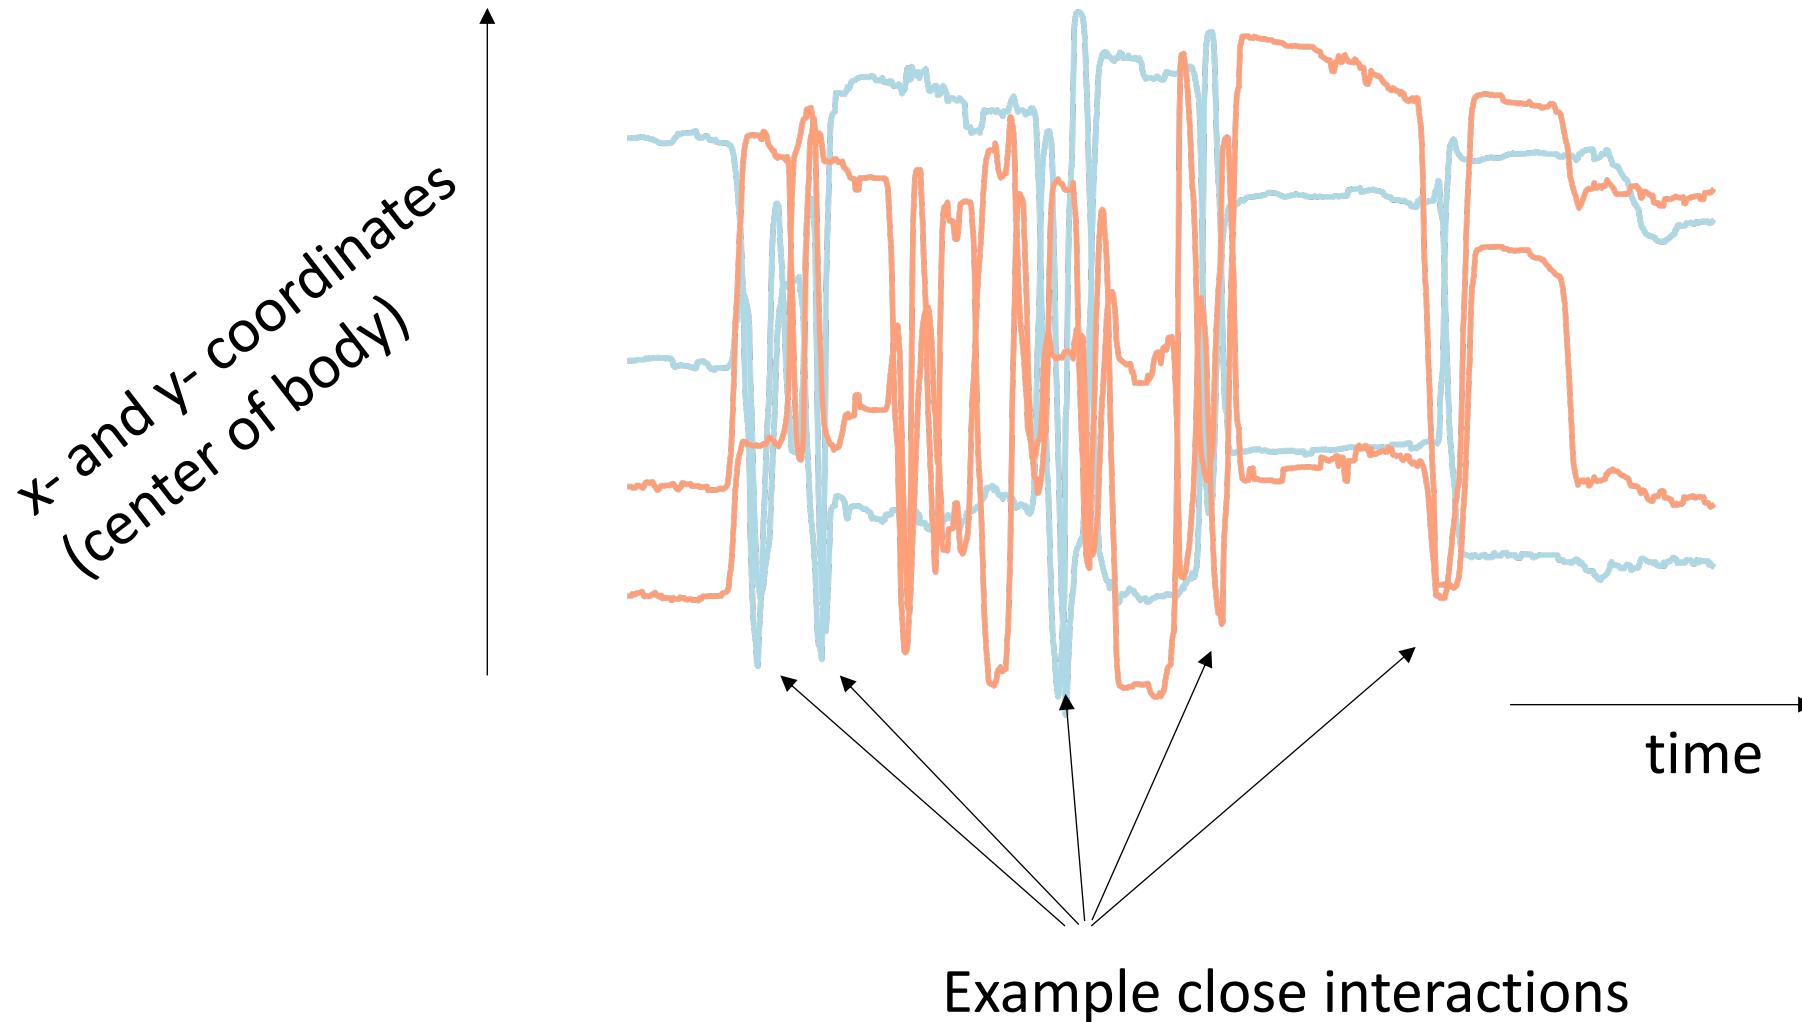
## Refine tracklets GUI: deeplabcut.refine\_tracklets()



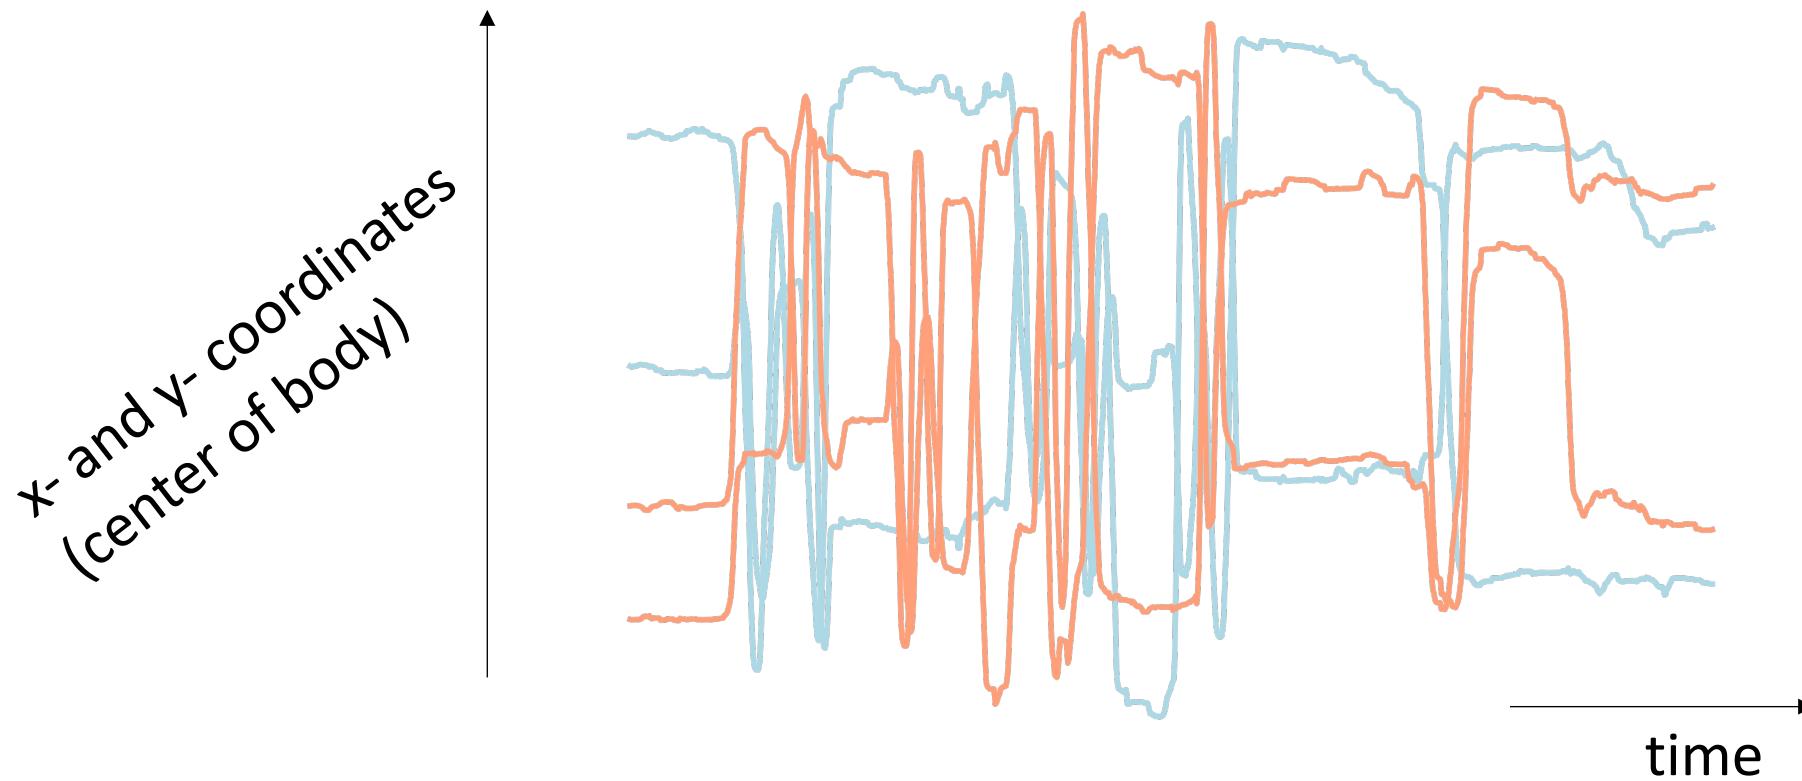
# Re-Identification is crucial for tracking in long-term scenarios



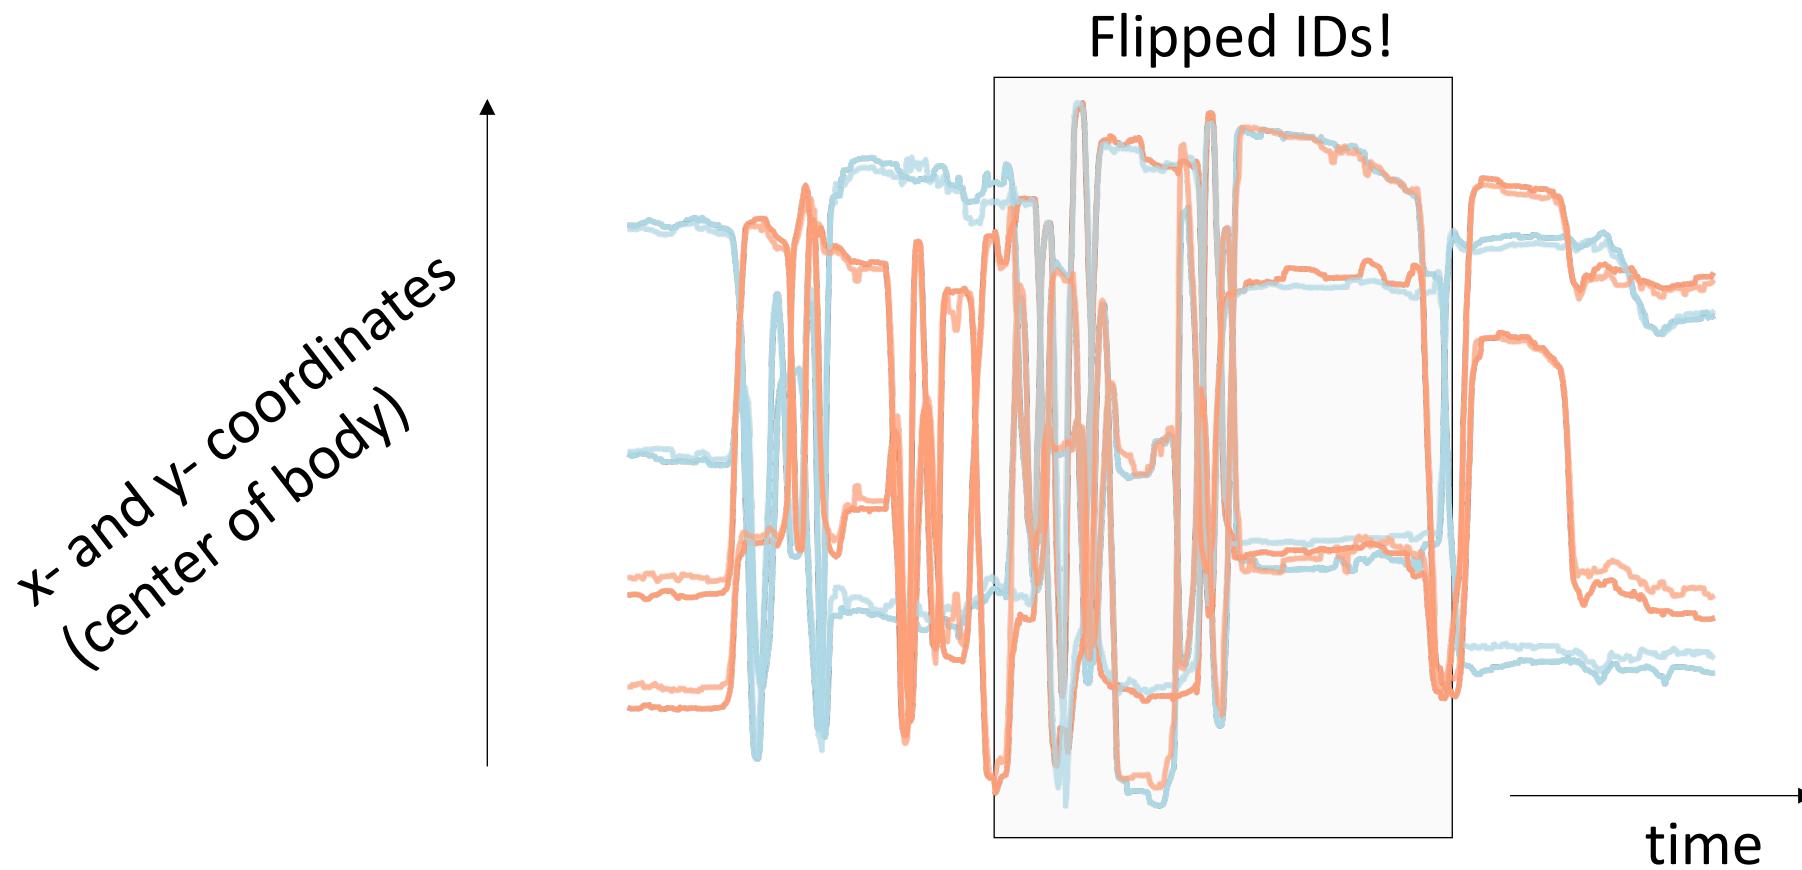
# Ground truth data



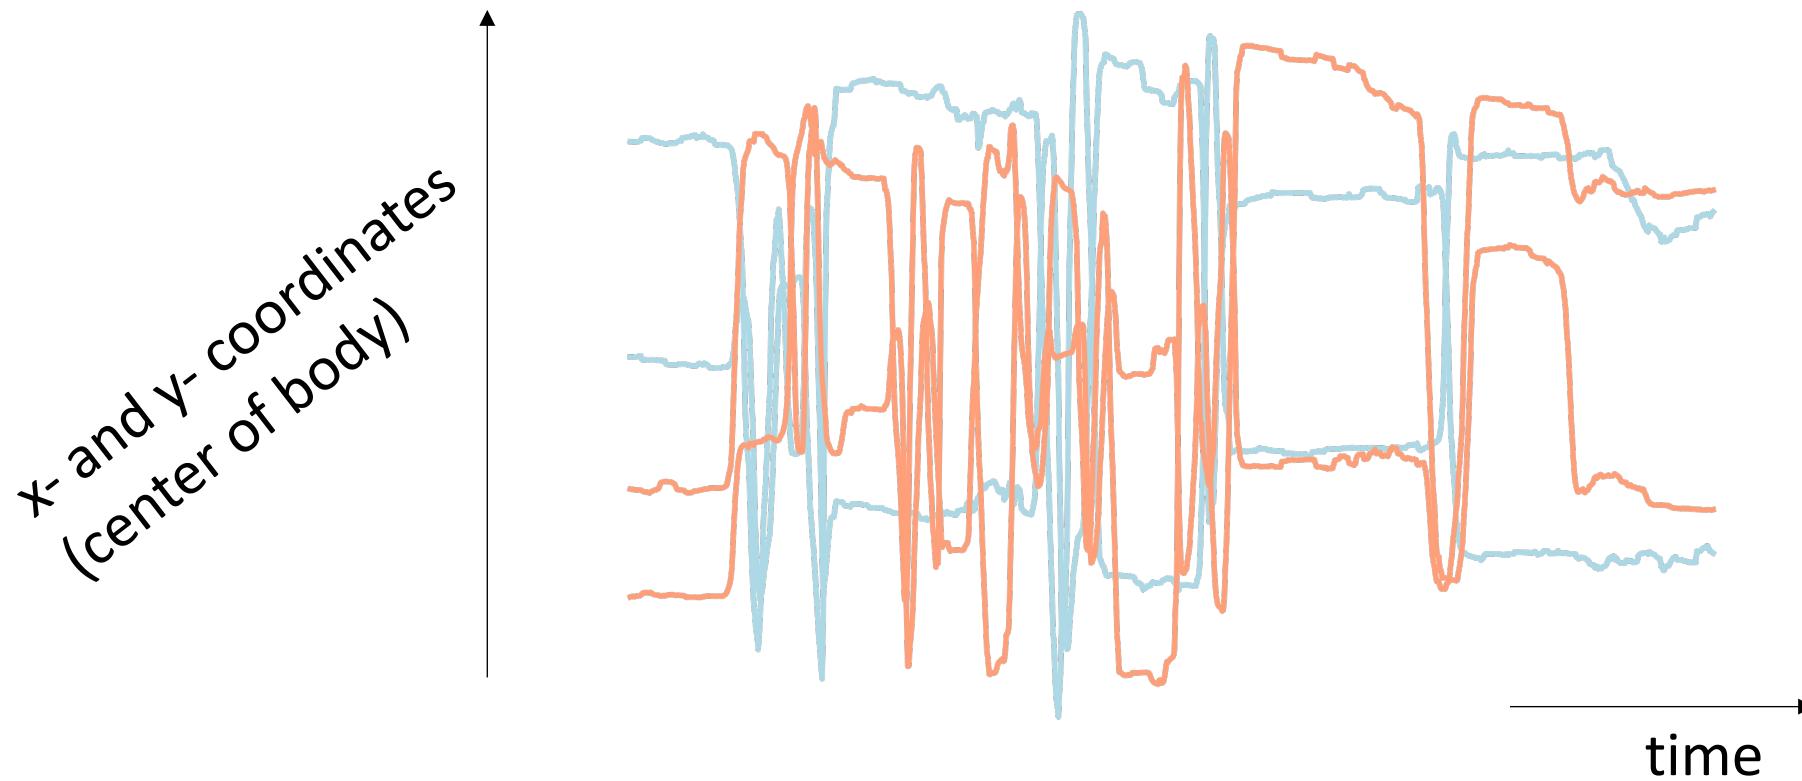
# ID-based (temporally) \*local\* tracking



# ID-based (temporally) \*local\* tracking

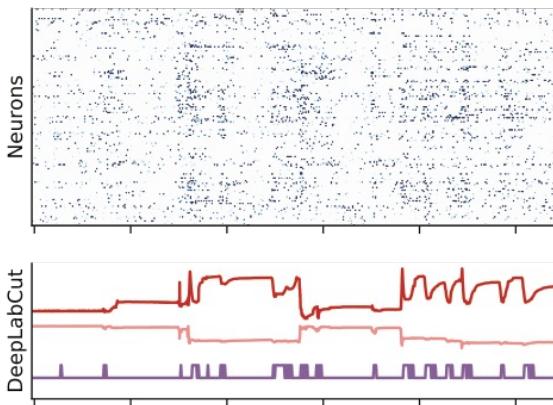


# ReID based tracking (frame wise!)

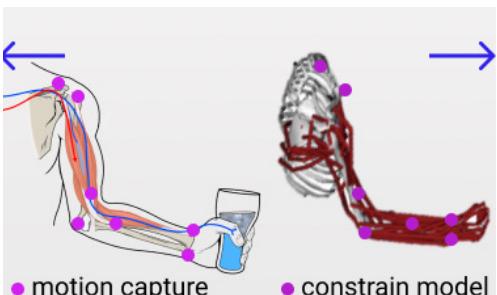


# From pose to actions....

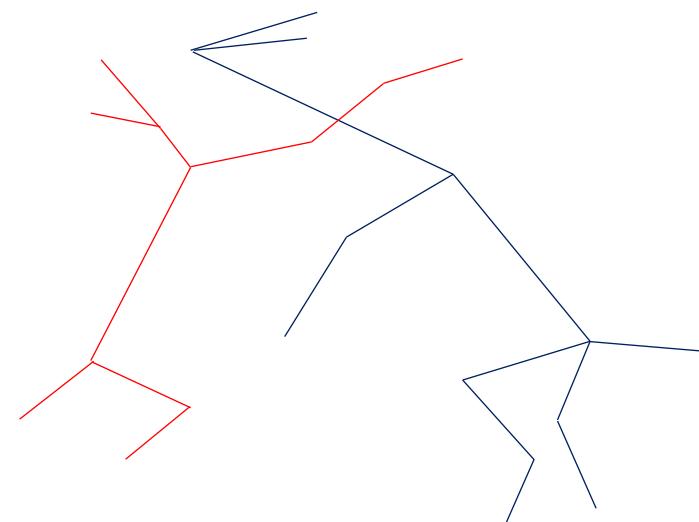
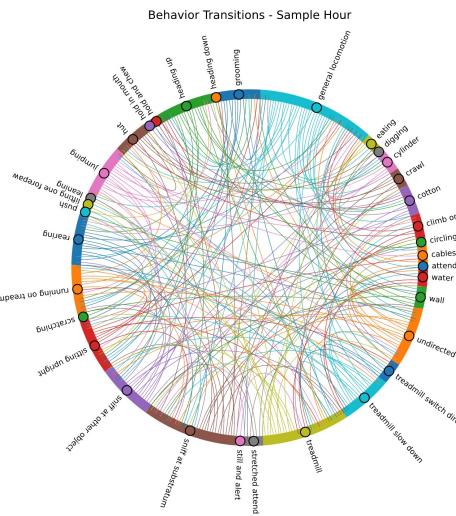
kinematics



biomechanics  
/ modeling



actions



Primer

## A Primer on Motion Capture with Deep Learning: Principles, Pitfalls, and Perspectives

Alexander Mathis,<sup>1,2,3,\*</sup> Steffen Schneider,<sup>3,4</sup> Jessy Lauer,<sup>1,2,3</sup> and Mackenzie Weygandt Mathis<sup>1,2,3,\*</sup>

<sup>1</sup>Center for Neuroprosthetics, Center for Intelligent Systems, Swiss Federal Institute of Technology (EPFL), Lausanne, Switzerland

<sup>2</sup>Brain Mind Institute, School of Life Sciences, Swiss Federal Institute of Technology (EPFL), Lausanne, Switzerland

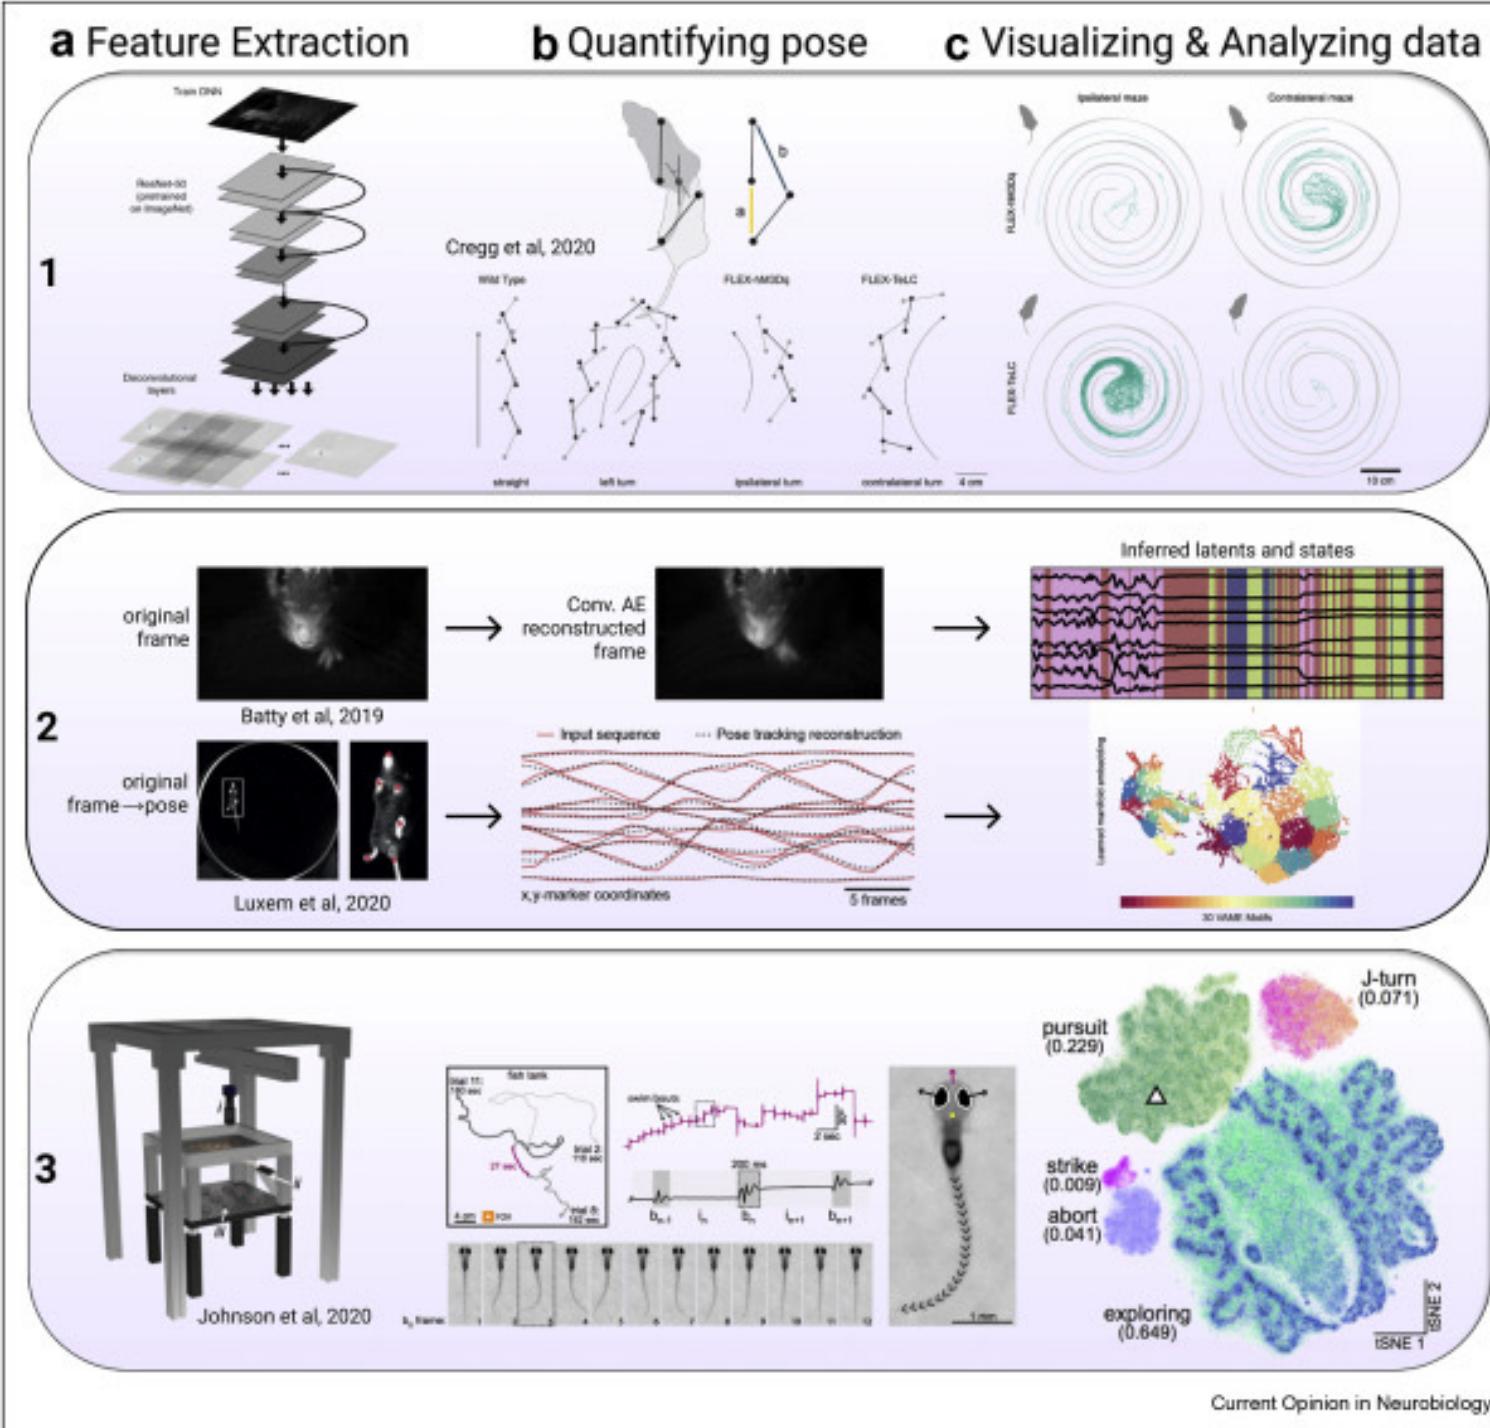
<sup>3</sup>The Rowland Institute at Harvard, Harvard University, Cambridge, MA, USA

<sup>4</sup>University of Tübingen and International Max Planck Research School for Intelligent Systems, Tübingen, Germany

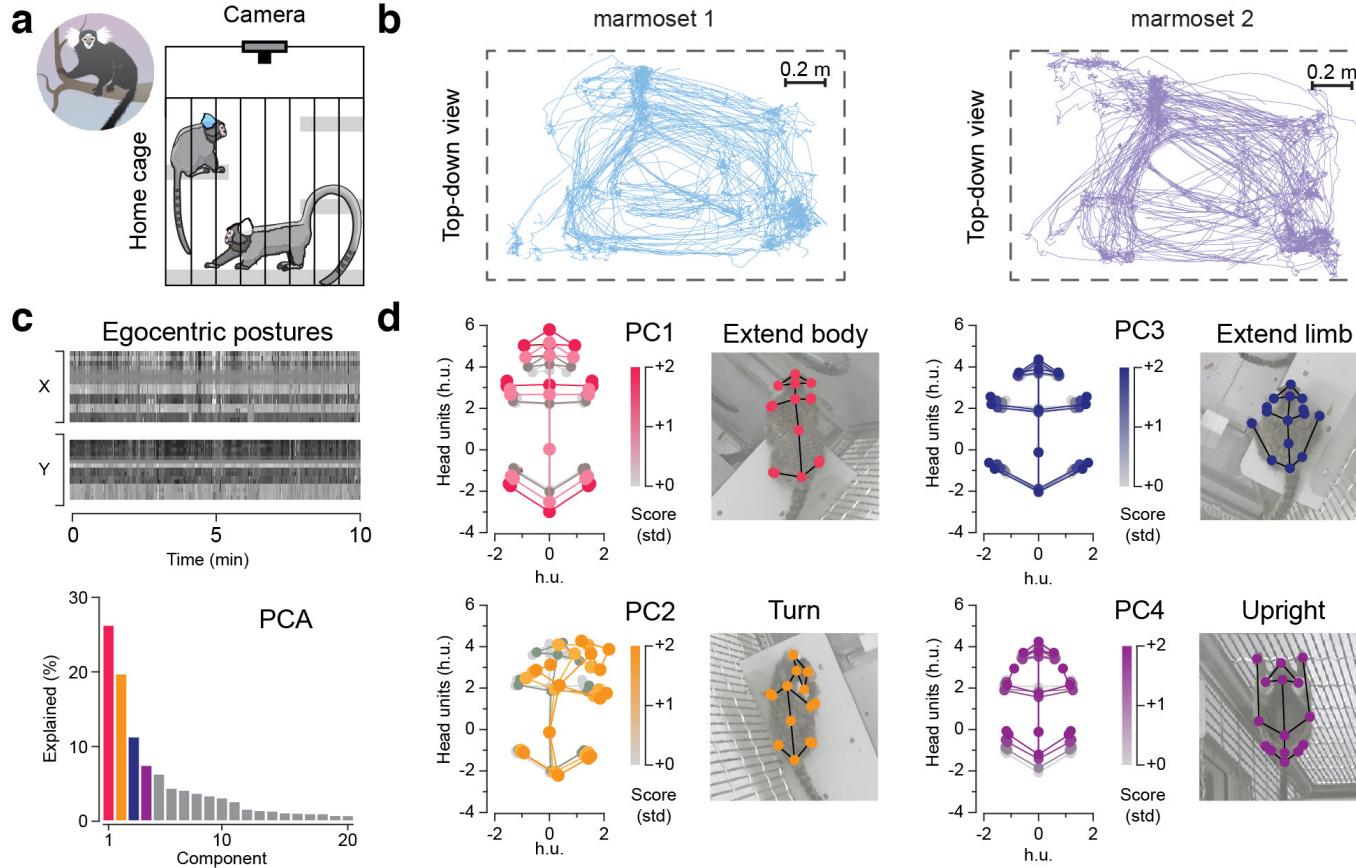
\*Correspondence: alexander.mathis@epfl.ch (A.M.), mackenzie.mathis@epfl.ch (M.W.M.)

<https://doi.org/10.1016/j.neuron.2020.09.017>

# What to do with the pose estimation outputs?



# Unsupervised Behavioral analysis based on DLC

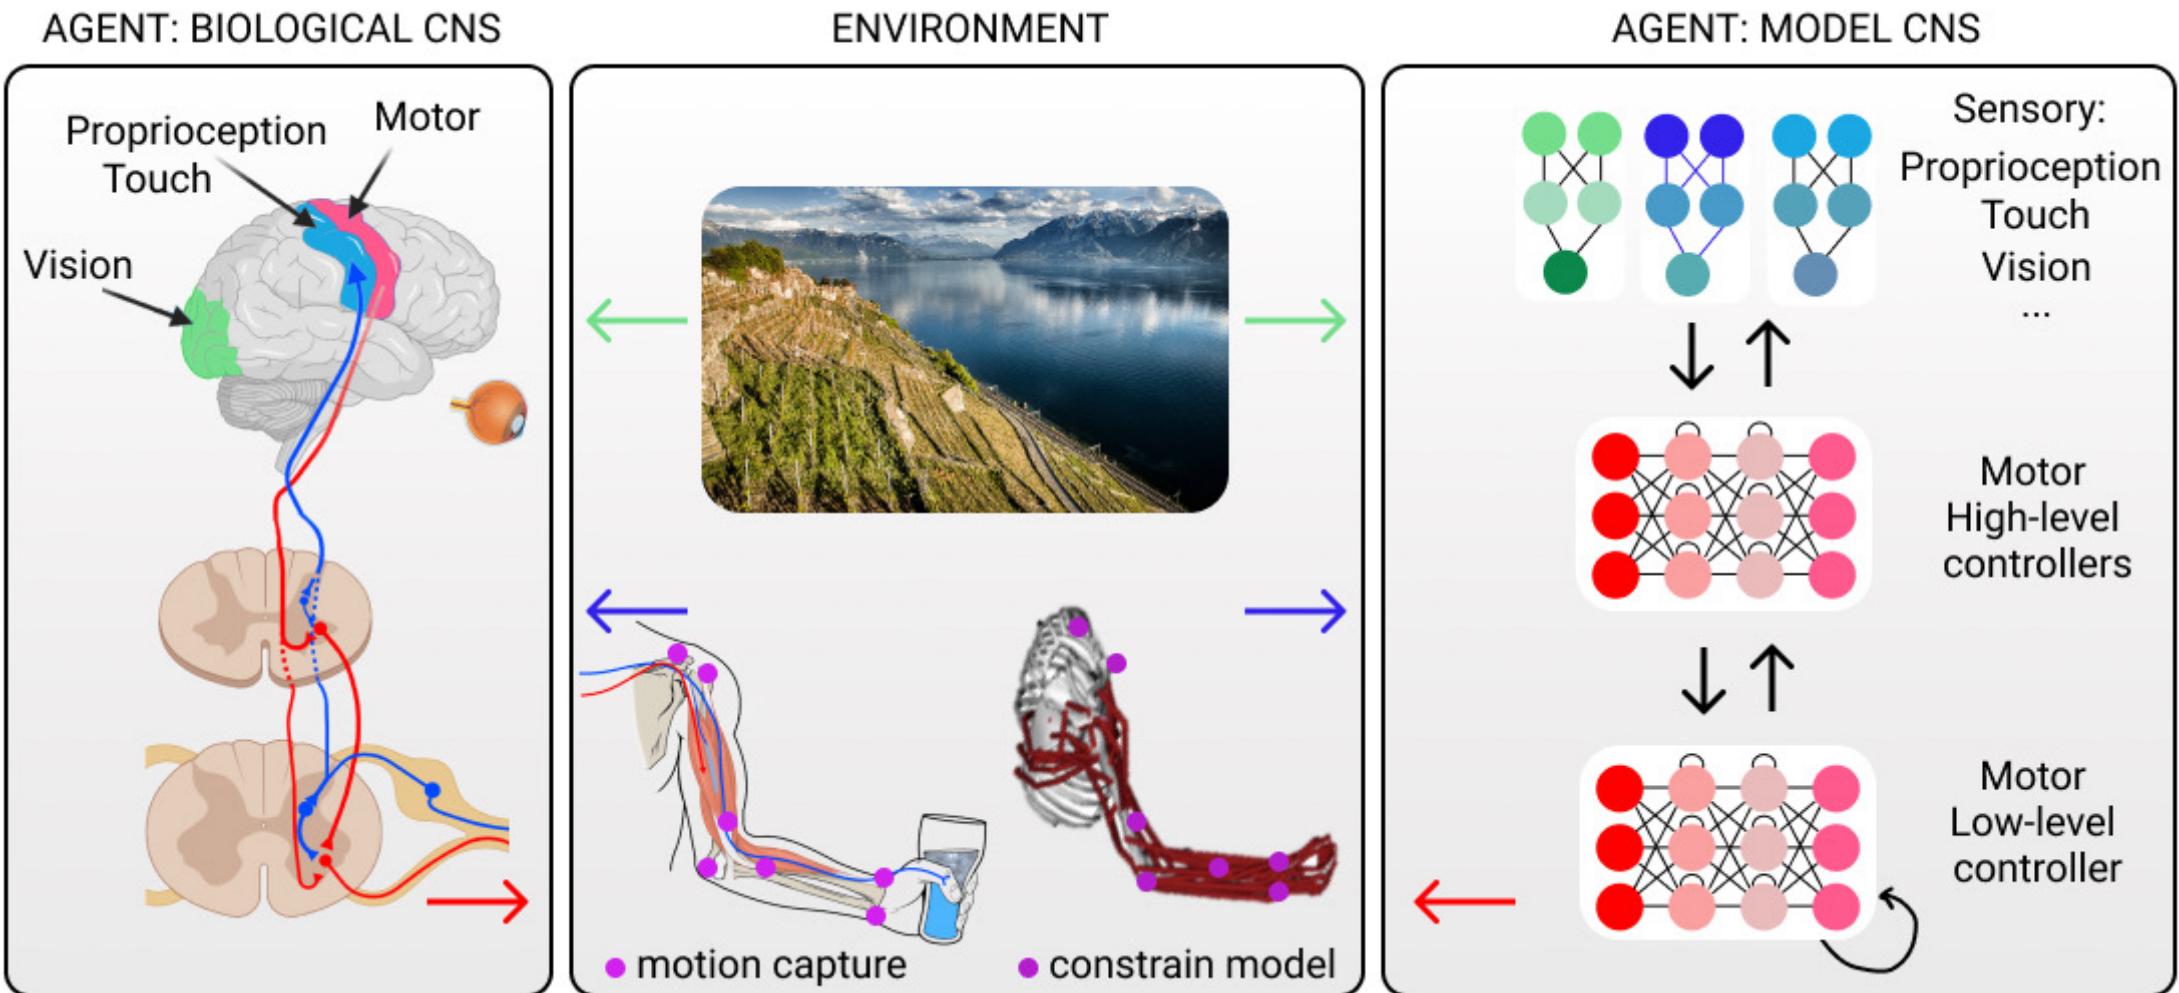


9h of continuous recording

With Dr. Will Menegas &  
Prof. Guoping Feng (MIT)

Lauer et al. BioRxiv 2021  
Nature Methods 2022

# Modeling of behavior with Deep Learning



**Box 2 | Parameters of interest in the network configuration file, *pose\_cfg.yaml***

Please note, there are more parameters that typically never need to be adjusted; they can be found in the default *pose\_cfg.yaml* file at [https://github.com/AlexEMG/DeepLabCut/blob/master/deeplabcut/pose\\_cfg.yaml](https://github.com/AlexEMG/DeepLabCut/blob/master/deeplabcut/pose_cfg.yaml).

- **display\_iters**: An integer value representing the period with which the loss is displayed (and stored in *log.csv*).
- **save\_iters**: An integer value representing the period with which the checkpoints (weights of the network) are saved. Each snapshot has >90 MB, so not too many should be stored.
- **init\_weights**: The weights used for training. Default: <DeepLabCut\_path>/Pose\_Estimation\_Tensorflow/pretrained/resnet\_v1\_50.ckpt. For ResNet-50 or 101, -- this will be automatically created. The weights can also be changed to restart from a particular snapshot if training is interrupted, e.g., <full path>-snapshot-5000 (with no file-type ending added). This would re-start training from the loaded weights (i.e., after 5,000 training iterations, the counter starts from 0).
- **multi\_step**: These are the learning rates and number of training iterations to perform at the specified rate. If the users want to stop before 1 million, they can delete a row and/or change the last value to be the desired stop point.
- **max\_input\_size**: All images larger with size width × height > max\_input\_size\*max\_input\_size are not used in training. The default is 1500 to prevent crashing with an out-of-memory exception for very large images. This will depend on your GPU memory capacity. However, we suggest reducing the pixel size as much as possible; see Mathis and Warren<sup>27</sup>.

The following parameters allow one to change the resolution:

- **global\_scale**: All images in the dataset will be rescaled by the following scaling factor to be processed by the convolutional neural network. You can select the optimal scale by cross-validation (see discussion in Mathis et al.<sup>12</sup>). Default is 0.8.
- **pos\_dist\_thresh**: All locations within this distance threshold (measured in pixels) are considered positive training samples for detection (see discussion in Mathis et al.<sup>12</sup>). Default is 17.

The following parameters modulate the data augmentation. During training, each image will be randomly rescaled within the range [scale\_jitter\_lo, scale\_jitter\_up] to augment training:

- **scale\_jitter\_lo**: 0.5 (default).
- **scale\_jitter\_up**: 1.5 (default).
- **mirror**: If the training dataset is symmetric around the vertical axis, this Boolean variable allows random augmentation. Default is False.
- **cropping**: Allows automatic cropping of images during training. Default is True.
- **cropratio**: Fraction of training samples that are cropped. Default is 40%.
- **minsize**, **leftwidth**, **rightwidth**, **bottomheight**, **topheight**: These define dimensions and limits for auto-cropping.

# Questions?

Please get in touch!

[alexander.mathis@epfl.ch](mailto:alexander.mathis@epfl.ch)

---



Open Source Code: <https://github.com/DeepLabCut/>  
Tutorials, Data, Papers etc.: <https://deeplabcut.org>  
User forum: <https://forum.image.sc/tag/deeplabcut>