

DeepLabCut tutorial

Society of Integrative and Comparative Biology
Phoenix, Arizona

**NOTE: some videos have been
removed from this pdf version!**



Mathis Group
computational neuroscience & AI

Alexander Mathis

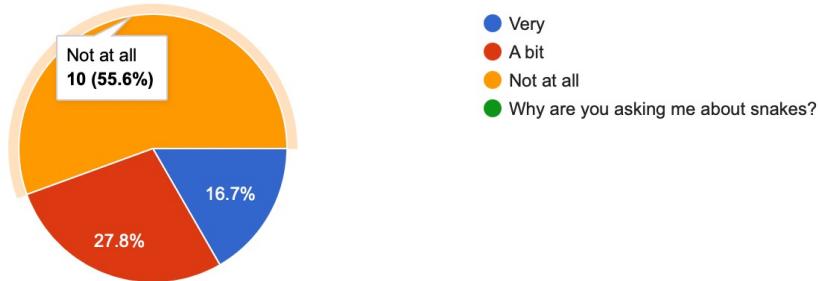
Swiss Federal Institute of Technology Lausanne /
École Polytechnique Fédérale de Lausanne (EPFL)

Jan 7th, 2022

Your feedback

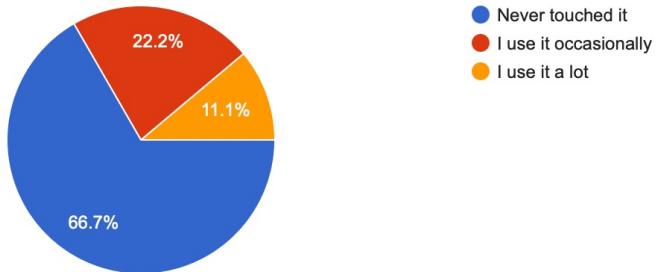
How experienced are you with Python?

18 responses



How experienced are you with DeepLabCut?

18 responses



What is your main application for DeepLabCut?

17 responses



Your goals ...

- I would like to get a basic understanding of how DeepLabCut works
- 3D implementation
- How to run DLC successfully and get animal tracking data from it
- Basic functions/introduction
- How to train it for very high accurate digitization.
- How to prepare files and train dataset
- Going through the entire process of setting up, training, collecting data
- Basics!
- Basic run-through of steps and especially how to fix poor training- do you need to start all over if it doesn't track well?
- General workflow, tips for beginners (if suitable for students). Also, suggestions on hardware
- Things I haven't tried yet, like real time and 3D
- How to use this software to track small, low-contrast animals in low-light settings.
- How to track/quantify specific movements and behaviors of a single animal
- How to use COLAB, options when tracking two different organisms, best approach when trying to train multiple models from recordings that all have one organism in common but a different second interacting organism
-

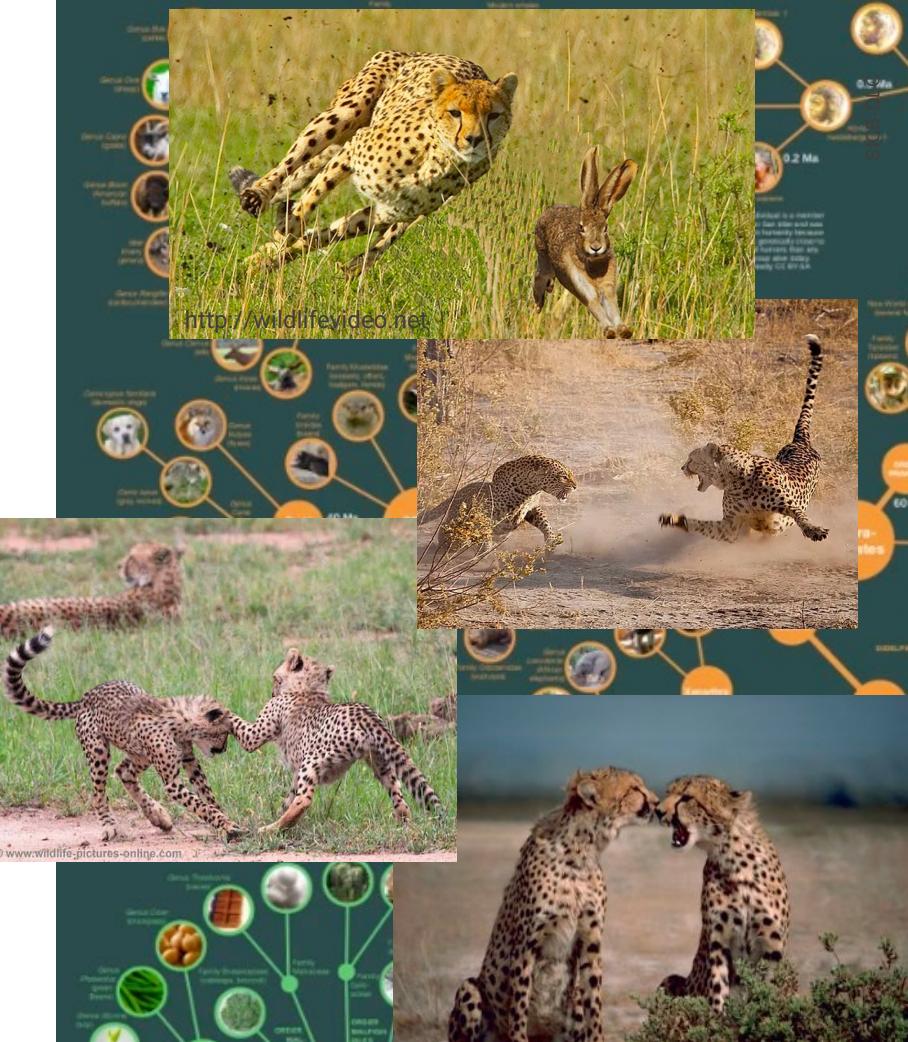
- I would like to get a basic understanding of how DeepLabCut works
- 3D implementation
- How to run DLC successfully and get animal tracking data from it
- Basic functions/introduction
- How to train it for very high accurate digitization
- How to prepare files and train dataset
- Going through the entire process of setting up, training, collecting data
- Basics
- Basic run-through of steps and especially how to fix poor training- do you need to start all over if it doesn't track well?
- General workflow, tips for beginners (if suitable for students). Also, suggestions on hardware
- Things I haven't tried yet, like real time and 3D
- How to use this software to track small, low-contrast animals in low-light settings
- How to track/quantify specific movements and behaviors of a single animal
- How to use COLAB options when tracking two different organisms, best approach when trying to train multiple models from recordings that all have one organism in common but a different second interacting organism
-

Massive diversity of animals

Behavior is highly diverse

We need to reduce the problem to meaningful quantifications

1/7



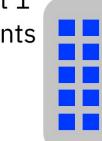
Measuring behavior: Pose estimation

Pixel Representation

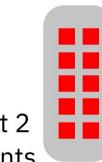


Pose Estimation
Algorithm

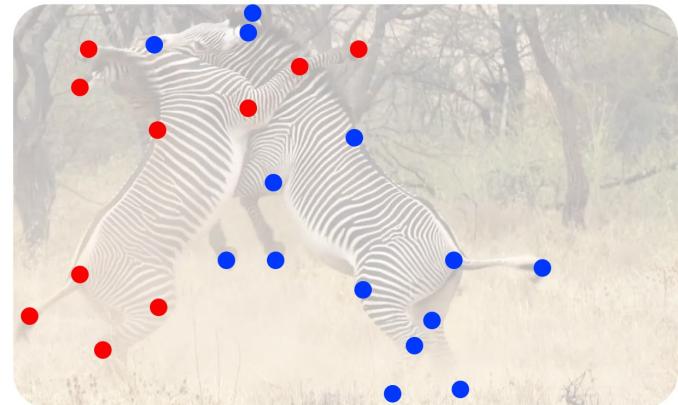
Subject 1
Keypoints



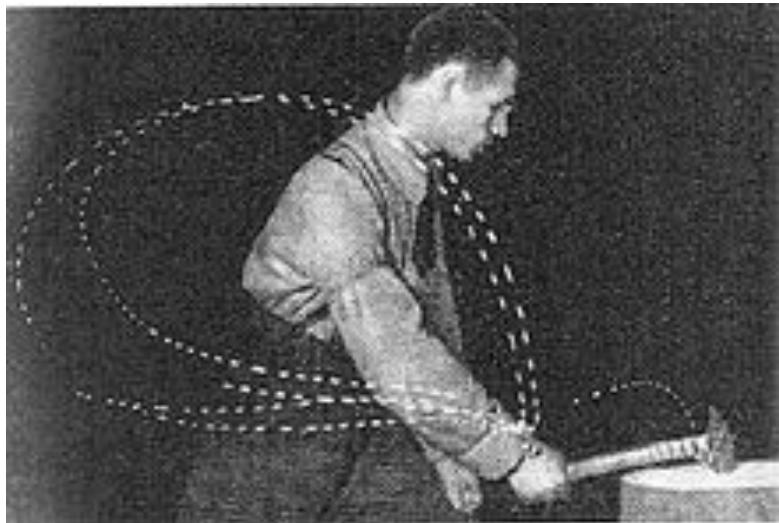
Subject 2
Keypoints



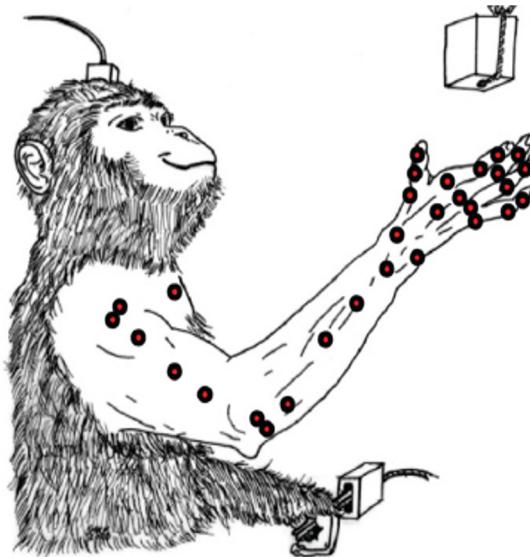
Keypoint Representation



Marker-based pose estimation



Bernstein, 1920



Vargas-Irwin et al. 2010

Deep learning for human pose estimation



MoDeep – starting 2014
DeepPose
Conv. PoseMachines

...
DeeperCut
OpenPose

...
> 10,000 papers on human pose estimation with Deep Learning

Appealing properties:

- Work “in the wild”
- Robust (backgrounds, individuals, etc)
- Relatively fast
- No body model
- No (manual) parameter tuning

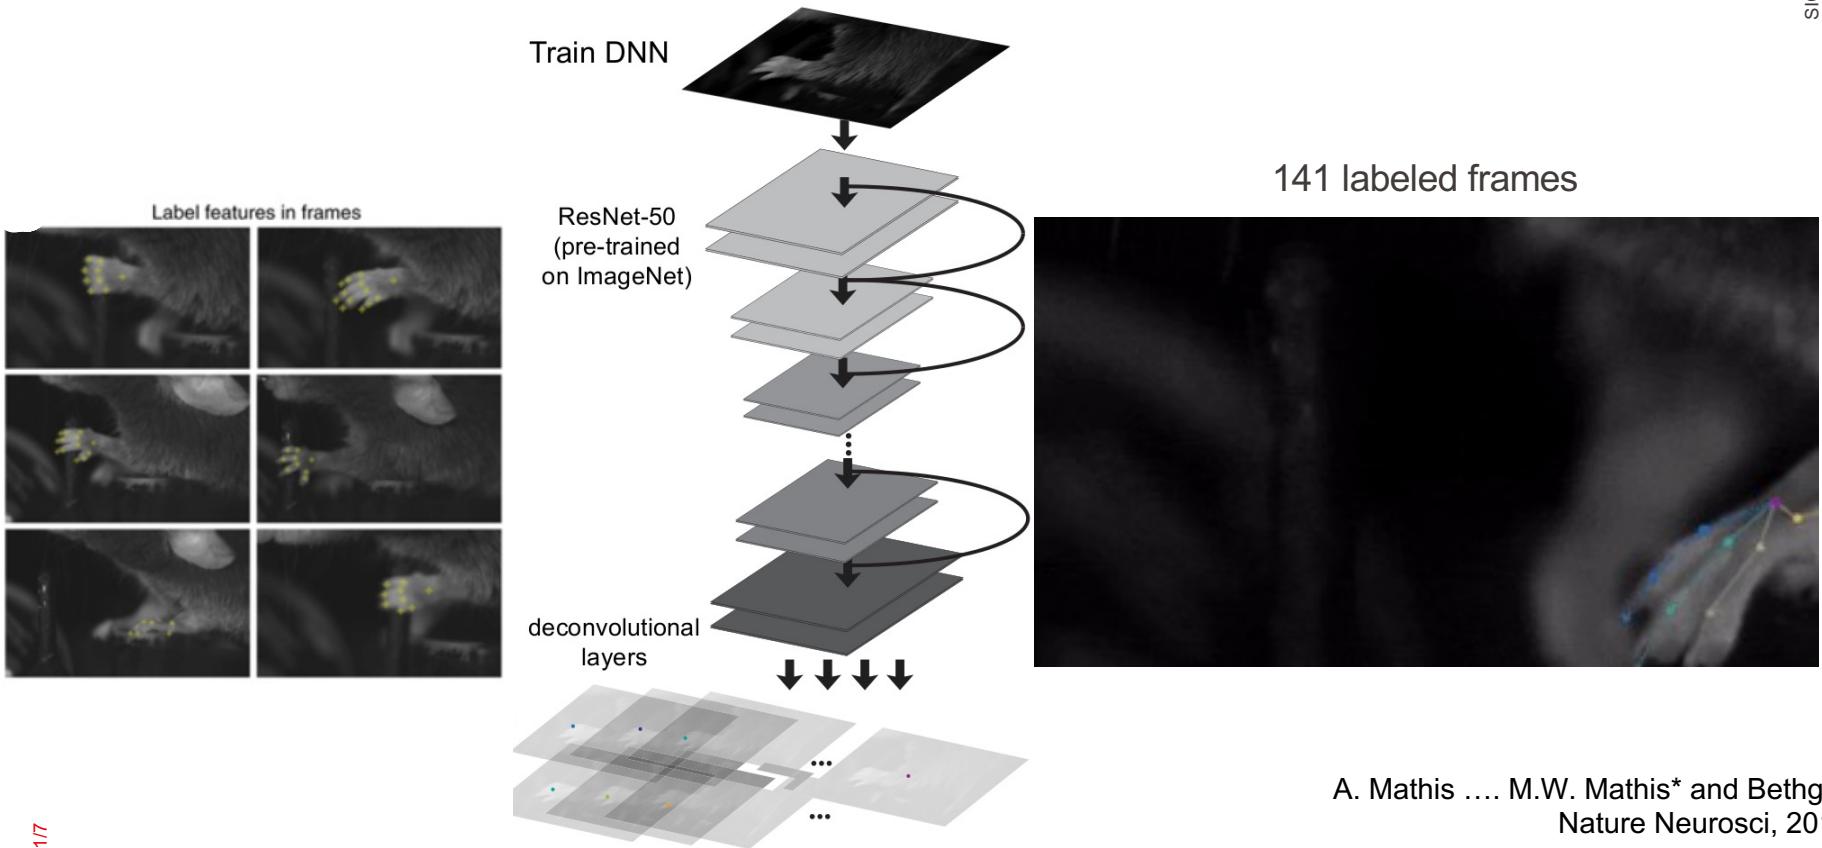
deep neural networks

image → Predictor → pose

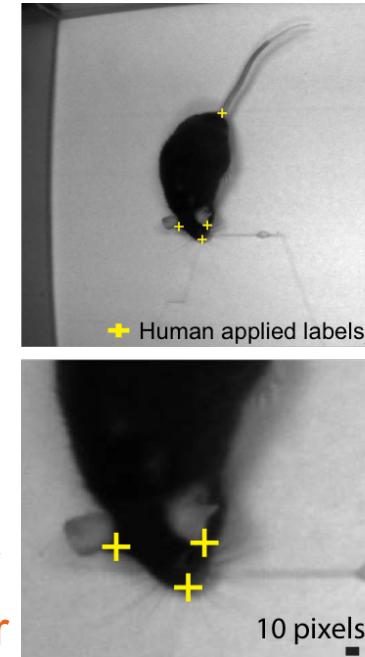
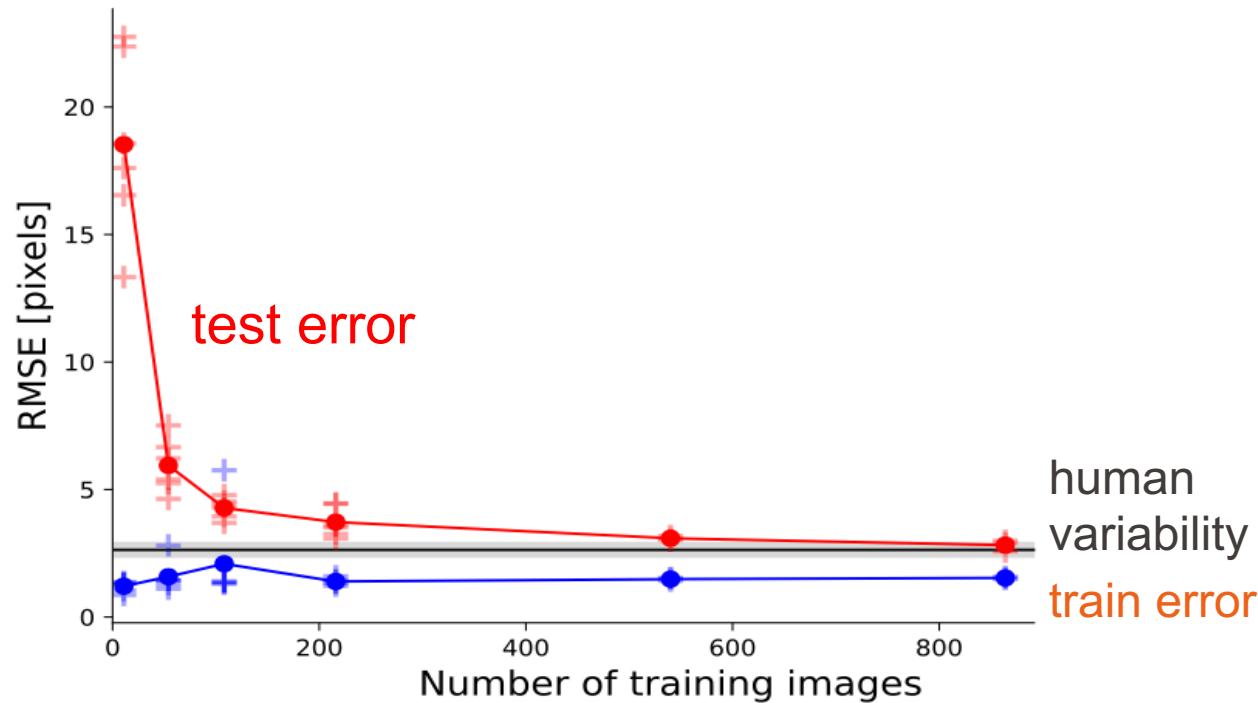
train

A lot of labeled images ($> 10^6$ joints!)

DeepLabCut: a toolbox for markerless pose estimation via transfer learning



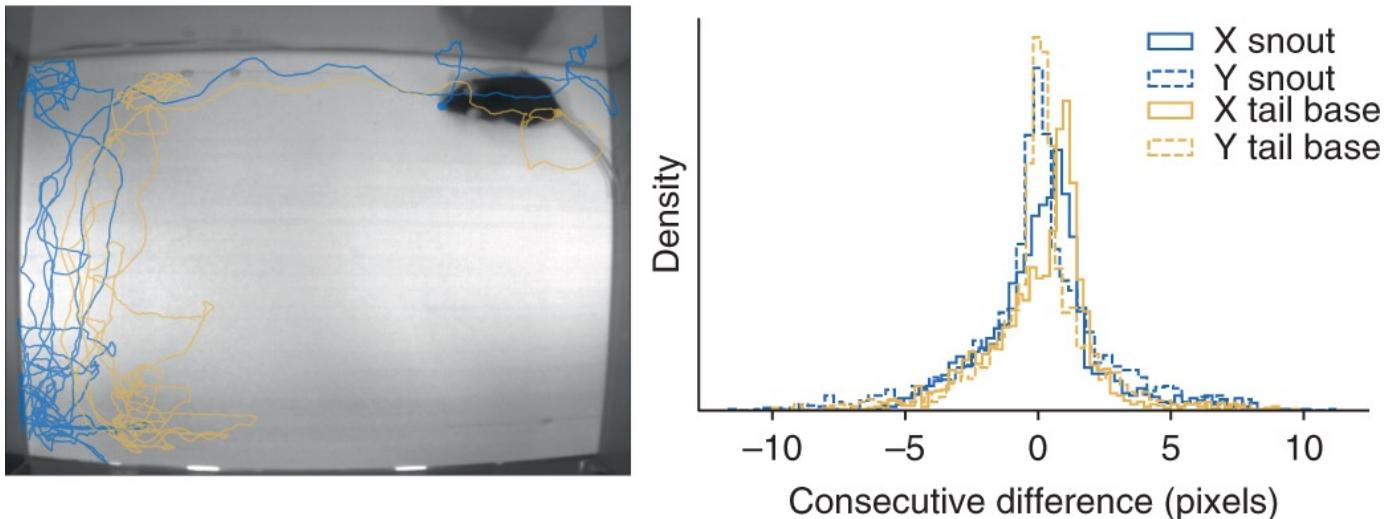
How little training data is required?



A. Mathis M.W. Mathis* and Bethge*
Nature Neurosci, 2018

Generalization to novel mice

Example of generalization: novel mouse tracking



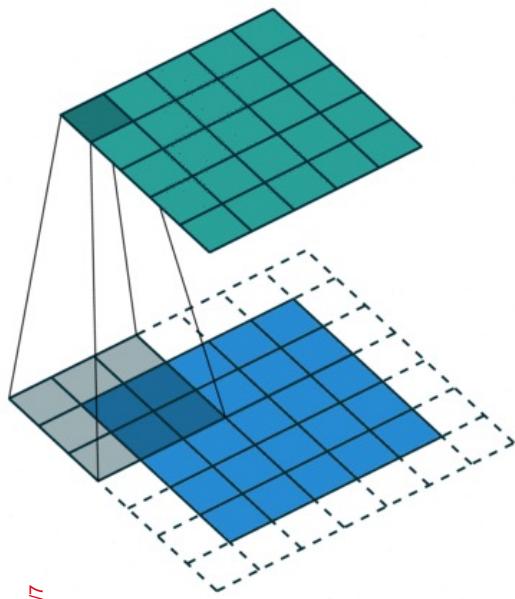
A. Mathis M.W. Mathis* and Bethge*
Nature Neurosci, 2018

Let's dig into the nitty gritty

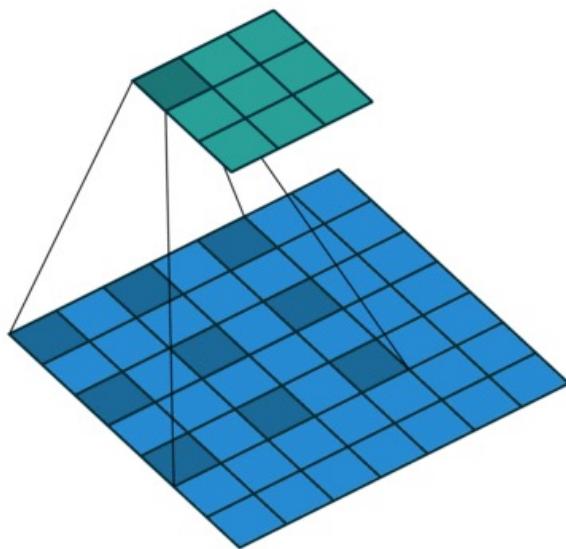
Machine learning system:

- Architecture
- Objective (task)
- Learning algorithm
- Data

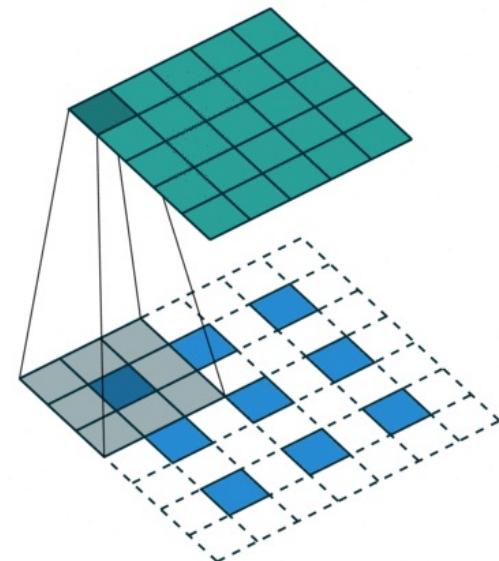
Basic convolution (“same”)



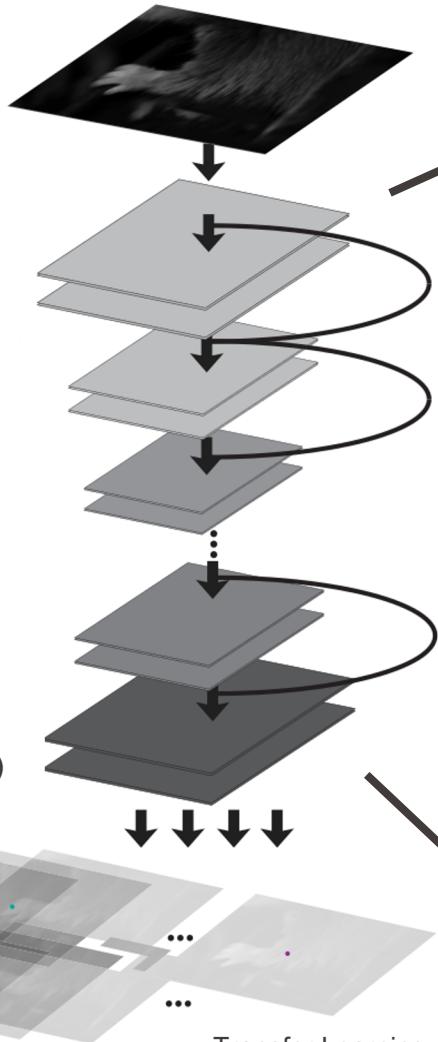
Strided convolution



Strided deconvolution



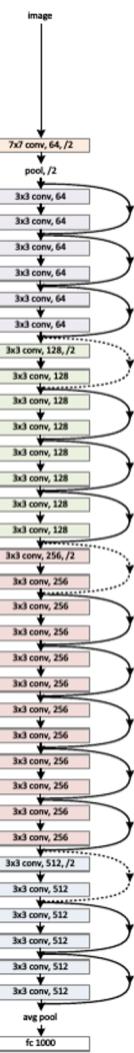
DeepLabCut
Based on
DeeperCut
Insafutdinov et al
CVPR 2017



ResNet

He et al.,
CVPR 2016

convolutions



Imagenet: 1.2 million images
for 1000 classes

Deng et al. CVPR 2009

Train to predict hammer (for those inputs)
As well as other object classes

Transfer Learning

EPFL Keypoint Detector (Segmentation mask)

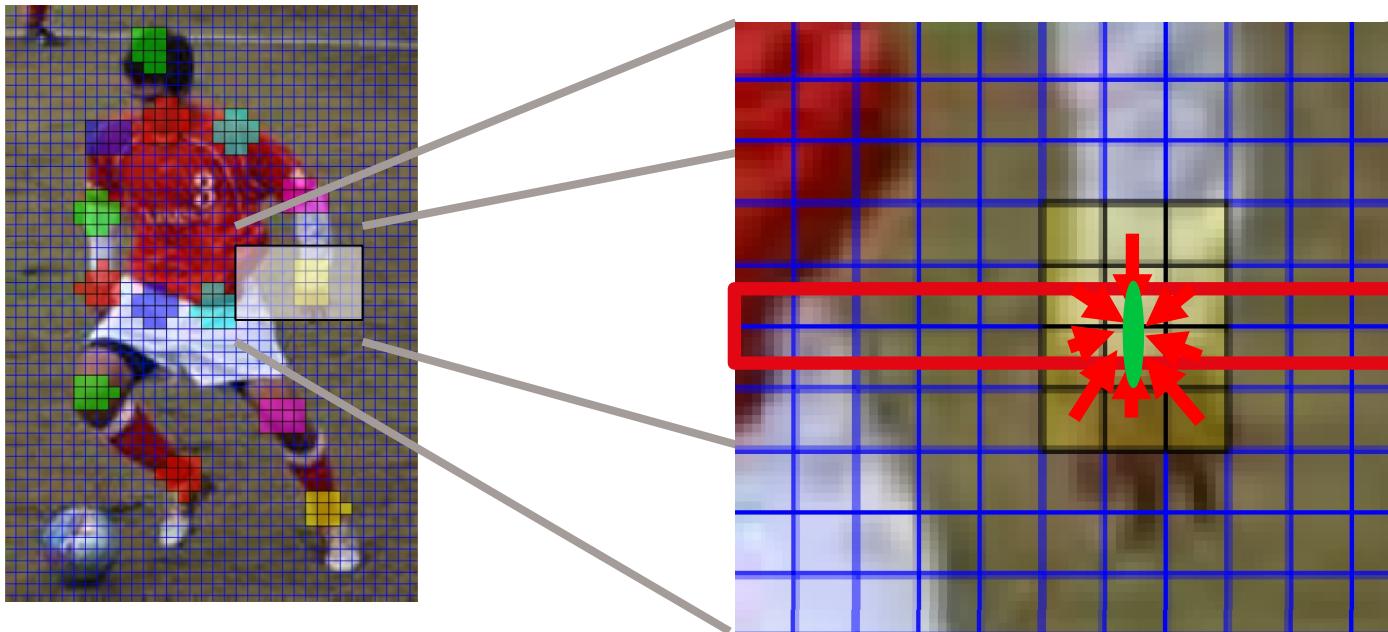
- Given: a ground truth coordinate (x_j, y_j) for the body joint j
- Construct a ground truth scoremap $s_j(x, y)$



$$s_j(x, y) = \begin{cases} 1, & \text{if } \|(x, y) - (x_j, y_j)\|_2 \leq R \\ 0, & \text{otherwise} \end{cases}$$

Precise Keypoint Localization (Location Refinement)

- Inspired by bounding box regression in Overfeat [Sermanet et al, 2014] and Fast R-CNN [Girshick, 2015]

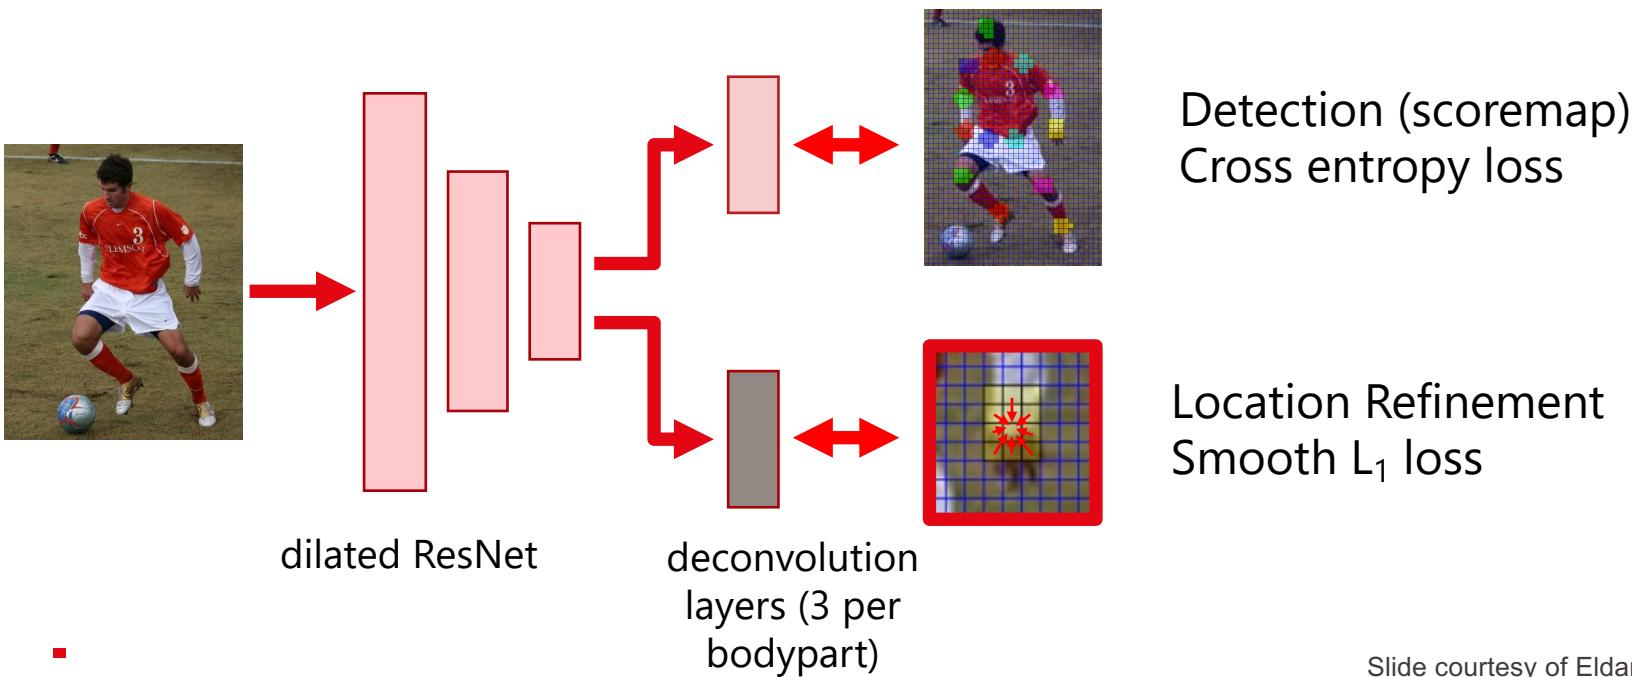


Regress the offsets

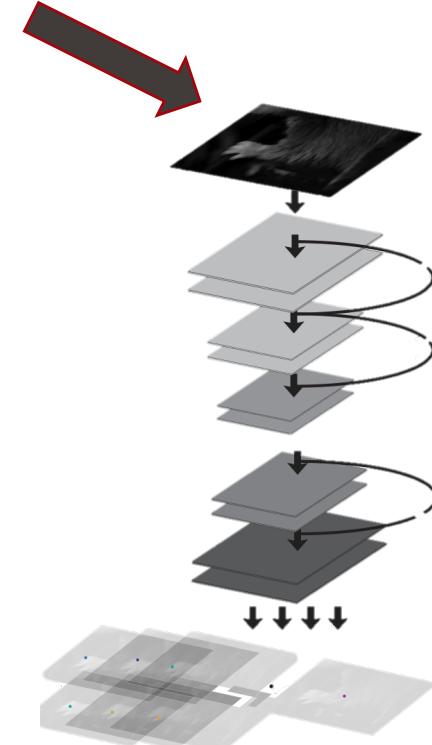
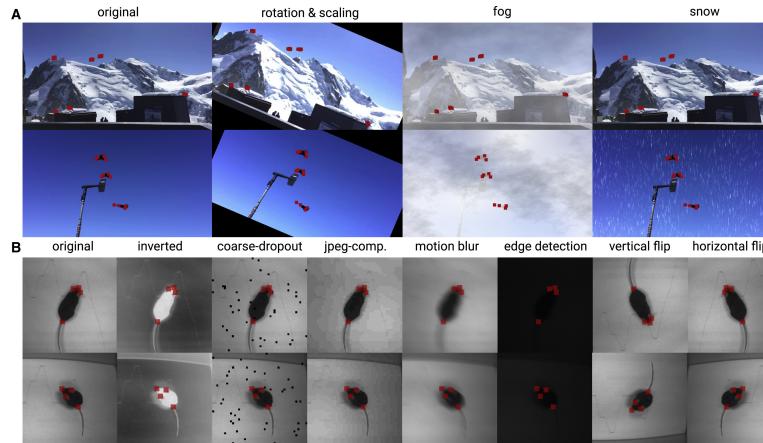
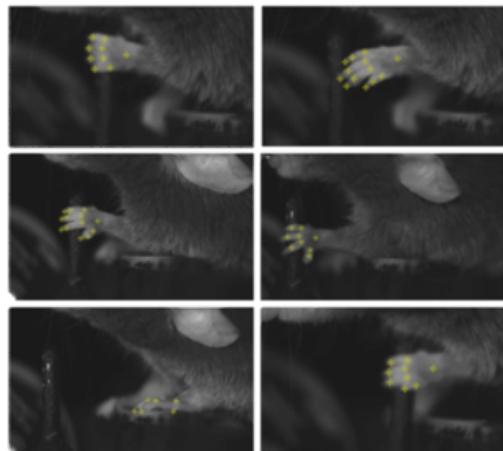
$$\Delta x = x_{part} - x$$
$$\Delta y = y_{part} - y$$

Multi-Task ConvNet

- Minimize the combined loss with SGD (with Momentum)
- Perform augmentation during training



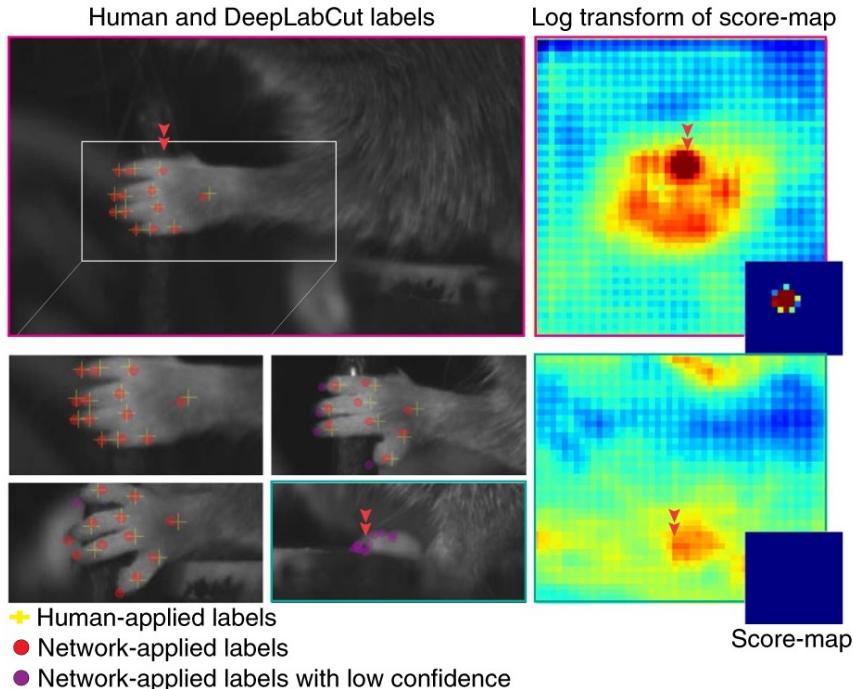
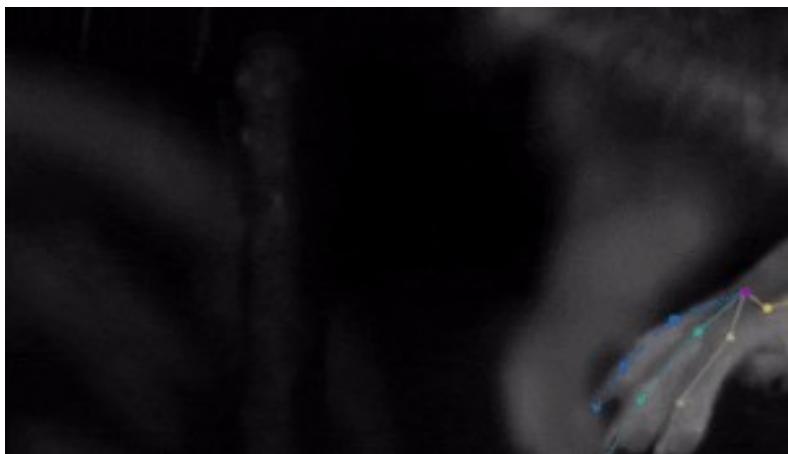
How does markerless pose estimation work?



Key Features:

- Data augmentation
- Model architecture
- Optimization

Confidence readout (score-maps)



A. Mathis M.W. Mathis* and Bethge*
Nature Neurosci, 2018

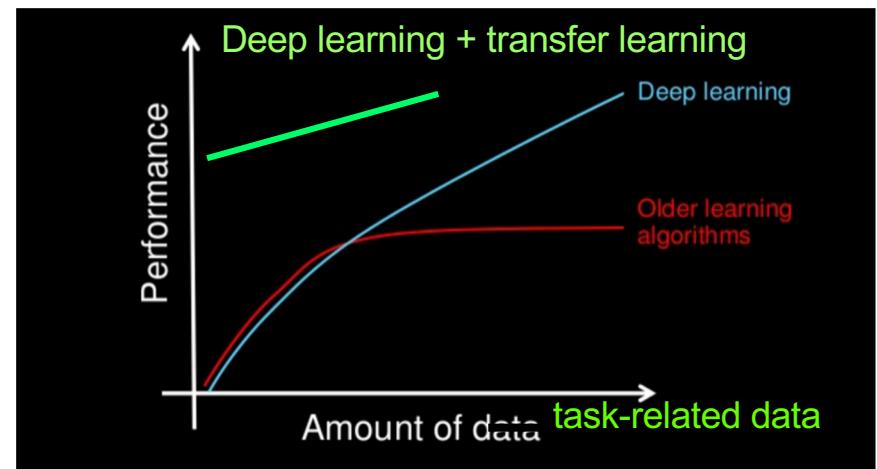
Transfer learning enables pose estimation with less data

Pre-trained! (i.e. on ImageNet)

image → Predictor → pose

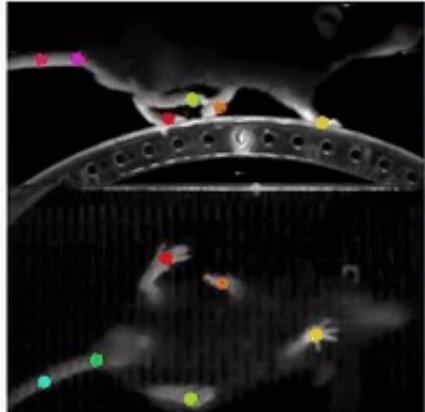
train

Only a few examples (10-200)
for most applications

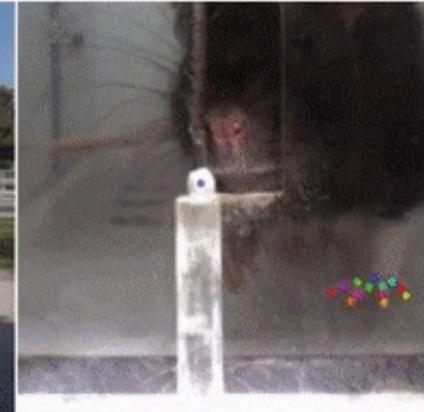
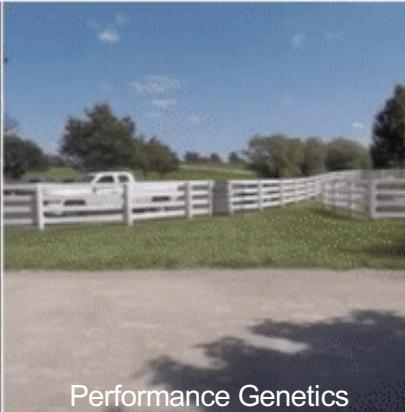


Andrew Ng

Locomotion
3D tracking of limbs



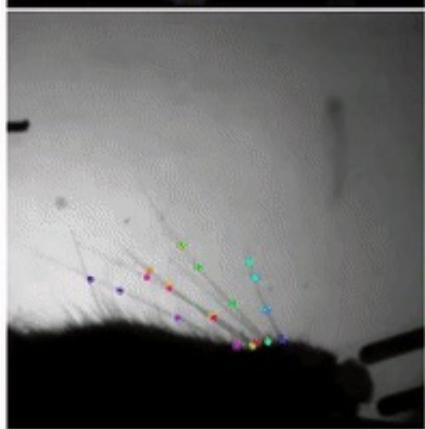
R. Warren (Sawtell lab)
Columbia University



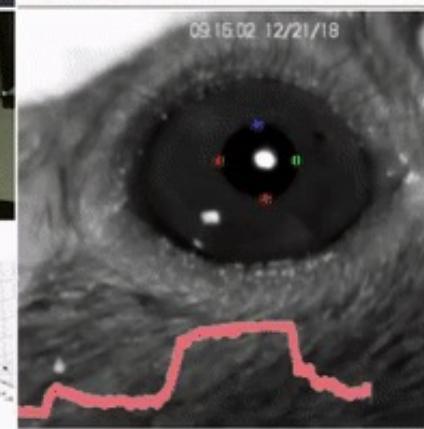
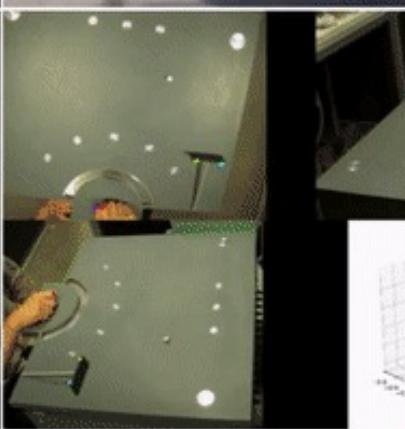
Pellet reaching task

D. Leventhal Lab
University of Michigan

Performance Genetics



Whisker tracking
A Erskine (Hires lab)
USC



Pupil tracking
T. Vaissiere
Scripps

J Bonaiuto (Ferrari lab) CNRS

DeepLabCut – Deep Learning toolbox for pose-estimation
■ Open source & free, > 1,000 labs, >200,000 installations

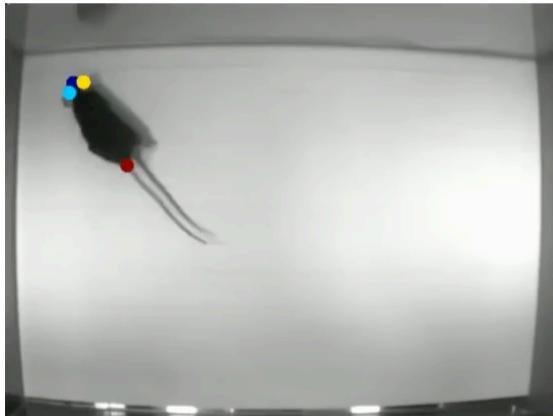


DeepLabCut:
a software package for
animal pose estimation

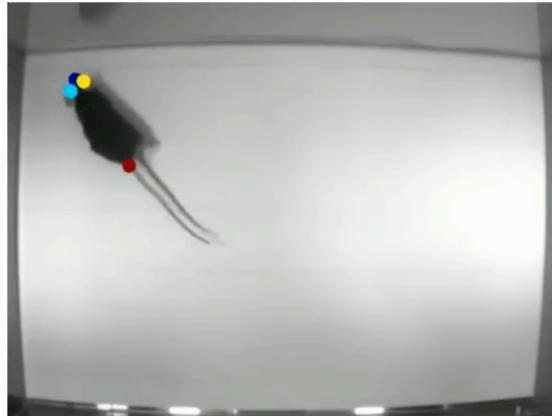
use our Jupyter Notebooks, Google Colab, or work in the terminal!



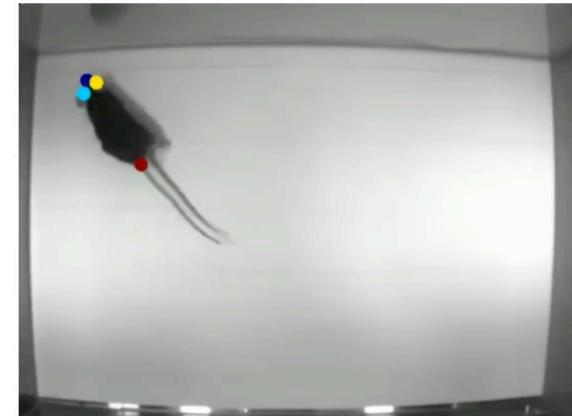
Data Augmentation: how to get the most out of your data!



tensorpack



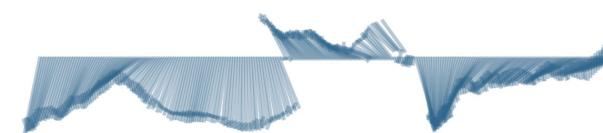
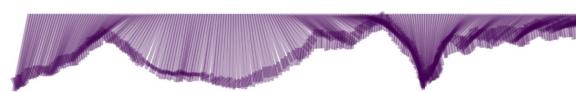
imgaug



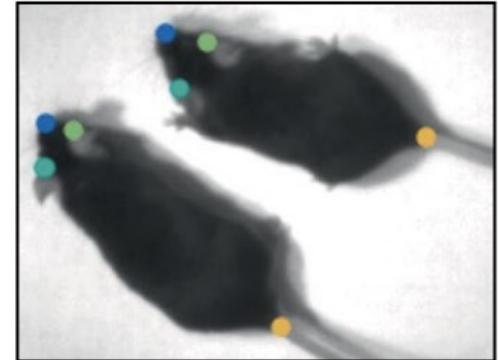
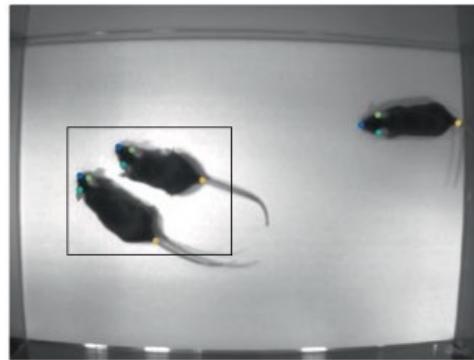
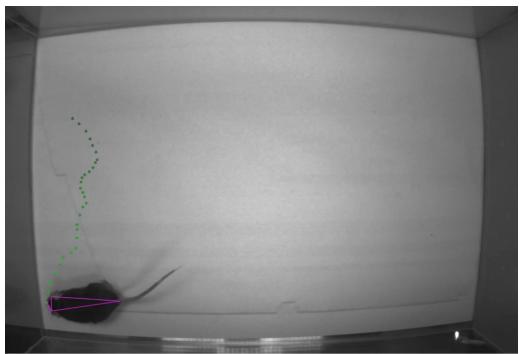
scale-crop

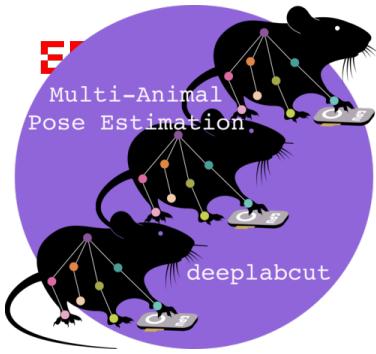


1/17

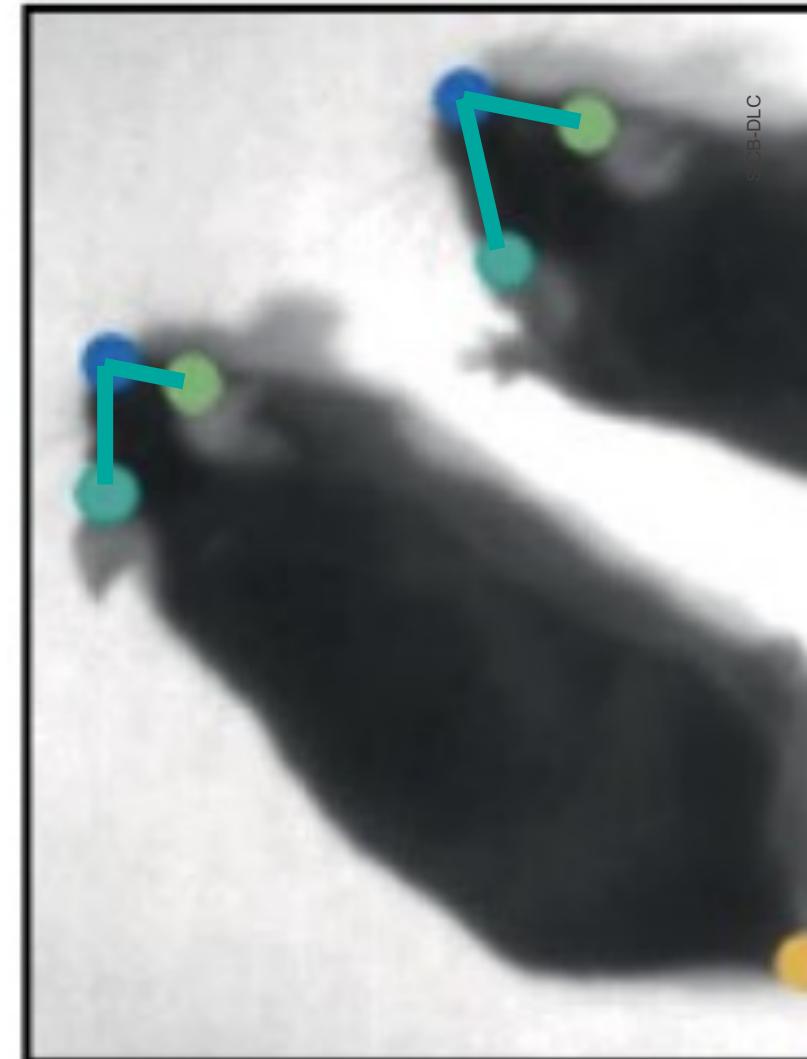


Pose estimation for multiple animals



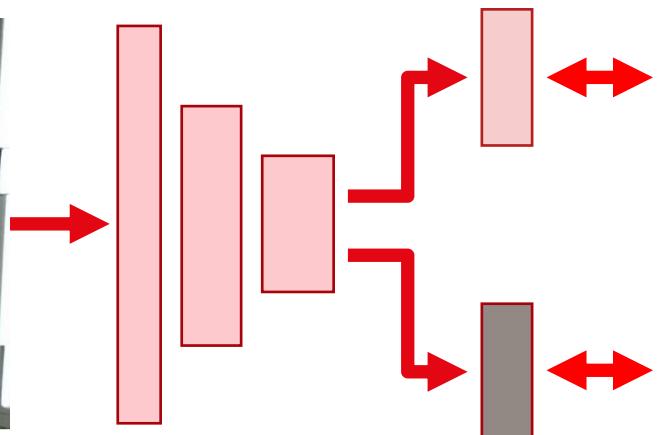


The idea: learn to predict "limbs"



EPFL Multi-Task Deep Convolutional Network

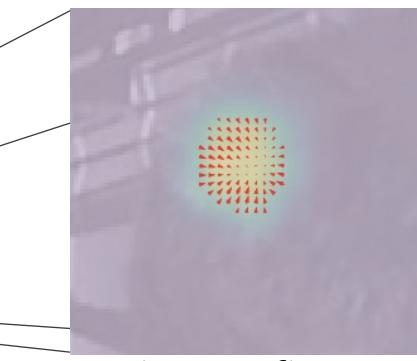
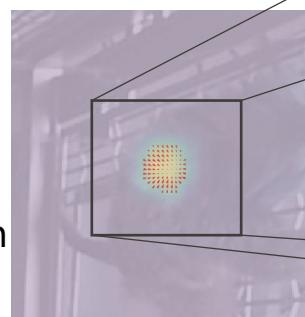
- Minimize the combined loss; segmentation + vector field
- Transfer learning (ConvNet pre-trained on object recognition)
- Perform augmentation during training



deconvolution
layers (3 per
bodypart)

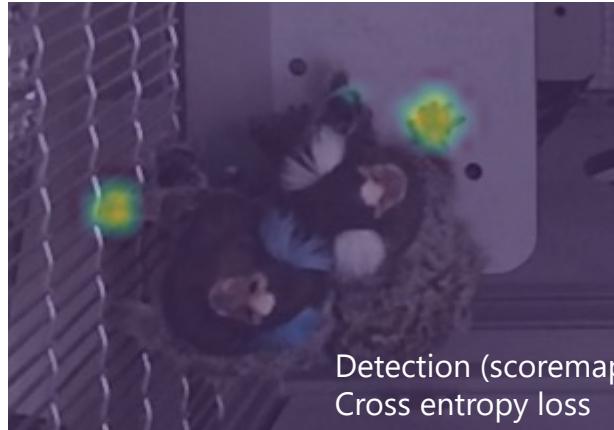
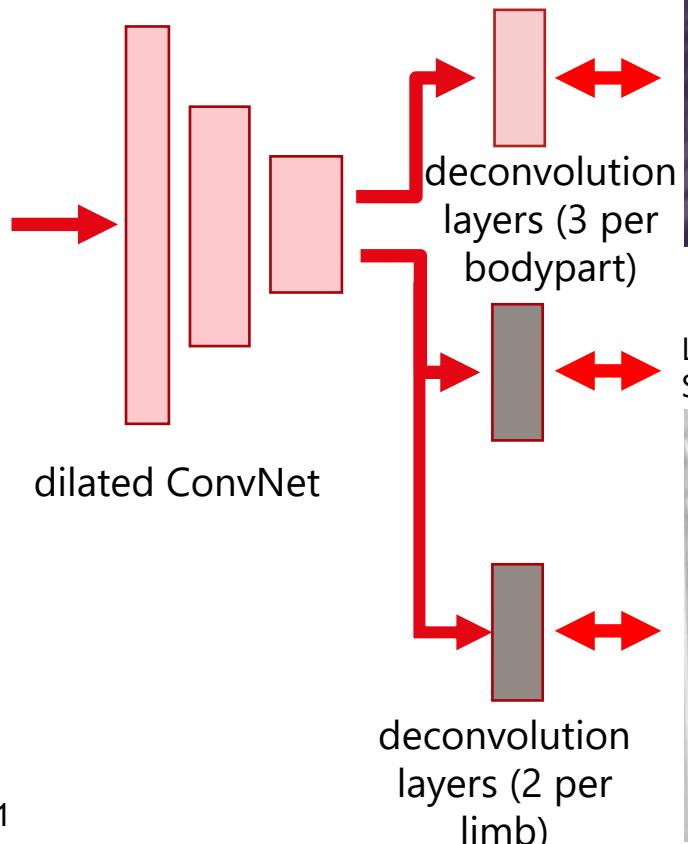


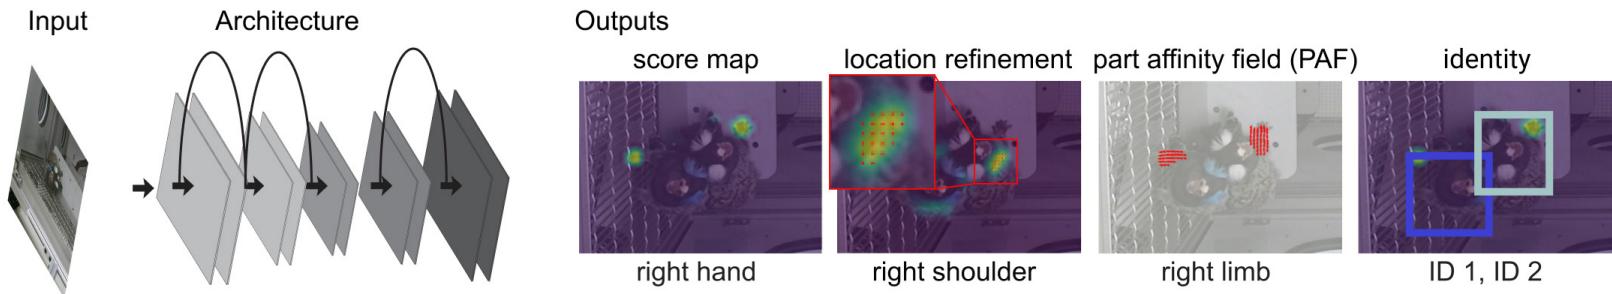
Detection (scoremap)
Cross entropy loss



Location Refinement
Smooth L₁ loss

Multi-Task DCN (DLC 2.2)



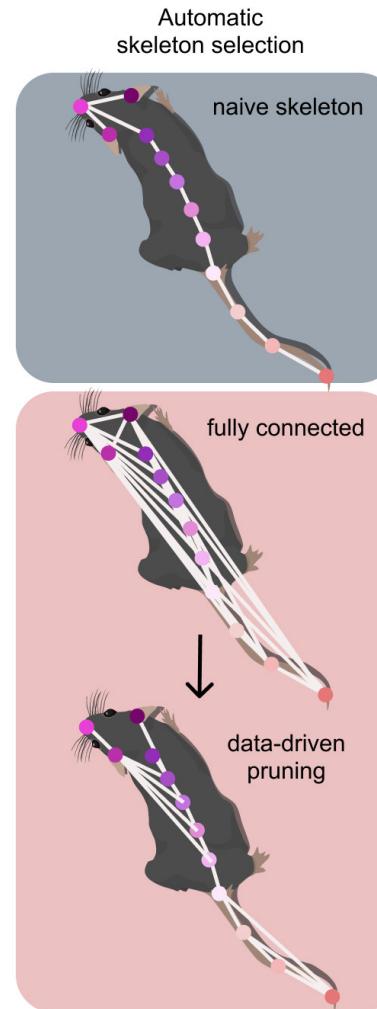


Bodyplan agnostic graph optimization

Learn to predict:

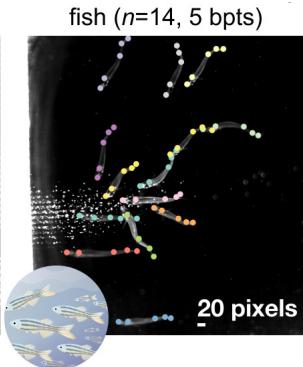
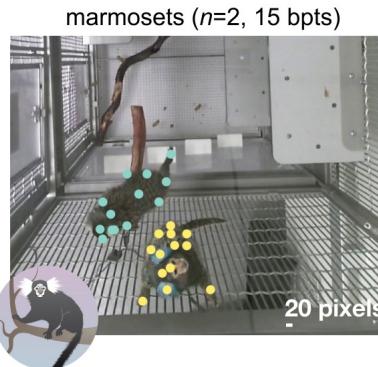
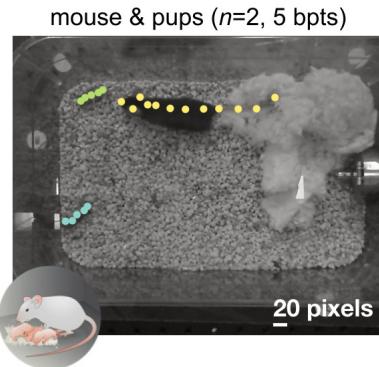
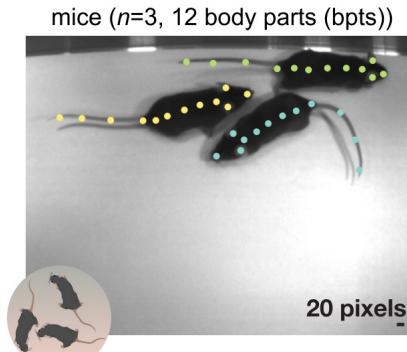
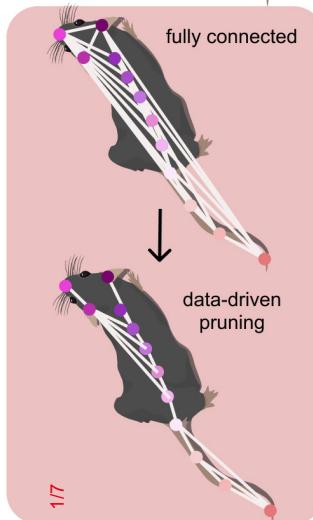
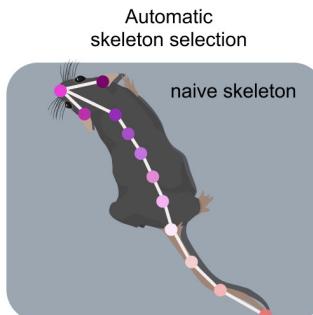
- Bodyparts
- Limbs (= costs for associations)

Which limbs should be learned?

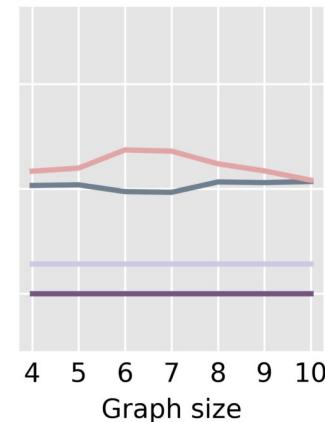
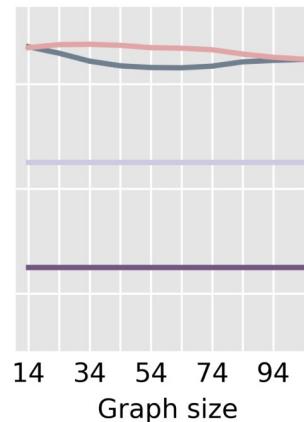
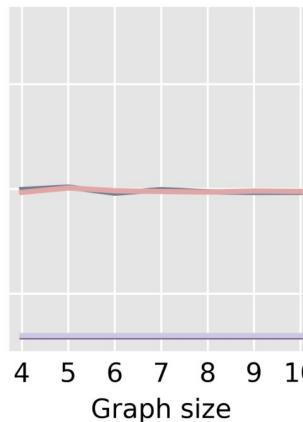
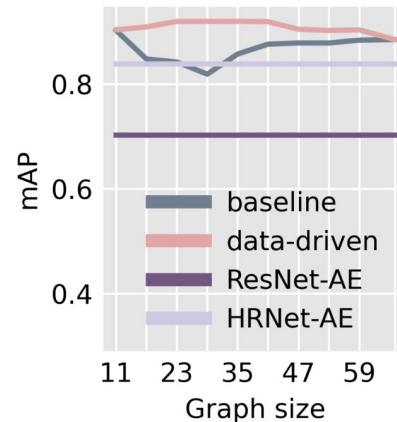


EPFL Novel data driven method outperforms baseline methods (giving SOTA perf)

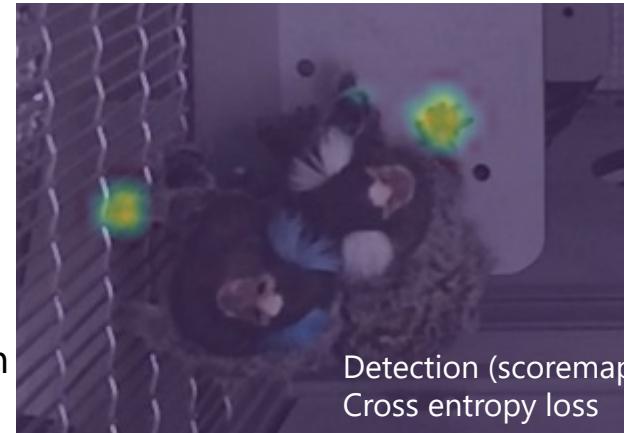
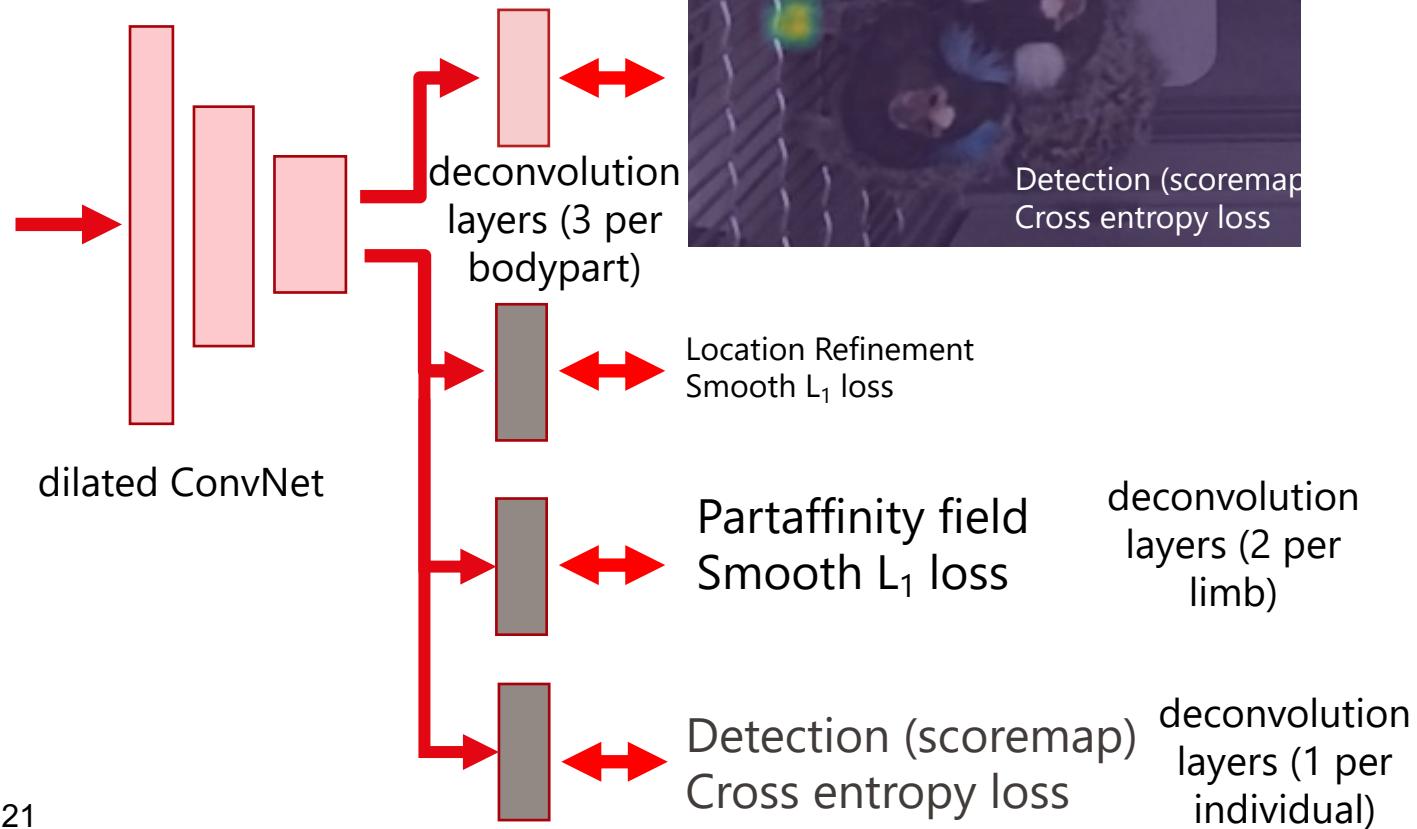
ICB-DLC



mAP performance

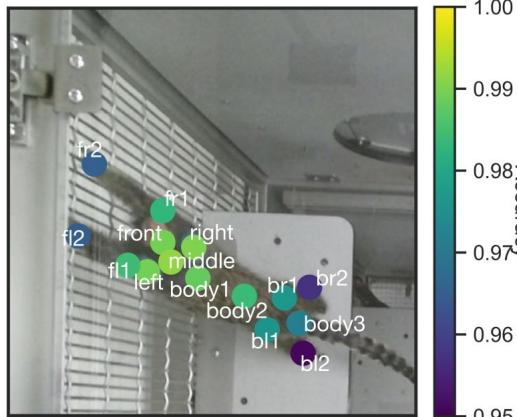
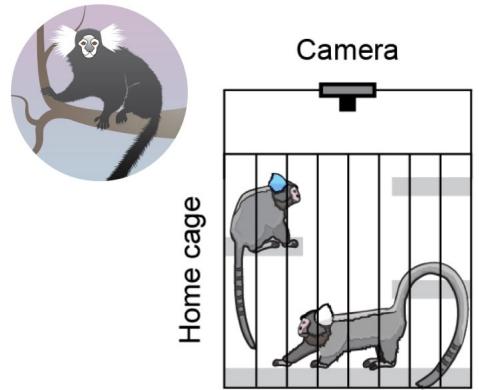


Multi-Task DCN (DLC 2.2)



Detection (scoremap)
Cross entropy loss

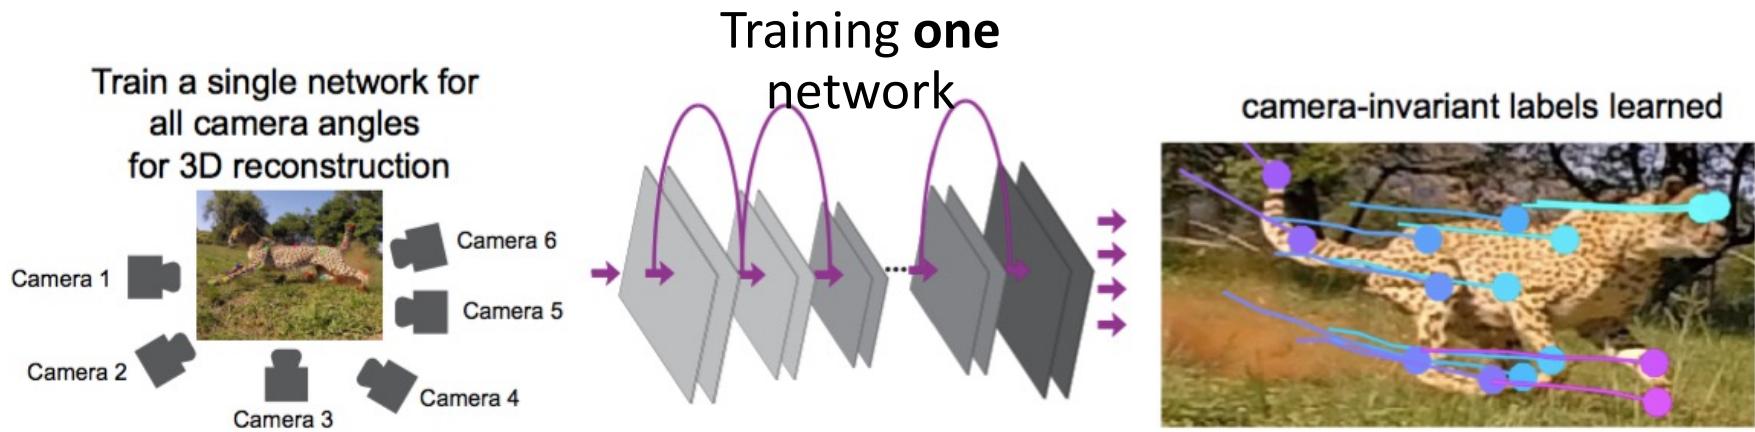
Individual identification performance



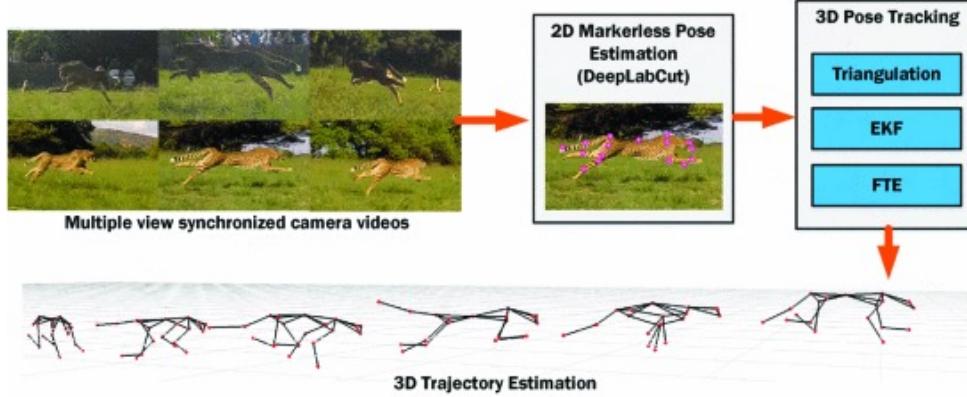
Identification of individuals

3D pose estimation?

3D-pose estimation of Cheetahs in the “wild”



AcinoSet: A 3D Pose Estimation Dataset and Baseline Models for Cheetahs in the Wild



Pose Type	Metric	TRI	EKF	FTE
Run	RMSE	28.24	3.40	2.76
	SEM	0.26	0.03	0.03
	NRMSE	0.17	0.02	0.02
Dive	RMSE	76.35	39.40	38.44
	SEM	0.70	0.36	0.35
	NRMSE	0.56	0.29	0.28

EKF: extended Kalman filter
FTE: full trajectory estimation

Can DeepLabCut be applied in real-time?

Providing real-time feedback in experiments

- <https://github.com/DeepLabCut/DeepLabCut-live>

#Initialize model:

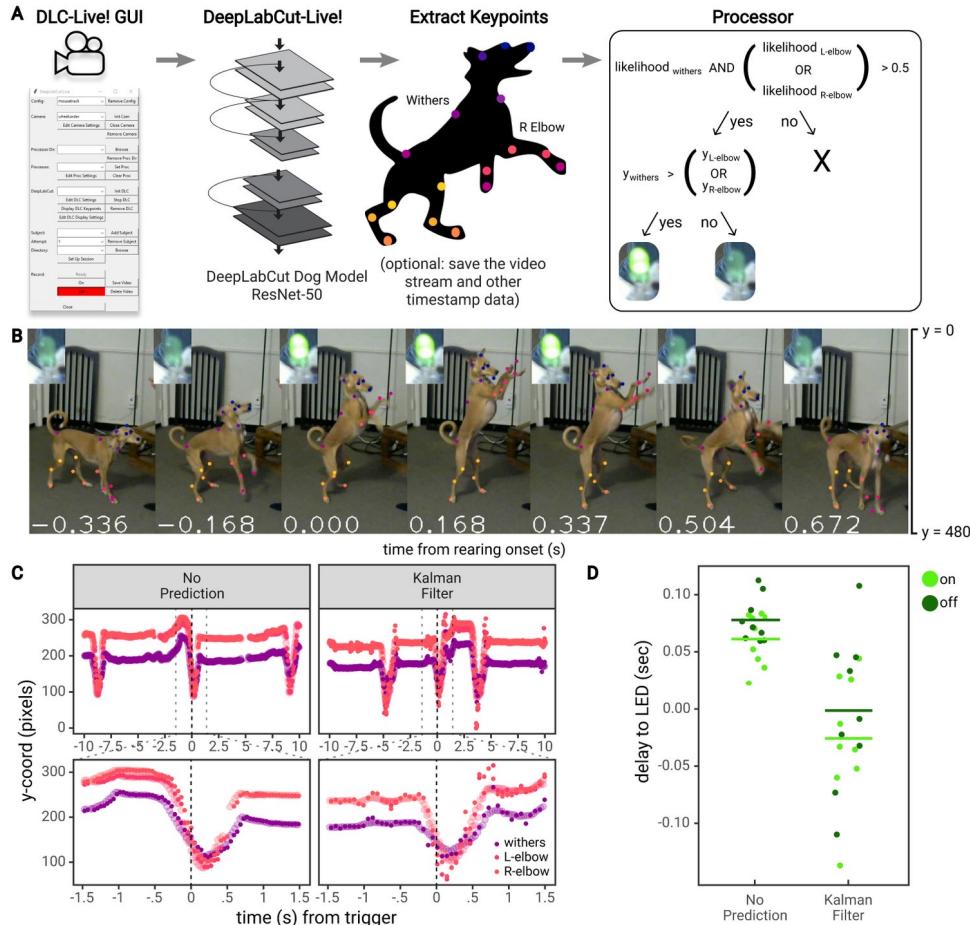
```
from dlclive import DLCLive  
my_live_object = DLCLive("/exportedmodel/directory")
```

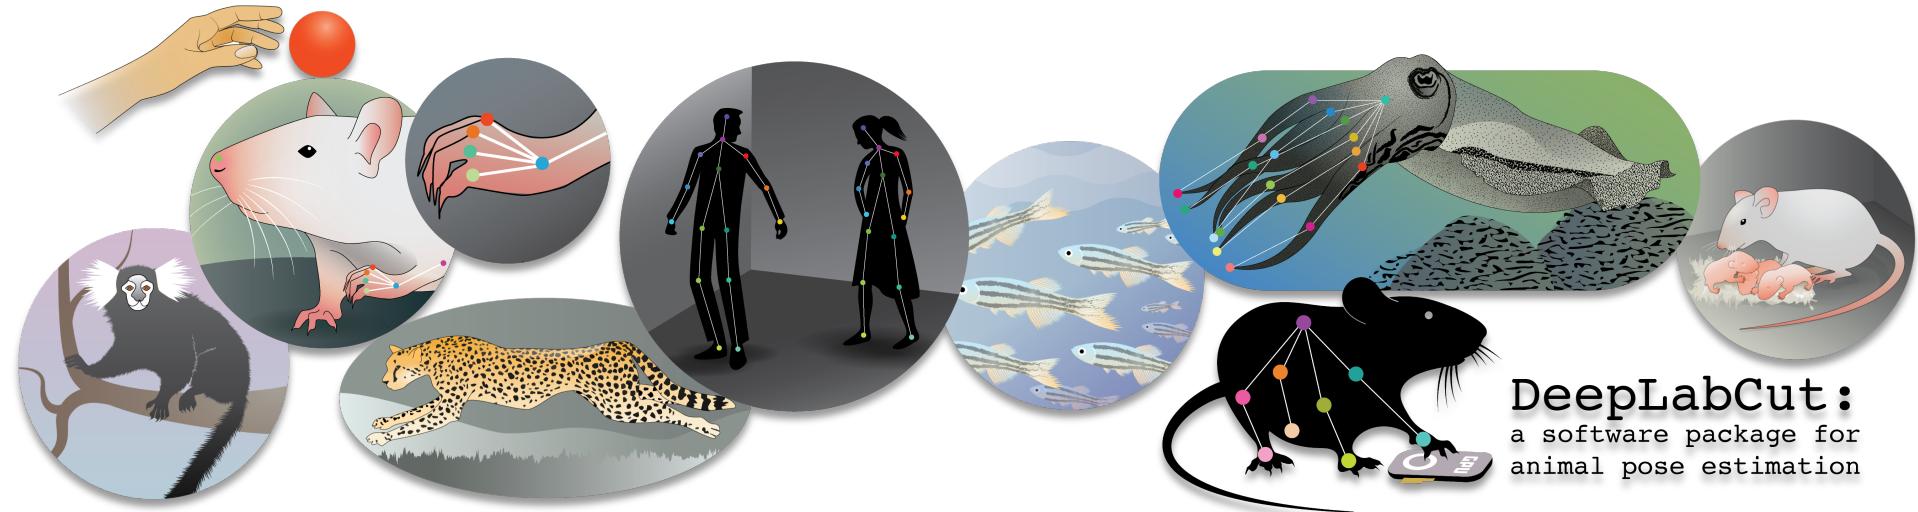
run inference:

```
my_live_object.init_inference(my_image)  
pose = my_live_object.get_pose(my_image)
```



Example experiment





DeepLabCut:
a software package for
animal pose estimation

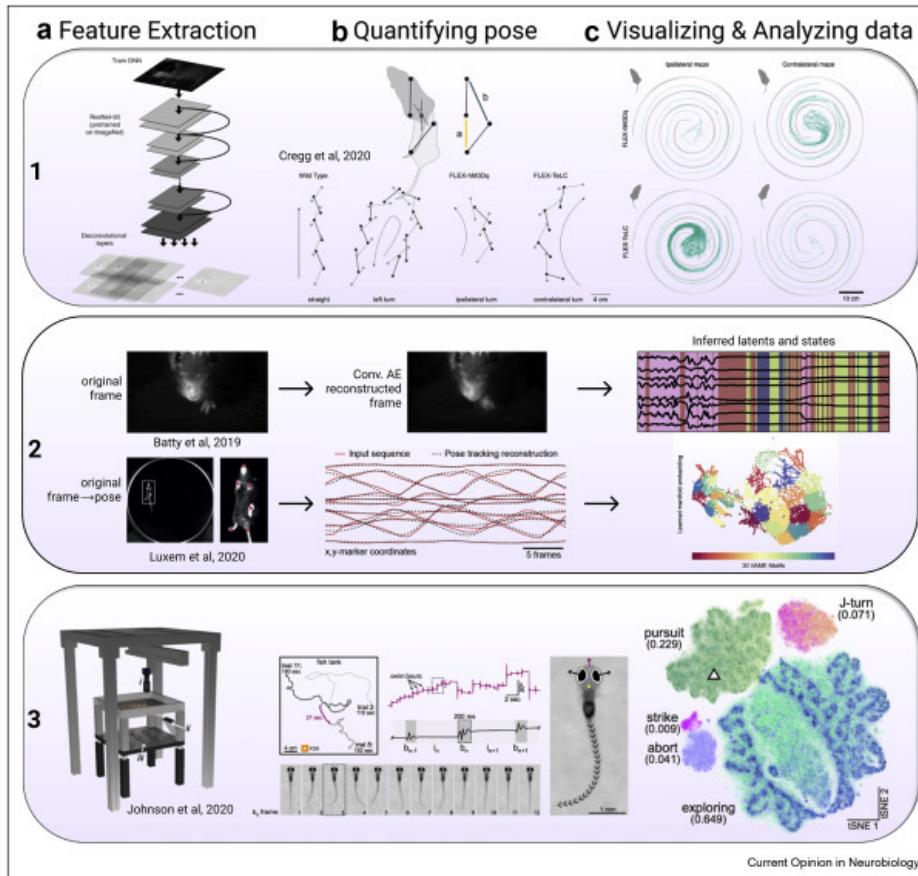
modelzoo.deeplabcut.org

Example model in the DLC model zoo

(added 3 days after the pre-print!)



What to do with DLC output?



How to use DeepLabCut?

DeepLabCut workflow

Train DNN

Create a project,
extract frames, +
GUIs to label your data

Select + Train your
deep neural network

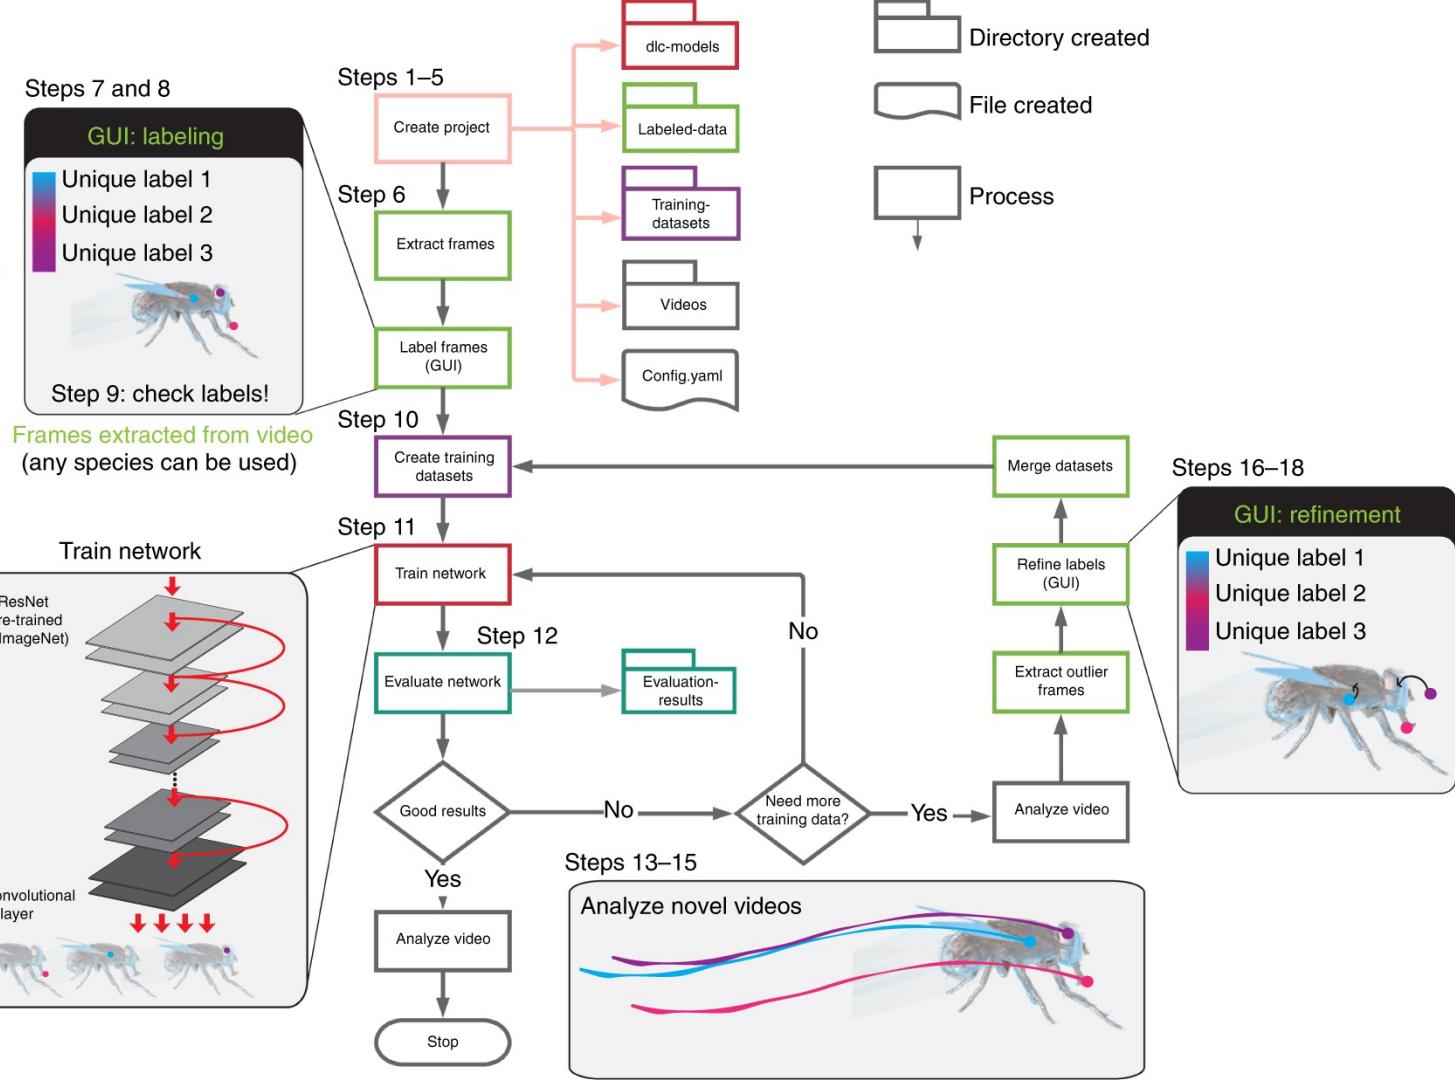
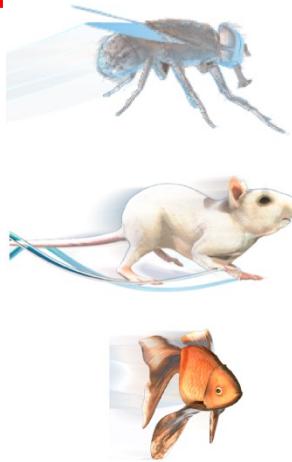
Evaluate network
performance
(active learning + GUIs
if improvement needed)

Run inference on
new videos,
create labeled videos,
+ plot your results!



refine?





nature
protocols

PROTOCOL

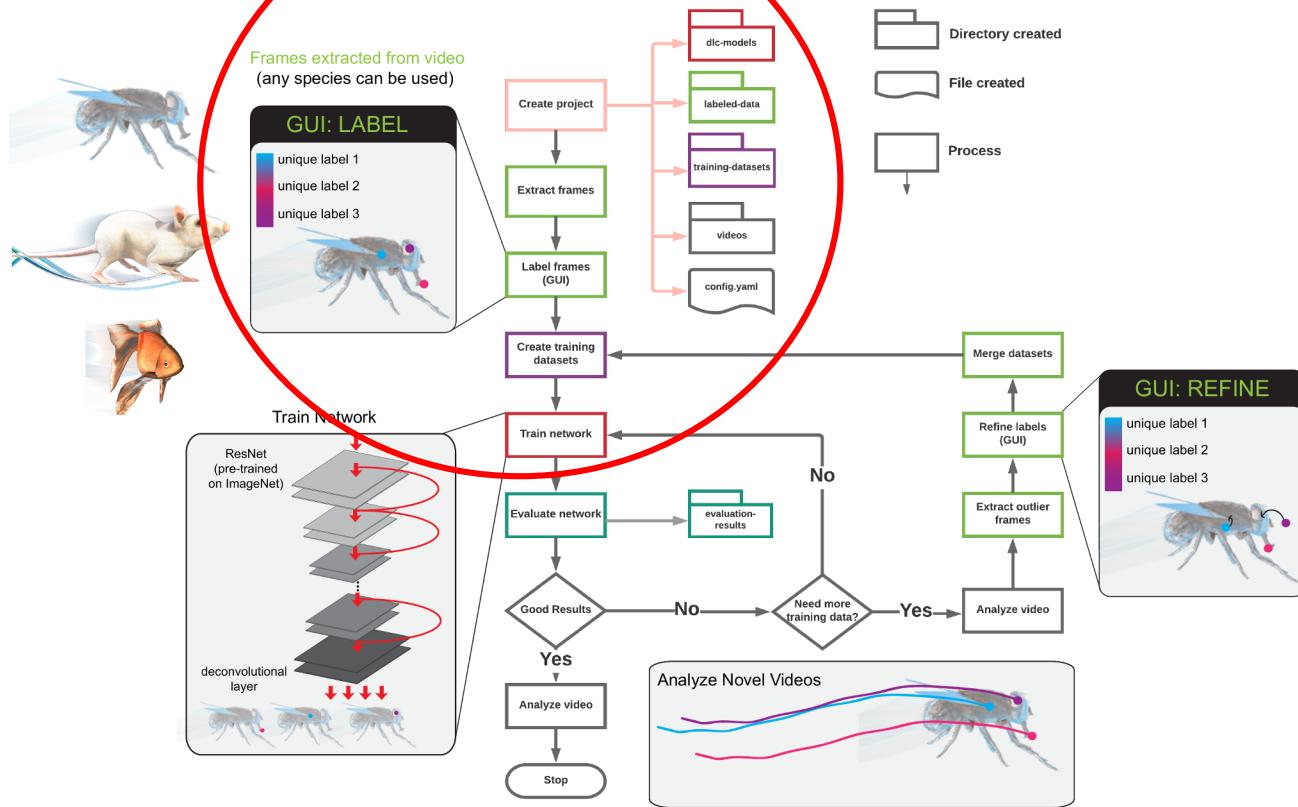
<https://doi.org/10.1038/s41596-019-0176-0>

Using DeepLabCut for 3D markerless pose estimation across species and behaviors

Key commands

Operation	Command
Open IPython and import DeepLabCut (Step 1)	<code>ipython import deeplabcut</code>
Create a new project (Step 2)	<code>deeplabcut.create_new_project('project_name', 'experimenter', ['path of video 1', 'path of video2', ...])</code>

DeepLabCut 2.0 workflow



2D Project Folders



- ▶ training data,
trained networks,
`config.yaml` file

This is the master folder created when
you create a project (Step 1)

New (2D) videos for analysis



- ▶ batch process videos
- ▶ analyzed data

you can place new videos here, and then run:

```
deeplabcut.analyze_videos(config_path,  
                           folderpath, videotype='mp4')
```

Your 2D or 3D project “entry point” is through the `config.yaml` file

When you want to work on your project:

```
activate DLCenvName  
ipython  
import deeplabcut  
config_path = '/home/yourprojectfolder/config.yaml'
```

Frame extraction: deeplabcut.extract_frames(..)

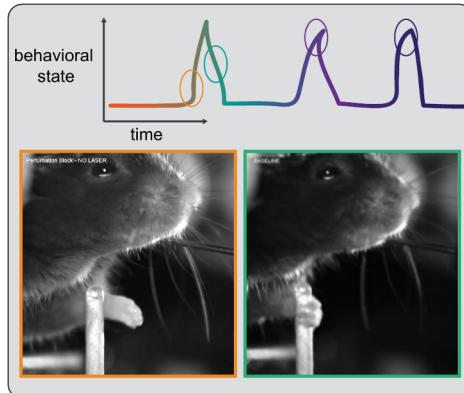
Select videos to grab frames:

Use videos with image from:

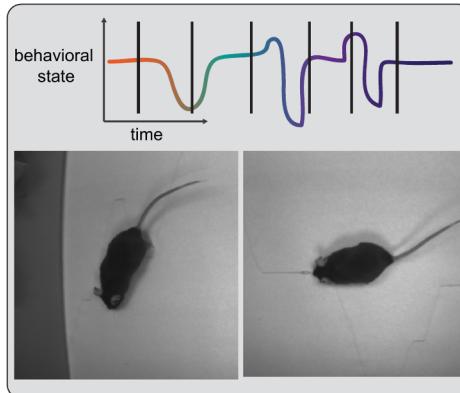
- different sessions reflecting (if the case) varying light conditions, backgrounds, setups, and camera angles (etc).
- different individuals, especially if they look different (i.e. brown + black mice)

3 methods for frame extraction to create a labeled train/test set

Image based clustering (k-means)



Random temporal sampling (uniform)

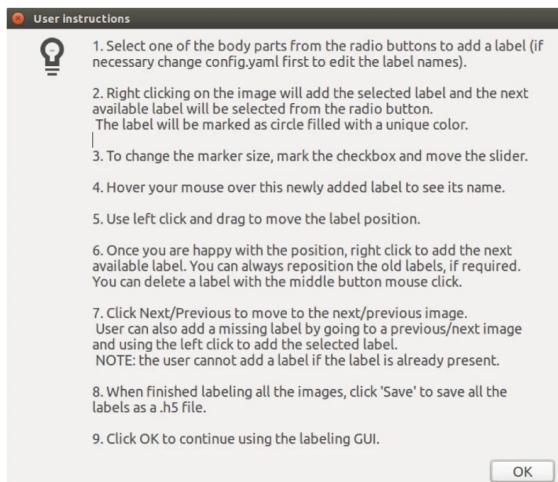
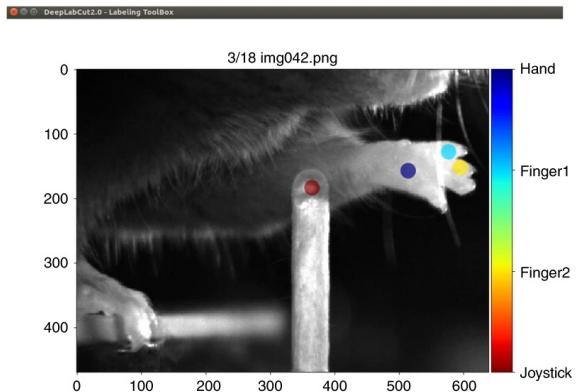
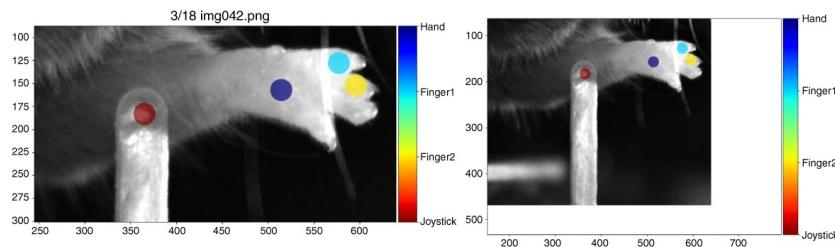
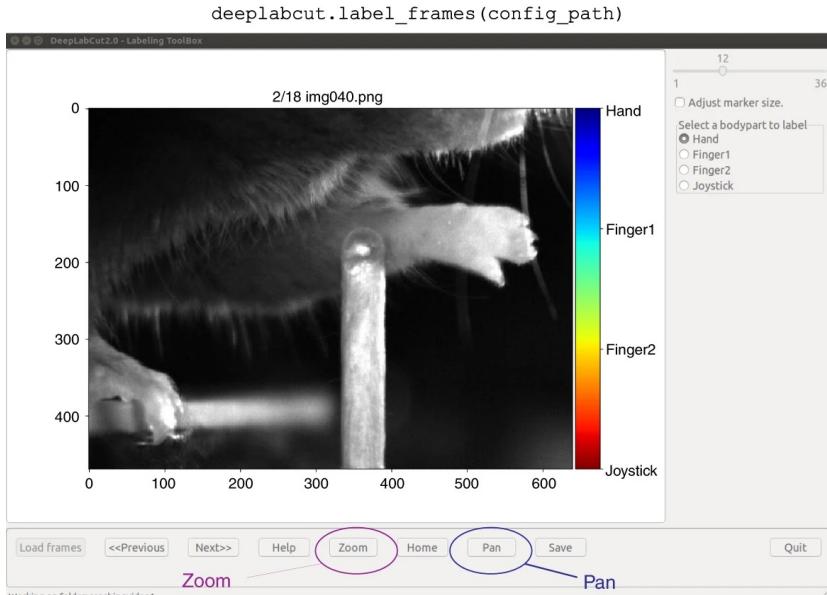


GUI for manual frame grabbing



Label frames: deeplabcut.label_frames(..)

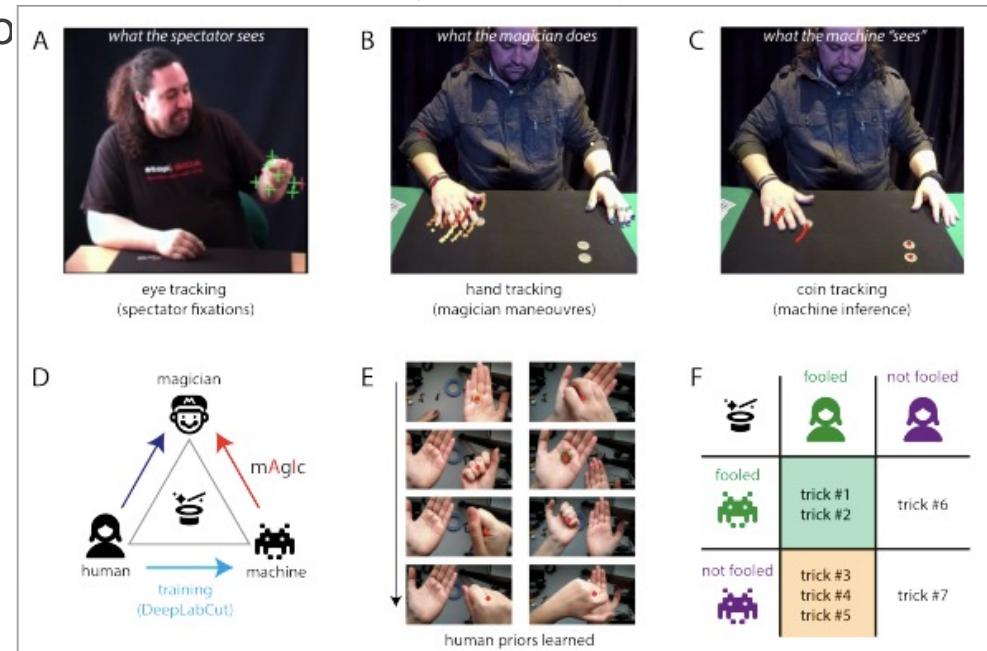
Label frames using the interactive GUI:



What bodyparts can be detected?

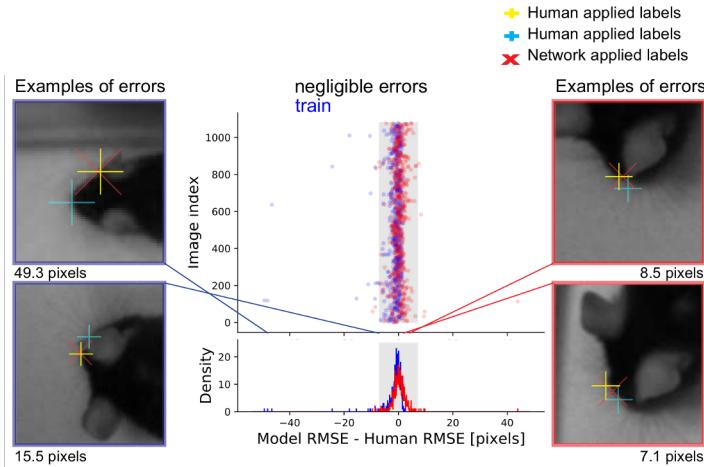
DLC has very large receptive fields (due to deep ResNet), so many visual features + geometric layout can be co

- Salient points (e.g. snout/tail base)
- Texture of objects
- Can teach the network to guess
- Geometric relationship
 - (e.g. left/right ear; whiskers!)



Labeling considerations

- Label (feature locations) **consistently**
- Errors are bad! (especially when only a few images are labeled)
- Left vs. right body parts can be dist. if full animal visible (due to large rec. fields), but flipped labels are problematic
- You can check labels by plotting (do this before training!)
- You can correct wrong labels manually with the labeling GUI



What network & augmentation method?

Network architecture:

- Network backbones: MobileNets V2, ResNets, *DLC-Rnet* & EfficientNets
- Scale (level of deconvolution) -> decrease stride for closely interacting animals
and and when you want high resolution
- Differences: Inference & training speed, resolution

Pre-training (Weights):

- *ImageNet pretrained (the generalist)* -> see Mathis et al. 2021 WACV
- MPII-pose pretrained
- DeepLabCut model zoo

Training method:

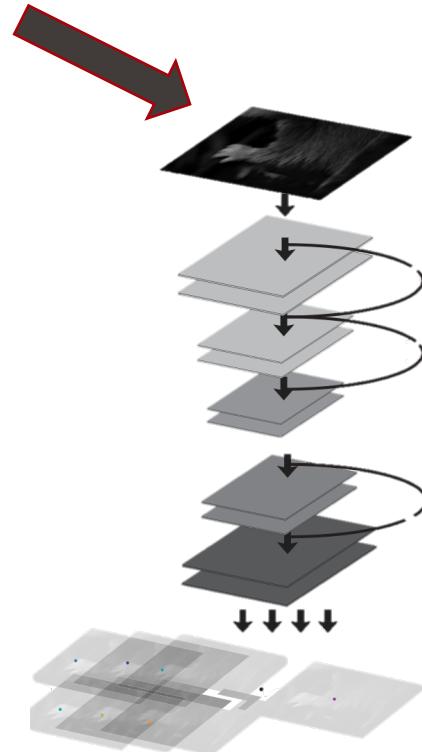
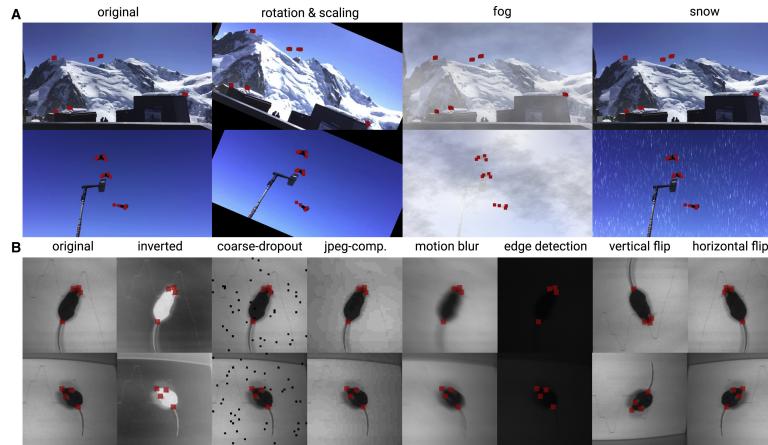
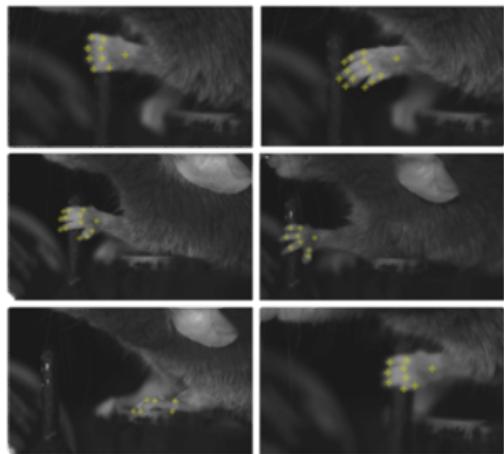
- Learning rate & optimizer (Adam/SGD, batch_size, ...)
- Augmentation transformations (rotation, scaling, ...)

Online demo ... let's start training

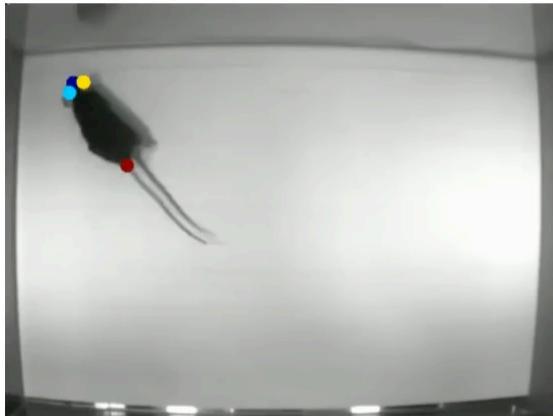
DEMO:

https://github.com/DeepLabCut/DeepLabCut/blob/master/examples/COLAB/COLAB_DEMO_mouse_open_field.ipynb

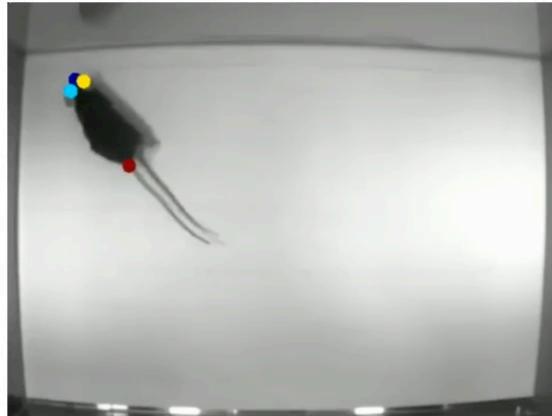
Augmentation



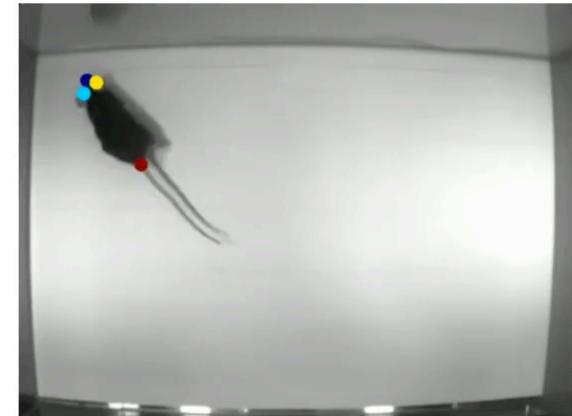
Data Augmentation: how to get the most out of your data!



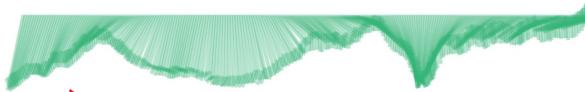
tensorpack



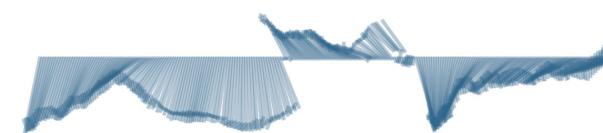
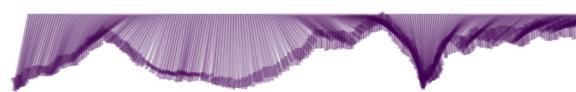
imgaug



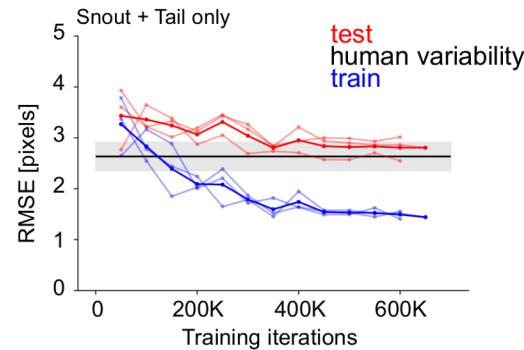
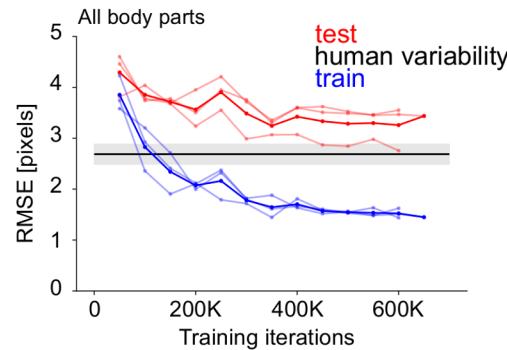
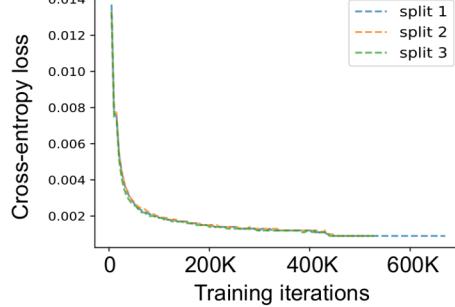
scale-crop



1/17



Evaluation: *Compare test RMSE and ideally your labeling error!*



Evaluation => look at test images!

a Example of the terminal output

```
deeplabcut.evaluate_network('<path of the proj. config file>', shuffle = [1], plotting=True)
Assessing accuracy of shuffle # 1 with 99 % training fraction.
Found the following training snapshots: [(200000, 0)]
You can choose among those for analysis of train/test performance.
Results for 200000 TRAINING iterations: 99 l train error: 4.81 pixels. Test error: 7.02 pixels.
With pccoff of 0.1 train error 3.3 pixels. Test error: 5.2 pixels.
```

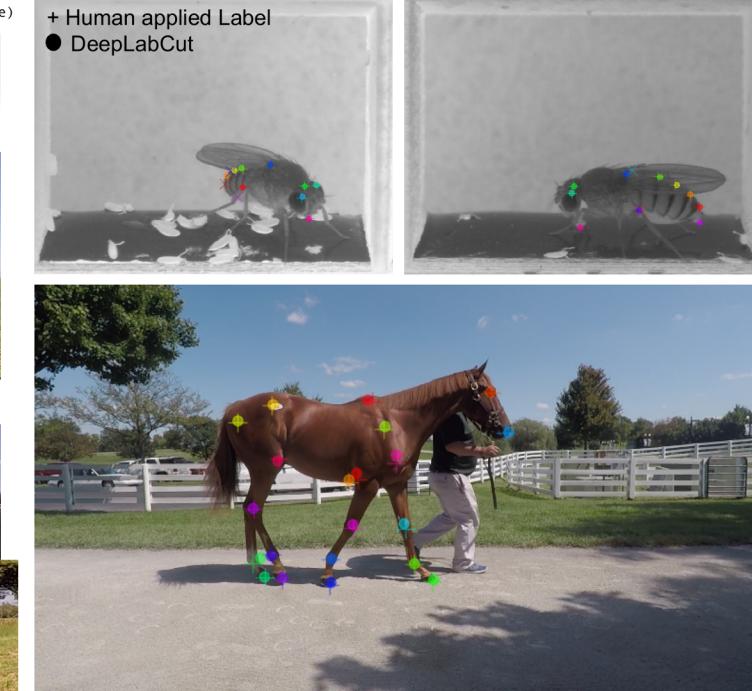
b Examples of **training images** with Human and DeepLabCut labels



c Examples of **test images** with Human and DeepLabCut labels



d Examples of **test images** with Human and DeepLabCut labels



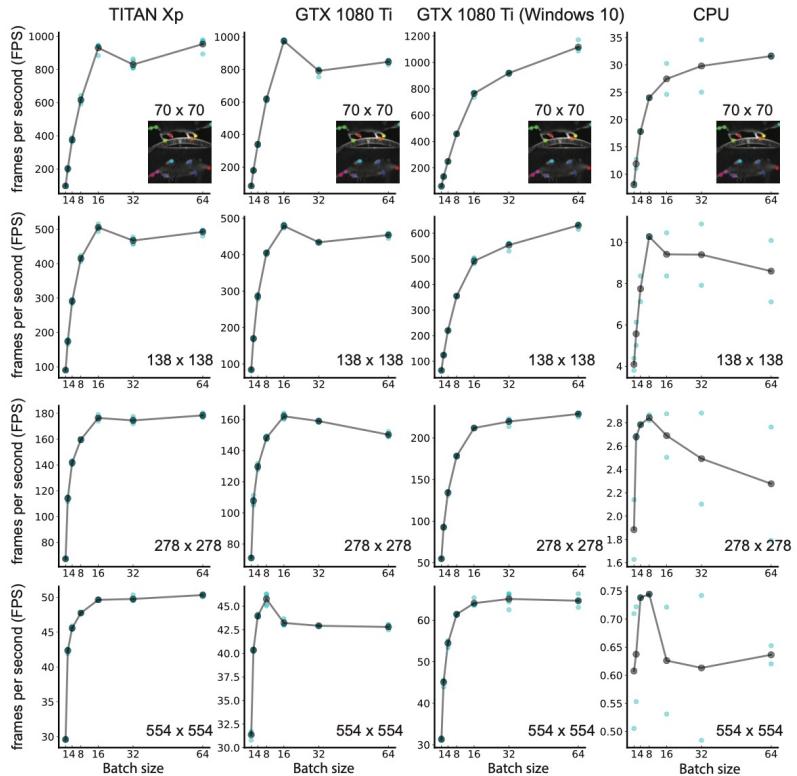
NATURE PROTOCOLS**PROTOCOL****Table 1 | Summary of commands**

Evaluate the trained network (Step 11)	deeplabcut.evaluate_network(config_path)
Video analysis and plotting results (Step 11)	deeplabcut.analyze_videos(config_path, ['path of video 1 or folder', 'path of video2', ...])
Video analysis and plotting results (Step 12)	deeplabcut.plot_trajectories(config_path, ['path of video 1', 'path of video2', ...])
Video analysis and plotting results (Step 13)	deeplabcut.create_labeled_video(config_path, ['path of video 1', 'path of video2', ...])
Refinement: extract outlier frames (Step 14)	deeplabcut.extract_outlier_frames(config_path, ['path of video 1', 'path of video 2'])
Refine labels (Step 15)	deeplabcut.refine_labels(config_path)
Combine datasets (Step 16)	deeplabcut.merge_datasets(config_path)

Video analysis

New videos need **not** be added to the config.yaml!

Set *batch_size* in the config.yaml!

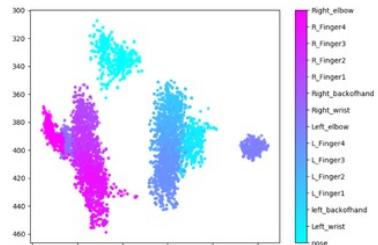


Trajectory visualization

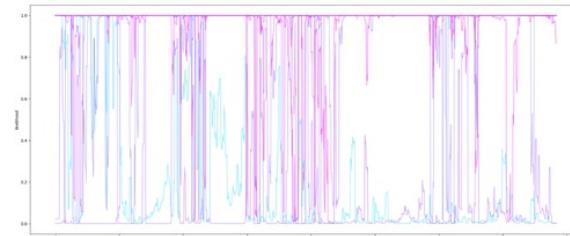
```
deeplabcut.plot_trajectories(config_path,['fullpath/analysis/project/videos/reachingvideo1.avi'])
```



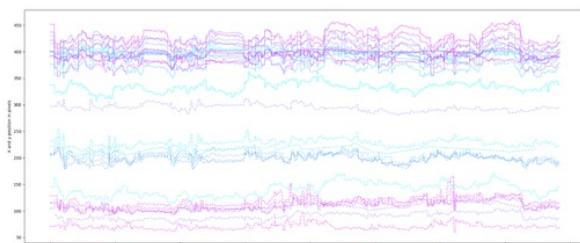
Body parts plotted in space
(over all the frames)



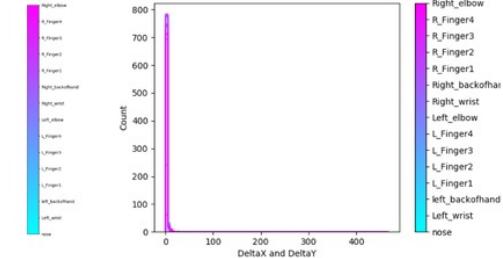
Every body part likelihood over time
(over all the frames)



All body parts across time (frames).
solid lines are Y and dashed lines are X



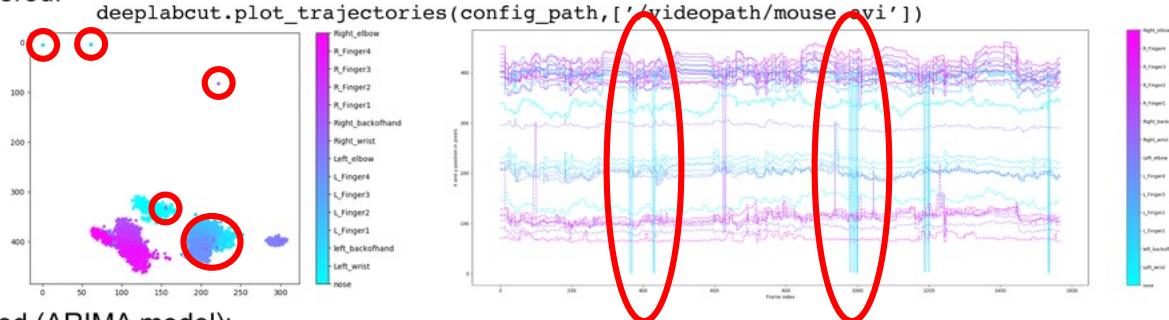
Consecutive coordinate differences
(low values = minimal jumps across frames)



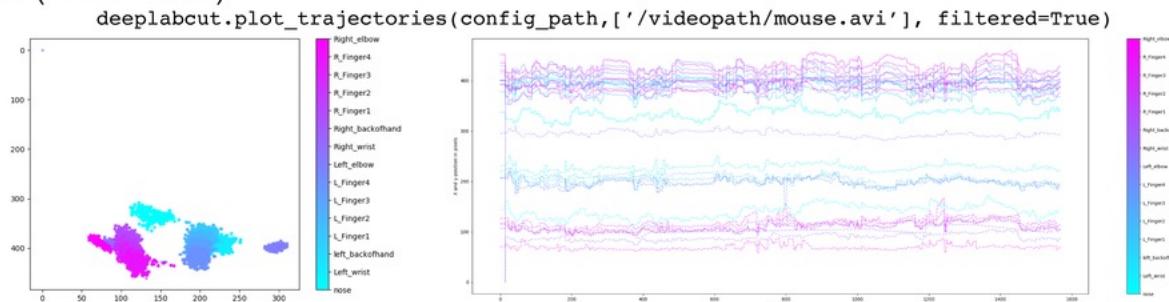
Trajectory filtering

```
deeplabcut.filterpredictions(config_path,['videopath/mouse.avi'], p_bound=0.1, ARdegree=5, MAdegree=1)
```

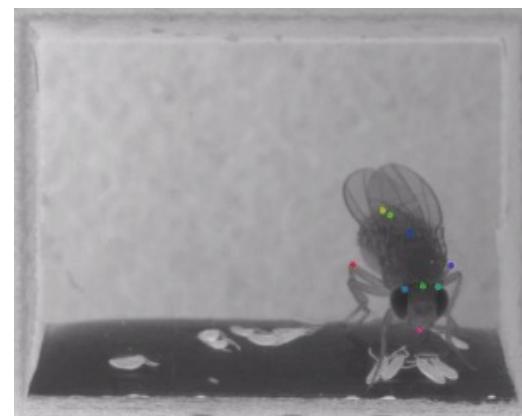
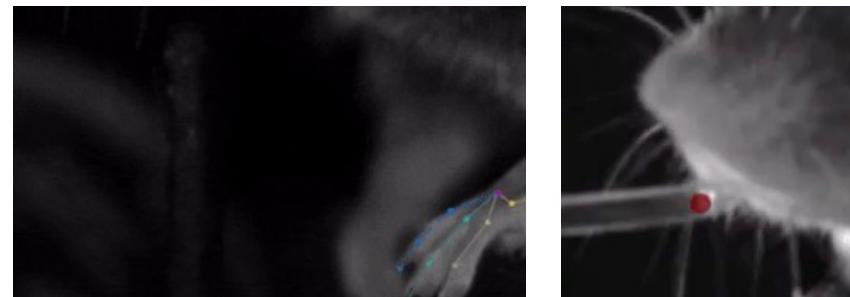
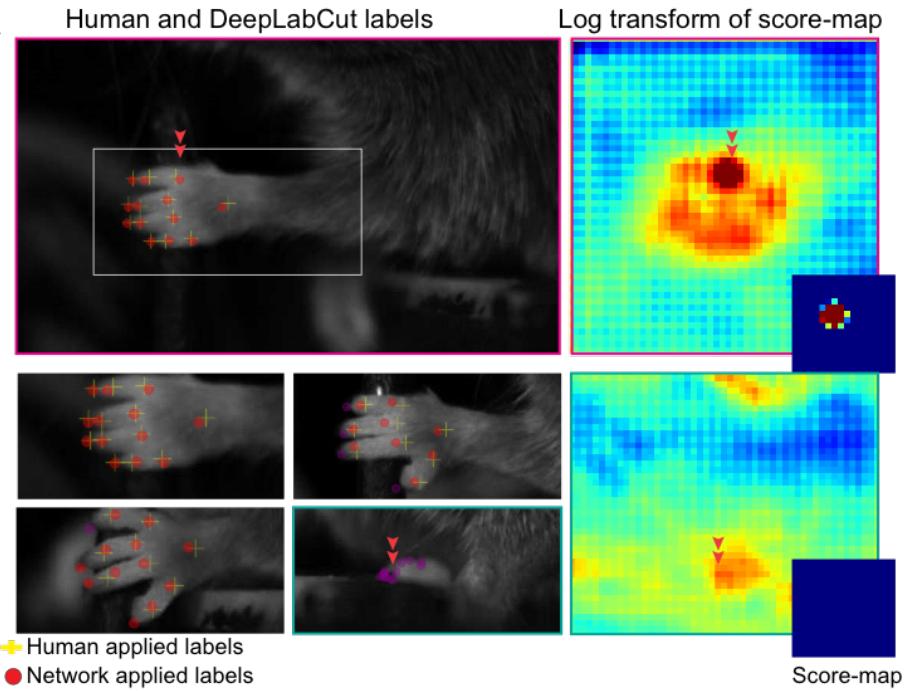
Unfiltered:



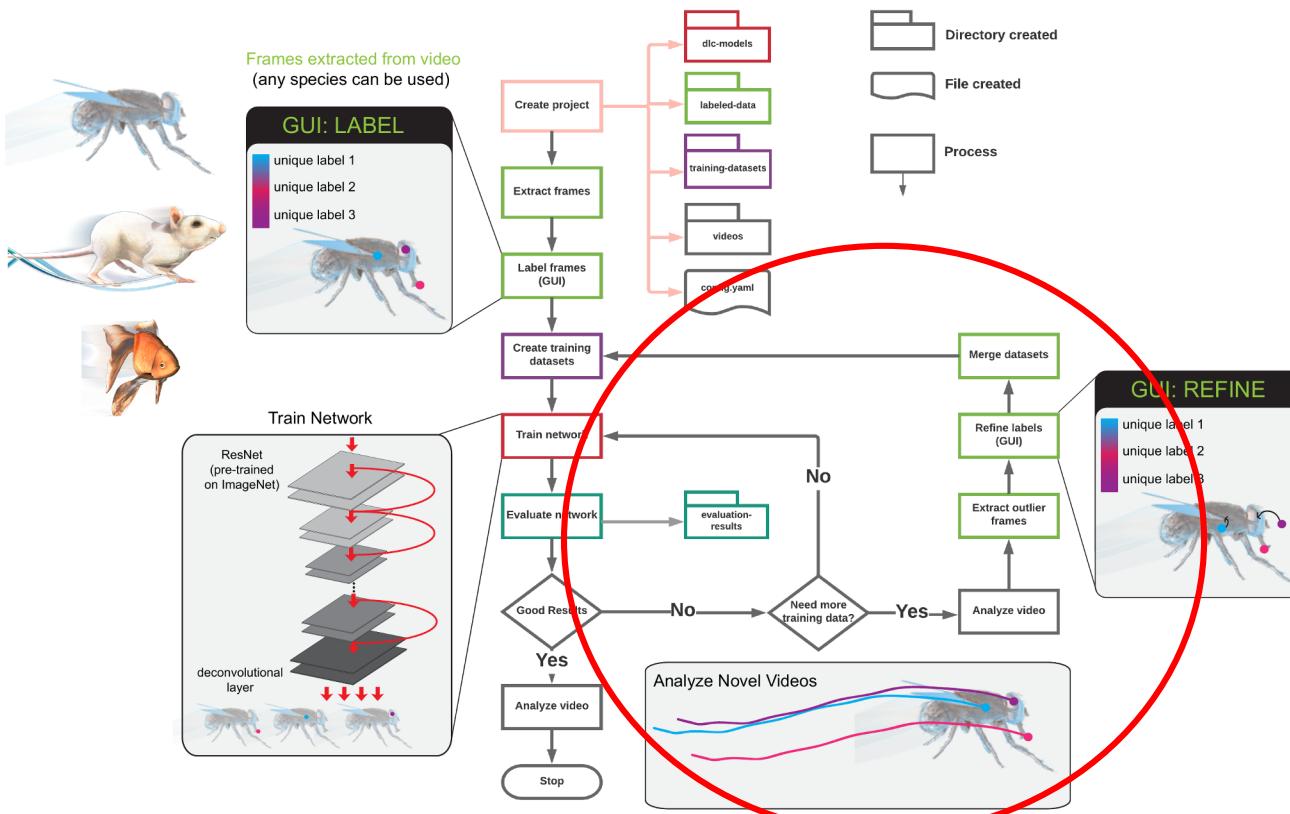
Filtered (ARIMA model):



Score-maps provide network confidence readout



DeepLabCut 2.0 workflow



Active learning

- If there are errors in analyzed videos, you can correct those manually and re-train!
- Note you do not need to correct all errors (presumably some are fixed due to generalization)

Active learning

```
deeplabcut.extract_outlier_frames(config_path,['videofile_path'])
```

• Key arguments

```
outlieralgorithm: 'fitting', 'jump', or 'uncertain'
```

Select frames with jumps between consecutive frames > ϵ_{pixel}



Select frames with confidences < p_{bound}

Select frames deviating from
a statistical model fit to the data

Random sampling

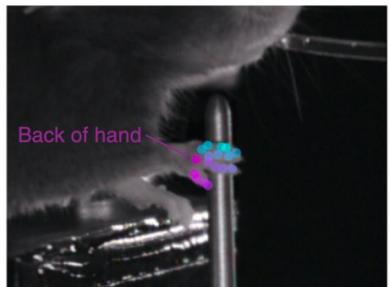
Clustering

```
extractionalgorithm='uniform'
```

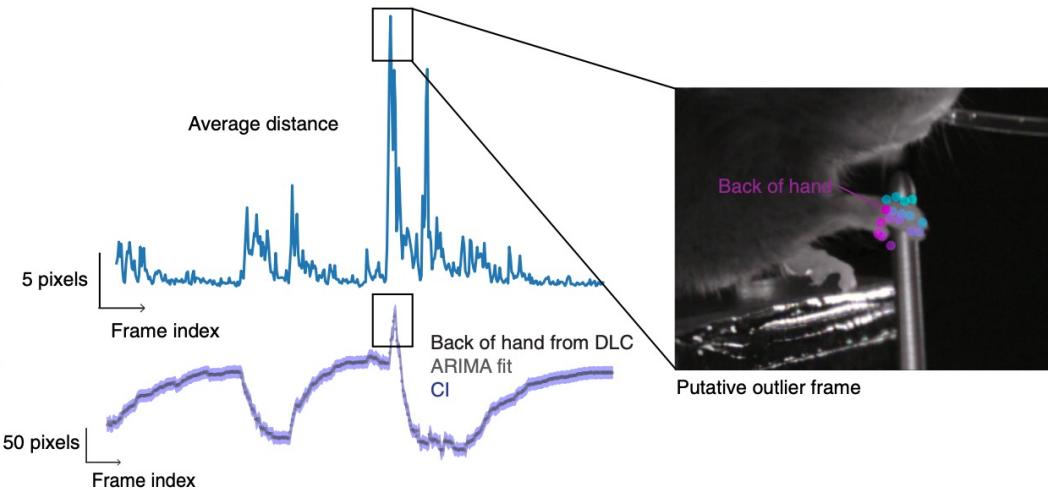
```
extractionalgorithm='k-means'
```

a

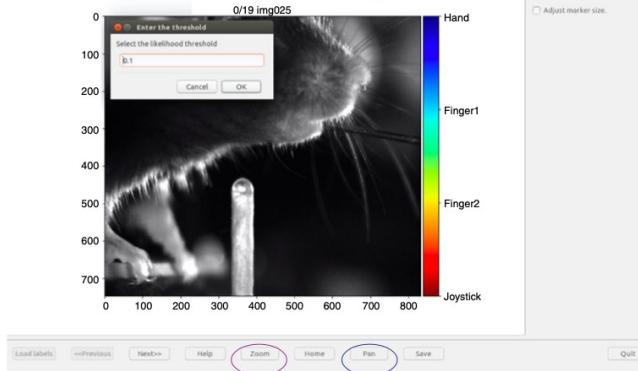
Identification of outlier frames



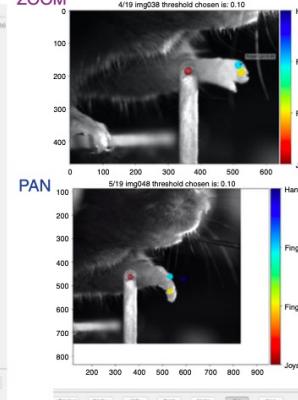
Good frame

**b**

Refinement of DeepLabCut-applied label location(s)

`deeplabcut.refine_labels(config_path)`

ZOOM

**c**

Label removal, i.e., from an occluded point

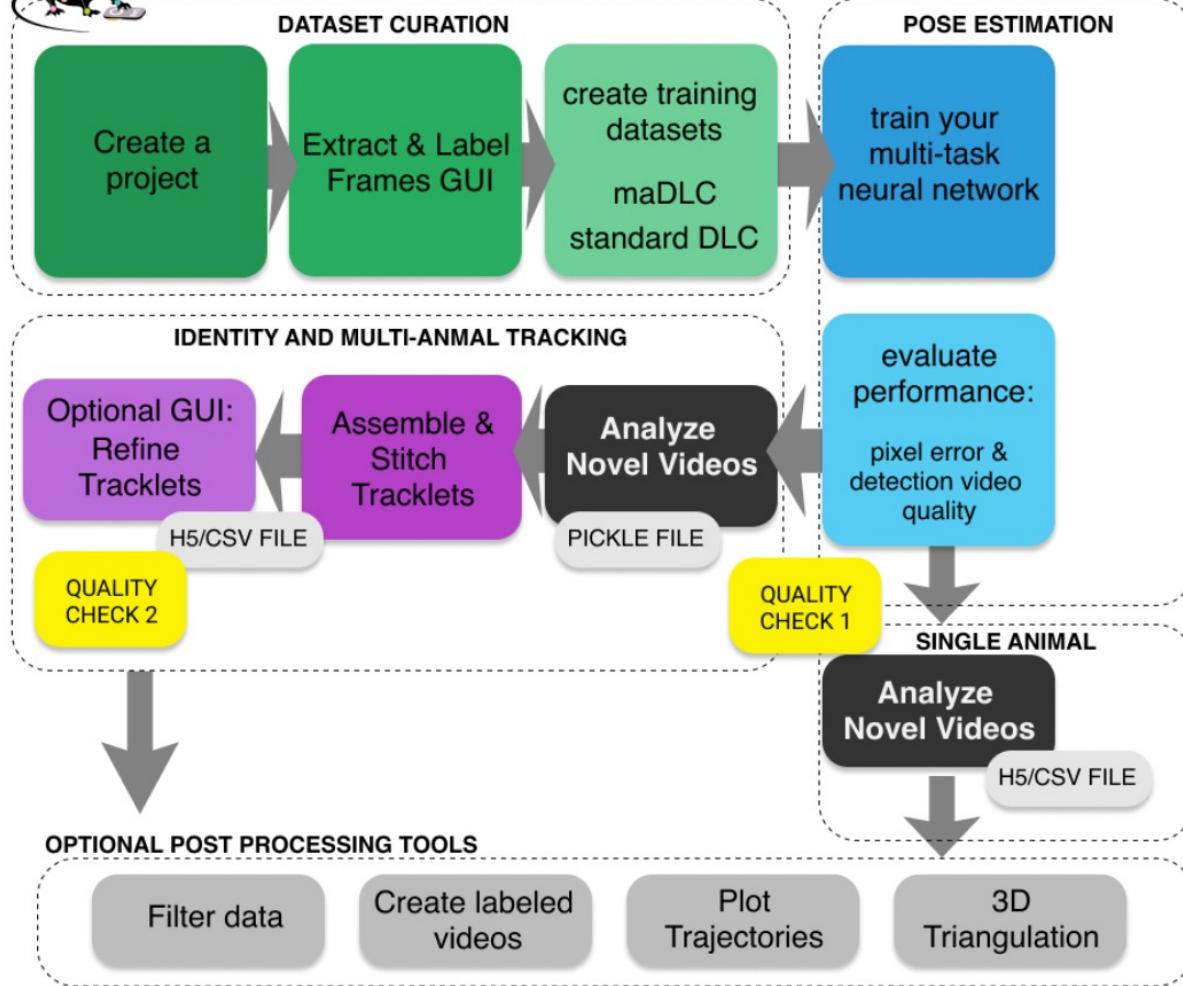


Multi-animal DLC

- **DEMO:**
https://github.com/DeepLabCut/DeepLabCut/blob/master/examples/COLAB/COLAB_3miceDemo.ipynb

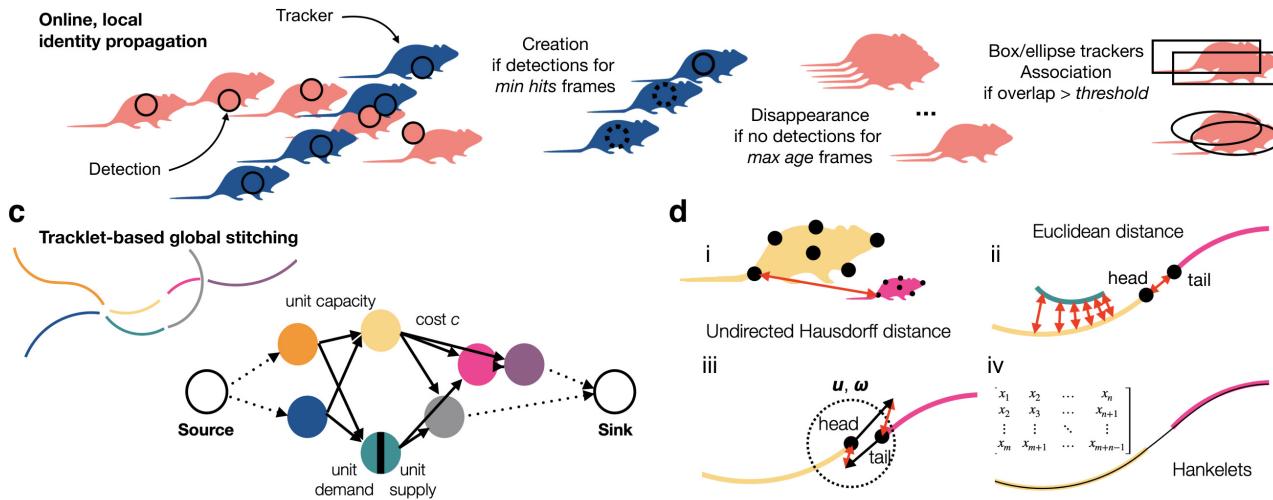


DeepLabCut 2.2+ workflow



Tracking

Stitching



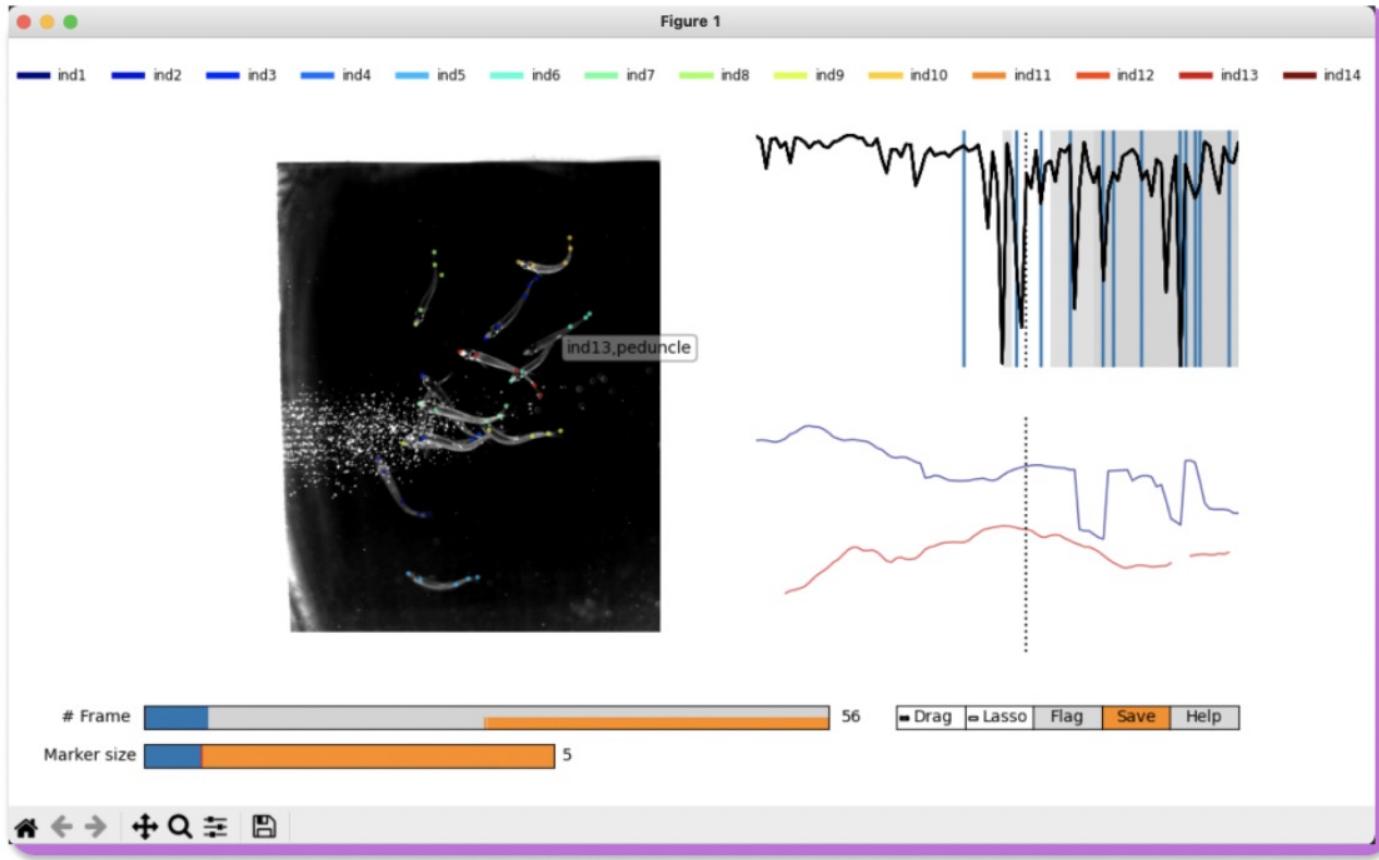
Corresponding commands:

```
dlc.analyze_videos( ... )
```

```
dlc.convert_detections2tracklets(...)
```

```
dlc.stitch_tracklets(...)
```

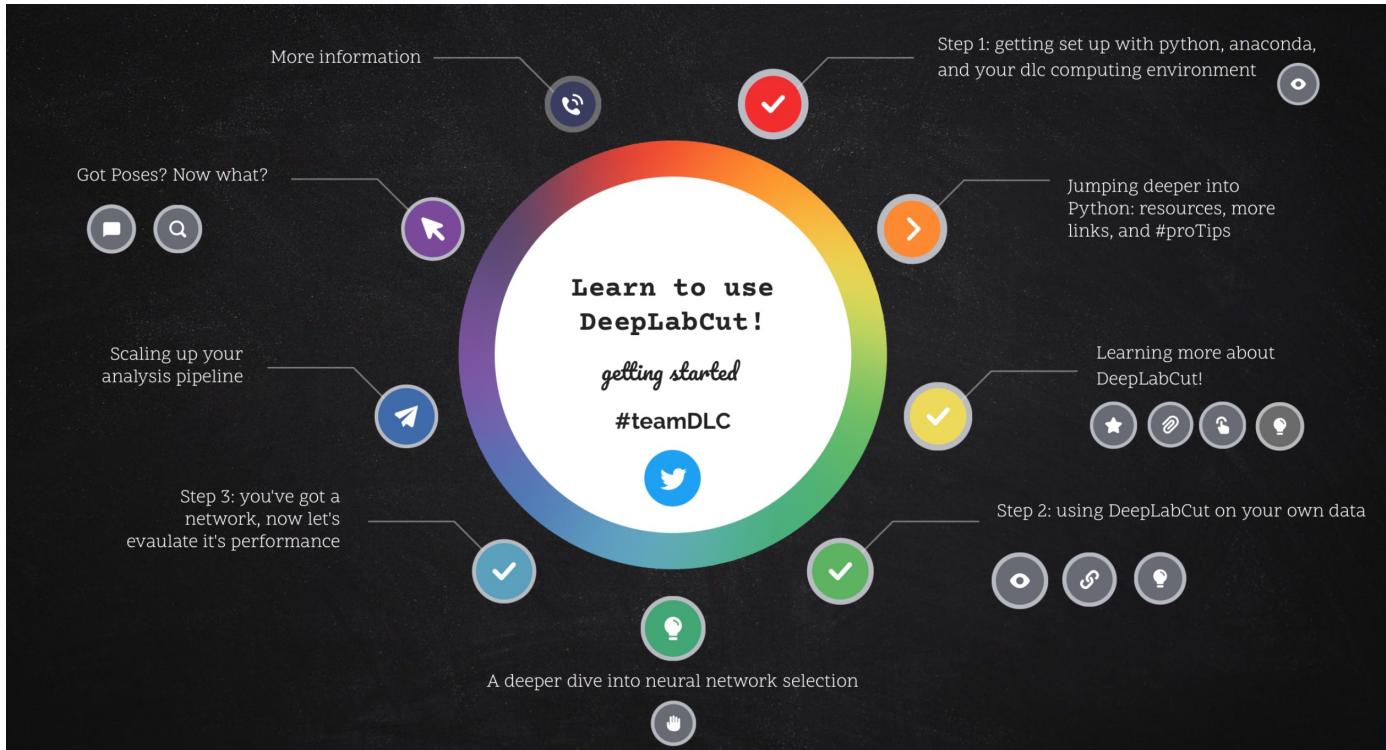
Refine tracklets GUI: deeplabcut.refine_tracklets()



Project GUI demo!



How to use DeepLabCut?



<https://github.com/DeepLabCut/DeepLabCut-Workshop-Materials>

- I would like to get a basic understanding of how DeepLabCut works
- 3D implementation
- How to run DLC successfully and get animal tracking data from it
- Basic functions/introduction
- How to train it for very high accurate digitization
- How to prepare files and train dataset
- Going through the entire process of setting up, training, collecting data
- Basics
- Basic run-through of steps and especially how to fix poor training- do you need to start all over if it doesn't track well?
- General workflow, tips for beginners (if suitable for students). Also, suggestions on hardware
- Things I haven't tried yet, like real time and 3D
- How to use this software to track small, low-contrast animals in low-light settings
- How to track/quantify specific movements and behaviors of a single animal
- How to use COLAB options when tracking two different organisms, best approach when trying to train multiple models from recordings that all have one organism in common but a different second interacting organism
-

Acknowledgements

**Chan
Zuckerberg
Initiative** 

EPFL

THE  KAVLI FOUNDATION

Experimental collaborators:

William Menegas, Feng lab, MIT (Marmoset data),
 Venkatesh N. Murthy, Harvard (Trail tracking, 3 mouse data)
 Catherine Dulac, Harvard (Social behaviors)
 Valentina Di Santo, Stockholm University & George Lauder, Harvard (Fish)
 Amir Patel, U of Cape Town (Cheetah)

Mathis Group, EPFL:

PhD students: Lucas Stoffl (Poet), Alessandro Marin-Vargas (DeepDraw), Alberto Chiappa, **Mu Zhou (DLC)**, Haozhe Qi
 Master's Students/interns: Axel Bisi (Deep Draw), Liza Kozlova, Maxime Vidal (Poet)

DeepLabCut collaborators:

Mackenzie Mathis, EPFL
Jessy Lauer, EPFL (DLC 2.2 / Mathis labs)
Mu Zhou, EPFL (Mathis group)
Ye Shaokai, EPFL (Mathis lab)
 Gary Kane, Harvard (DLC-Live! / M. Mathis lab)
Tanmay Nath, Harvard (DLC 2.0 / M. Mathis lab)
 Matthias Bethge, Uni Tübingen
 Steffen Schneider, Uni Tübingen (M Mathis/Bethge labs)
 Tom Biasi, Harvard (Robustness)
 Mert Yüksekgönül, Uni Tübingen (Faster Inference, Robustness)
+ >65 people on GitHub (Thanks!)

Please get in touch!

alexander.mathis@epfl.ch

Open Source Code: <https://github.com/DeepLabCut/>
Tutorials, Data, Papers etc.: <https://deeplabcut.org>
User forum: <https://forum.image.sc/tag/deeplabcut>

Box 2 | Parameters of interest in the network configuration file, *pose_cfg.yaml*

Please note, there are more parameters that typically never need to be adjusted; they can be found in the default *pose_cfg.yaml* file at https://github.com/AlexEMG/DeepLabCut/blob/master/deeplabcut/pose_cfg.yaml.

- **display_iters**: An integer value representing the period with which the loss is displayed (and stored in *log.csv*).
- **save_iters**: An integer value representing the period with which the checkpoints (weights of the network) are saved. Each snapshot has >90 MB, so not too many should be stored.
- **init_weights**: The weights used for training. Default: <DeepLabCut_path>/Pose_Estimation_Tensorflow/pretrained/resnet_v1_50.ckpt. For ResNet-50 or 101, -- this will be automatically created. The weights can also be changed to restart from a particular snapshot if training is interrupted, e.g., <full path>-snapshot-5000 (with no file-type ending added). This would re-start training from the loaded weights (i.e., after 5,000 training iterations, the counter starts from 0).
- **multi_step**: These are the learning rates and number of training iterations to perform at the specified rate. If the users want to stop before 1 million, they can delete a row and/or change the last value to be the desired stop point.
- **max_input_size**: All images larger with size width × height > max_input_size*max_input_size are not used in training. The default is 1500 to prevent crashing with an out-of-memory exception for very large images. This will depend on your GPU memory capacity. However, we suggest reducing the pixel size as much as possible; see Mathis and Warren²⁷.

The following parameters allow one to change the resolution:

- **global_scale**: All images in the dataset will be rescaled by the following scaling factor to be processed by the convolutional neural network. You can select the optimal scale by cross-validation (see discussion in Mathis et al.¹²). Default is 0.8.
- **pos_dist_thresh**: All locations within this distance threshold (measured in pixels) are considered positive training samples for detection (see discussion in Mathis et al.¹²). Default is 17.

The following parameters modulate the data augmentation. During training, each image will be randomly rescaled within the range [scale_jitter_lo, scale_jitter_up] to augment training:

- **scale_jitter_lo**: 0.5 (default).
- **scale_jitter_up**: 1.5 (default).
- **mirror**: If the training dataset is symmetric around the vertical axis, this Boolean variable allows random augmentation. Default is False.
- **cropping**: Allows automatic cropping of images during training. Default is True.
- **cropratio**: Fraction of training samples that are cropped. Default is 40%.
- **minsize**, **leftwidth**, **rightwidth**, **bottomheight**, **topheight**: These define dimensions and limits for auto-cropping.